

EMC[®] Documentum[®] Foundation Classes

Version 7.1

Development Guide

EMC Corporation
Corporate Headquarters:
Hopkinton, MA 01748-9103
1-508-435-1000
www.EMC.com

Legal Notice

Copyright © 1994-2013 EMC Corporation. All Rights Reserved.

EMC believes the information in this publication is accurate as of its publication date. The information is subject to change without notice.

THE INFORMATION IN THIS PUBLICATION IS PROVIDED "AS IS." EMC CORPORATION MAKES NO REPRESENTATIONS OR WARRANTIES OF ANY KIND WITH RESPECT TO THE INFORMATION IN THIS PUBLICATION, AND SPECIFICALLY DISCLAIMS IMPLIED WARRANTIES OF MERCHANTABILITY OR FITNESS FOR A PARTICULAR PURPOSE.

Use, copying, and distribution of any EMC software described in this publication requires an applicable software license.

For the most up-to-date listing of EMC product names, see EMC Corporation Trademarks on EMC.com. Adobe and Adobe PDF Library are trademarks or registered trademarks of Adobe Systems Inc. in the U.S. and other countries. All other trademarks used herein are the property of their respective owners.

Documentation Feedback

Your opinion matters. We want to hear from you regarding our product documentation. If you have feedback about how we can make our documentation better or easier to use, please send us your feedback directly at IIGDocumentationFeedback@emc.com

Table of Contents

Preface	11
Chapter 1 Getting Started with DFC	13
What Is DFC?	13
Where Is DFC?	14
DFC programming resources	15
DFC documentation set	15
DFC developer support	16
DFC online reference documentation	16
Using dfc.properties to configure DFC	16
BOF and global registry settings	17
Connecting to the global registry	17
Performance tradeoffs	17
Diagnostic settings	18
Diagnostic mode	18
Configuring docbrokers	18
dfc.data.dir	18
Tracing options	18
XML processing options	19
Search options	19
Storage policy options	19
Performance tradeoffs	19
Registry emulation	19
Microsoft Object Linking and Embedding (OLE)	20
Client file system locations	20
New dfc.properties Key to Turn Off Trusted Login	20
Using DFC logging	20
Using DFC from application programs	21
Java	21
Java 2 Security	21
Chapter 2 DFC Programming Basics	23
Client/Server model	23
Packages	23
Interfaces	24
IDfPersistentObject	24
Processing a repository object	25
Failover enabled methods	27
Chapter 3 Sessions and Session Managers	29
Sessions	29
Sharable and Private Sessions	30
Session Managers	30

Getting session managers and sessions.....	30
Instantiating a Session Manager	30
Setting Session Manager Identities.....	31
Getting and releasing sessions.....	32
When you can and cannot release a managed session.....	33
Sessions can be released only once.....	33
Sessions cannot be used once released.....	34
Objects disconnected from sessions.....	34
Related sessions (subconnections).....	34
Original vs. object sessions.....	35
Transactions	35
Configuring sessions using IDfSessionManagerConfig	36
Getting sessions using login tickets	36
Methods for getting login tickets	37
Generating a login ticket using a superuser account.....	37
User authentication	38
Kerberos SSO	38
Overview	39
Procedure to enable Kerberos SSO	39
Configuring the DFC application's service principal name and *.keytab file	39
Mapping the SPN to a user name.....	40
Configuring the DFC application for Kerberos.....	41
Setting up the Kerberos configuration file.....	41
Setting up JAAS configuration.....	42
Performance best practices	43
Known Performance Issue with QUEST's TCP Protocol.....	43
Principal authentication support.....	44
Implementing principal support in your custom application	44
Default classes for demonstrating principal support implementation	45
Maintaining state in a session manager	46
Chapter 4 Creating a Test Application	47
The DfcTutorialFrame class	47
The DfcTutorialApplication class	52
The TutorialGetDirectory class	53
The TutorialSessionManager class.....	54
Running the tutorial application	55
Chapter 5 Working with Objects	57
Understanding repository objects	57
The DFC Object Tutorial Frame	58
Creating a cabinet.....	58
Creating a folder.....	60
Creating a document object	62
Accessing attributes	64
Dumping Attributes	64
Getting a single attribute by name	65
Getting a single attribute by number	67
Setting attributes	68

Setting a single attribute.....	69
Setting an attribute by number	71
Appending a repeating attribute.....	74
Removing an attribute value	76
Getting object content	79
Destroying an object	80
Real-time content deletion.....	82
Chapter 6 Working with Document Operations	83
Understanding documents	83
Virtual documents	83
XML Documents	85
Understanding operations.....	85
Types of operation	86
Basic steps for manipulating documents	87
Steps for manipulating documents.....	87
Details of manipulating documents	88
Obtaining the operation	88
Setting parameters for the operation	88
Adding documents to the operation.....	88
Executing the Operation	89
Processing the results.....	89
Working with nodes	89
Operations for manipulating documents	90
Checking out.....	92
Special considerations for checkout operations	94
Checking out a virtual document.....	94
Checking in.....	96
Special considerations for checkin operations	97
Setting up the operation	98
Processing the checked in documents.....	98
Behavior Change after Failed Save or Checkin	98
Cancelling checkout.....	99
Special considerations for cancel checkout operations	101
Cancel checkout for virtual document.....	101
Importing	103
Special Considerations for Import Operations	104
Setting up the operation	105
XML processing	105
Processing the imported documents	105
Exporting.....	106
Special considerations for export operations.....	108
Copying.....	109
Special considerations for copy operations	111
Moving.....	112
Special considerations for move operations	113
Deleting	114
Special considerations for delete operations	116
Predictive caching	117
Special considerations for predictive caching operations	117
Validating an XML document against a DTD or schema	118
Special considerations for validation operations	118
Performing an XSL transformation of an XML document	119
Special considerations for XML transform operations.....	121
Handling document manipulation errors	122

	The add Method Cannot Create a Node	122
	The execute Method Encounters Errors	122
	Examining Errors After Execution	122
	Using an Operation Monitor to Examine Errors	123
	Operations and transactions	124
Chapter 7	Using the Business Object Framework (BOF)	125
	Overview of BOF	125
	BOF infrastructure	126
	Modules and registries	126
	Dynamic delivery mechanism	127
	Global registry	127
	Global registry user	128
	Accessing the global registry	128
	Service-based Business Objects (SBOs)	128
	SBO introduction	128
	SBO architecture	129
	Implementing SBOs	129
	Stateful and stateless SBOs	130
	Managing Sessions for SBOs	130
	Overview	130
	Structuring Methods to Use Sessions	131
	Managing repository names	131
	Maintaining State Beyond the Life of the SBO	131
	Obtaining Session Manager State Information	132
	Using Transactions With SBOs	132
	SBO Error Handling	134
	SBO Best Practices	134
	Follow the Naming Convention	134
	Don't Reuse SBOs	135
	Make SBOs Stateless	135
	Rely on DFC to Cache Repository Data	135
	Type-based Business Objects (TBOs)	135
	Use of Type-based Business Objects	135
	Creating a TBO	136
	Create a custom repository type	136
	Create the TBO interface	136
	Define the TBO implementation class	137
	Implement methods of IDfBusinessObject	138
	getVersion method	139
	getVendorString method	139
	isCompatible method	139
	supportsFeature method	139
	Code the TBO business logic	139
	Using a TBO from a client application	140
	Using TBOs from SBOs	141
	Getting sessions inside TBOs	142
	Inheritance of TBO methods by repository subtypes without TBOs	142
	Dynamic inheritance	143
	Exploiting dynamic inheritance with TBO reuse	144
	Signatures of Methods to Override	145
	Calling TBOs and SBOs	149
	Calling SBOs	149
	Returning a TBO from an SBO	150
	Calling TBOs	150
	Sample SBO and TBO implementation	151

ITutorialSBO	151
TutorialSBO	151
ITutorialTBO	152
TutorialTBO	153
Deploying the SBO and TBO	154
Calling DFS, JAXB API, or other components from a BOF module	154
Aspects	155
Examples of usage	155
General characteristics of aspects	156
Creating an aspect	156
Creating the aspect interface	156
Creating the aspect class	157
Deploy the customer service aspect	158
TestCustomerServiceAspect	158
Using aspects in a TBO	161
Using DQL with aspects	162
Enabling aspects on object types	162
Default aspects	162
Referencing aspect attributes from DQL	163
Full-text index	163
Object replication	164
Chapter 8 Working with Virtual Documents	165
Understanding virtual documents	165
Setting version labels	166
Getting version labels	169
Creating a virtual document	171
Traversing the virtual document structure	174
Binding to a version label	176
Clearing a version label binding	179
Removing a virtual document child	180
Chapter 9 Support for Other Documentum Functionality	183
Security Services	183
XML	184
Virtual Documents	184
Workflows	184
Document Lifecycles	185
Validation Expressions in Java	185
Unified Login	186
Search Service	186
WORM Support for Filestores	186

List of Figures

Figure 1.	The DFC base tutorial frame.....	56
Figure 2.	The DFC object tutorial frame.....	58
Figure 3.	The DFC operations tutorial frame.....	91
Figure 4.	Basic TBO design.....	137
Figure 5.	TBO design with extended intervening class	138
Figure 6.	Inheritance by object subtype without associated TBO	143
Figure 7.	Design-time dynamic inheritance hierarchies	144
Figure 8.	Runtime dynamic inheritance hierarchies.....	144
Figure 9.	Design-time dynamic inheritance with TBO reuse	145
Figure 10.	Runtime dynamic inheritance with TBO reuse.....	145

List of Tables

Table 1.	DFC operation types and nodes.....	86
Table 2.	Methods to override when implementing TBOs.....	145

Preface

This manual describes EMC Documentum Foundation Classes (DFC). It provides overview and summary information.

For an introduction to other developer resources, refer to [DFC developer support](#), page 16.

Intended audience

This manual is for programmers who understand how to use Java and are generally familiar with the principles of object oriented design.

Revision History

The following changes have been made to this document.

Revision Date	Description
November 2013	Initial publication.

Getting Started with DFC

This chapter introduces DFC. It contains the following major sections:

- [What Is DFC?, page 13](#)
- [Where Is DFC?, page 14](#)
- [DFC programming resources, page 15](#)
- [Using dfc.properties to configure DFC, page 16](#)
- [Using DFC logging, page 20](#)
- [Using DFC from application programs, page 21](#)

What Is DFC?

DFC is a key part of the Documentum software platform. While the main user of DFC is other Documentum software, you can use DFC in any of the following ways:

- Access Documentum functionality from within one of your company's enterprise applications.
For example, your corporate purchasing application can retrieve a contract from your Documentum system.
- Customize or extend products such as Webtop.
For example, you can modify Webtop functionality to implement one of your company's business rules.
- Write a method or procedure for Content Server to execute as part of a workflow or document lifecycle.
For example, the procedure that runs when you promote an XML document might apply a transformation to it and start a workflow to subject the transformed document to a predefined business process.

You can view Documentum functionality as having the following elements:

Repositories	One or more places where you keep the content and associated metadata of your organization's information. The metadata resides in a relational database, and the content resides in various storage elements.
Content Server	Software that manages, protects, and imposes an object oriented structure on the information in repositories. It provides tools for managing the lifecycles of that information and automating processes for manipulating it.
Client programs	Software that provides interfaces between Content Server and end users. The most common clients run on application servers (for example, Webtop).
End Users	People who control, contribute, or use your organization's information. They use a browser to access client programs running on application servers.

In this view of Documentum functionality, Documentum Foundation Classes (DFC) lies between Content Server and clients. Documentum Foundation Services are the primary client interface to the Documentum platform. Documentum Foundation Classes are used for server-side business logic and customization.

DFC is Java based. As a result, client programs that are Java based can interface directly with DFC.

When application developers use DFC, it is usually within the customization model of a Documentum client, though you can also use DFC to develop the methods associated with Content Server functionality, such as document lifecycles.

In the Java application server environment, Documentum client software rests on the foundation provided by the Web Development Kit (WDK). This client has a customization model that allows you to modify the user interface and also implement some business logic. However, the principal tool for adding custom business logic to a Documentum system is to use the Business Object Framework (BOF).

BOF enables you to implement business rules and patterns as reusable elements, called *modules*. The most important modules for application developers are type based objects (TBOs), service based objects (SBOs), and Aspects. Aspect modules are similar to TBOs, but enable you to attach properties and behavior on an instance-by-instance basis, independent of the target object's type.

BOF makes it possible to extend some of DFC's implementation classes. As a result, you can introduce new functionality in such a way that existing programs begin immediately to reflect changes you make to the underlying business logic.

The *Documentum Content Server Fundamentals* manual provides a conceptual explanation of the capabilities of Content Server. DFC provides a framework for accessing those capabilities. Using DFC and BOF makes your code much more likely to survive future architectural changes to the Documentum system.

Where Is DFC?

DFC runs on a Java virtual machine (JVM), which can be on:

- The machine that runs Content Server.

For example, to be called from a method as part of a workflow or document lifecycle, or to be executed on the Java Method Server on the Content Server host.

- A middle-tier system.

For example, on an application server to support WDK or to execute server methods.

For client machines, Documentum provides Documentum Foundation Services (DFS) as the primary support for applications communicating with the Documentum platform.

Note: Refer to the DFC release notes for the supported versions of the JVM. These can change from one minor release to the next.

The *DFC Installation Guide* describes the locations of files that DFC installs. The *config* directory contains several files that are important to DFC's operation.

DFC programming resources

This section provides an overview of the resources available to application developers to help them use DFC.

DFC documentation set

The following documents pertain to DFC or to closely related subjects

- *Documentum System Migration Guide*

Refer to the migration guide if you are upgrading from an earlier version of DFC.

- *DFC Installation Guide*

This guide explains how to install DFC. It includes information about preparing the environment for DFC.

- *DFC Development Guide*

A programmers guide to using DFC to access Documentum functionality. This guide focuses on concepts and overviews. Refer to the Javadocs for details.

- *DFC Online Reference (Javadocs)*

The Javadocs provide detailed reference information about the classes and interfaces that make up DFC. They contain automatically generated information about method signatures and interclass relationships, explanatory comments provided by developers, and a large number of code samples.

- *Documentum Content Server Fundamentals*

A conceptual description of the capabilities of Documentum Content Server and how to use them. The material in this manual is a key to understanding most DFC interfaces.

- *XML Application Development Guide*

A description of the XML-related capabilities of the Documentum server, and an explanation of how to design applications that exploit those capabilities.

DFC developer support

Application developers using DFC can find additional help from EMC Developer Network.

The EMC Developer Network website (<https://community.emc.com/community/edn>) is an extensive resource that every developer should take some time to explore and become familiar with.

The Developer section contains tips, sample code, downloads, a component exchange, and other resources for application developers using DFC, WDK, and other tools.

DFC online reference documentation

The public API for DFC is documented in the Javadocs that ship with the product. You have the option of deploying the Javadocs during DFC installation, or you can download the Javadocs from the EMC Developer Connection web site (<http://developer.emc.com> — search for “DFC Javadocs”). Direct access to undocumented classes or interfaces is not supported.

Using dfc.properties to configure DFC

You can adjust some DFC behaviors. This section describes the `dfc.properties` file, which contains properties compatible with the `java.util.Properties` class.

The `dfc.properties` file enables you to set preferences for how DFC handles certain choices in the course of its execution. The accompanying `dfcfull.properties` file contains documentation of all recognized properties and their default settings. Leave `dfcfull.properties` intact, and copy parts that you want to use into the `dfc.properties` file. The following sections describe the most commonly used groups of properties. Refer to the `dfcfull.properties` file for a complete list.

The DFC installer creates a simple `dfc.properties` file and places it in the classpath. Other EMC installers for products that bundle DFC also create `dfc.properties` file in the appropriate classpath. This file is located in `APP-INF/classes` if you are deploying the EAR file, or in `WEB-INF/classes` if you deploying the WAR file.

If you prefer, you can use a `#include` statement to point to a properties file outside of the web application on the local file system. This can make access to some settings more convenient and allows you to modularize your configuration settings:

```
#include C:\Documentum\config\dfc.properties
```

At a minimum, the `dfc.properties` file must include the following entries:

```
dfc.docbroker.host[0]
```

```
dfc.globalregistry.repository
```

```
dfc.globalregistry.username
```

```
dfc.globalregistry.password
```

The following sections provide some details on settings that your application may require.

BOF and global registry settings

All Documentum installations must designate a global registry to centralize information and functionality. Refer to [Global registry, page 127](#) for information about global registries.

Connecting to the global registry

The `dfc.properties` file contains the following properties that are mandatory for using a global registry.

- `dfc.globalregistry.repository`

The name of the repository. The repository must project to a connection broker that DFC has access to.

- `dfc.globalregistry.username`

The user name part of the credentials that DFC uses to access the global registry. Refer to [Global registry user, page 128](#) for information about how to create this user.

- `dfc.globalregistry.password`

The password part of the credentials that DFC uses to access the global registry. The DFC installer encrypts the password if you supply it. If you want to encrypt the password yourself, use the following instruction at a command prompt:

```
java com.documentum.fc.tools.RegistryPasswordUtils password
```

You can call the `getGlobalRegistryPassword` method to get the global registry password. For example, the following code snippet prints this password:

```
System.out.println(DfPreferences.getInstance().getGlobalRegistryPassword());
```

Note: This method returns the global registry password as it is presented in the `dfc.properties` file. If the password is encrypted in the file, it is encrypted in the return.

The `dfc.properties` file also provides an optional property to resist attempts to obtain unauthorized access to the global registry. For example, the entry

```
dfc.bof.registry.connect.attempt.interval=60
```

sets the minimum interval between connection attempts to the default value of 60 seconds.

Performance tradeoffs

You can adjust the amount of time DFC relies on cached information before checking for consistency between the local cache and the global registry. For example, the entry

```
dfc.bof.cache.currency_check_interval=60
```

sets that interval to the default value of 60 seconds. Because global registry information tends to be relatively static, you might be able to check less frequently in a production environment. On the other hand, you might want to check more frequently in a development environment. The check is inexpensive. If nothing has changed, the check consists of looking at one `vstamp` object. Refer to the *Content Server Object Reference* for information about `vstamp` objects. The name of this entry was formerly `dfc.bof.cacheconsistency.interval`.

Diagnostic settings

DFC provides a number of properties that facilitate diagnosing and solving problems.

Diagnostic mode

DFC can run in diagnostic mode. You can cause this to happen by including the following setting in `dfc.properties`:

```
dfc.diagnostics.resources.enable=true
```

The set of problems that diagnostic mode can help you correct can change without notice. Here are some examples of issues detected by diagnostic mode.:

- Session leaks
- Collection leaks

DFC catches the leaks at garbage collection time. If it finds an unreleased session or an unclosed collection, it places an appropriate message in the log.

Configuring docbrokers

You must set the repeating property `dfc.docbroker.host`, one entry per docbroker. For example,

```
dfc.docbroker.host[0]=docbroker1.yourcompany.com
dfc.docbroker.host[1]=docbroker2.yourcompany.com
```

These settings are described in the *dfcfull.properties* file.

dfc.data.dir

The `dfc.data.dir` setting identifies the directory used by DFC to store files. By default, it is a folder relative to the current working directory of the process running DFC. You can set this to another value.

Tracing options

DFC has extensive tracing support. Trace files can be found in a directory called *logs* under `dfc.data.dir`. For simple tracing, add the following line to `dfc.properties`:

```
dfc.tracing.enable=true
```

That property traces DFC entry calls, return values and parameters.

For more extensive tracing information, add the following properties to `dfc.properties`.

```
dfc.tracing.enable=true
dfc.tracing.verbose=true
dfc.tracing.include_rpcs=true
```

These properties include more details and RPCs sent to the server. Start with the simple trace because the verbose trace produces much more output.

XML processing options

DFC's XML processing is largely controlled by configuration files that define XML applications. Two properties provide additional options. Refer to the *XML Application Development Guide* for information about working with content in XML format.

Search options

DFC supports the search capabilities of Federated Search Services (FS2) with a set of properties. The FS2 installer sets some of these. Most of the others specify diagnostic options or performance tradeoffs.

Storage policy options

Some properties control the way DFC applies storage policy rules. These properties do not change storage policies. They provide diagnostic support and performance tradeoffs.

Performance tradeoffs

Several properties enable you to make tradeoffs between performance and the frequency with which DFC executes certain maintenance tasks.

DFC caches the contents of properties files such as `dfc.properties` or `dbor.properties`. If you change the contents of a properties file, the new value does not take effect until DFC rereads that file. The `dfc.config.check_interval` property specifies the interval in which DFC rereads the DFC configuration file changes in seconds. The default value for this property is 30.

DFC periodically reclaims unused resources. The `dfc.housekeeping.cleanup.interval` property specifies the interval between cleanups. The default value is 7 days.

Some properties described in the [BOF and global registry settings, page 17](#) and [Search options, page 19](#) sections also provide performance tradeoffs.

Registry emulation

DFC uses the `dfc.registry.mode` property to keep track of whether to use a file, rather than the Windows registry, to store certain settings. The DFC installer sets this property to *file* by default.

Microsoft Object Linking and Embedding (OLE)

When importing or exporting Microsoft OLE documents, DFC can automatically search for linked items and import them as individual system objects linked as nodes in a virtual document. On export, the document can be reassembled into a single document with OLE links. In order to enable this behavior, you must configure the registry emulation mode in `dfc.properties`.

By default, the setting is:

```
dfc.registry.mode = file
```

To enable OLE file handling, set the value to:

```
dfc.registry.mode = windows
```

Client file system locations

Several properties record the file system locations of files and directories that DFC uses. The DFC installer sets the values of these properties based upon information that you provide. There is normally no need to modify them after installation.

New `dfc.properties` Key to Turn Off Trusted Login

By default, applications running on the Content Server host are allowed to make repository connections as the installation owner without presenting a password. This is called a trusted login.

If an application, such as Documentum Administrator, that has an explicit login dialog box is installed on a Content Server host, a user can log in as the installation owner without a password using a trusted login.

Documentum 6.7 and later uses a new `dfc.properties` key to turn off trusted logins if you do not want to allow trusted logins through such applications. The key is:

```
dfc.session.allow_trusted_login
```

Setting this key to false requires users always to provide a password, even when logging in as the installation owner.

Using DFC logging

DFC provides diagnostic logging. The logging is controlled by a `log4j` configuration file, which is customarily named `log4j.properties`. The `dmcl` installer as well as other EMC installers place a default version of this file in the same directory where the `dfc.properties` file is created. Though you generally will not need to change this file, you can customize your logging by consulting the `log4j` documentation.

In past releases, DFC used the DMCL library to communicate with the server and provided support for integrating DMCL logging and DFC logging. Since DMCL is no longer used, the features to integrate its tracing into the DFC log are no longer needed.

Using DFC from application programs

You can use Java to interface directly to DFC. On a Windows system you can also use the Microsoft component object model (COM).

The DFC installer installs `dfc.dll` and, for backward compatibility, `dfc.tlb`. Always import `dfc.dll` rather than `dfc.tlb` into your COM programs.

The most convenient way to access DFC on a Windows system depends on the programming language.

Java

From Java, add `dctm.jar` to your classpath. This file contains a manifest, listing all files that your Java execution environment needs access to. The `javac` compiler does not recognize the contents of the manifest, so for compilation you must ensure that the compiler has access to `dfc.jar`. This file contains most Java classes and interfaces that you need to access directly. In some cases you may have to give the compiler access to other files described in the manifest. In your Java source code, import the classes and interfaces you want to use.

Ensure that the classpath points to the config directory.

Java 2 Security

If you plan to use the Java 2 security architecture for your installation, you will need to create security policy files to specify the permissions allowed to the classes in your installation. To help you develop your own security policy files, an example policy template file is included in the DFC installation. The file specifies the permissions required for the DFC classes. You will need to merge these permissions to your policy file, and replace the variables in the file with the values in your installation to allow the DFC classes to function. You will specify the policy file that you create at the startup of your Java Virtual Machine (JVM). For a stand-alone installation, you will specify the file in the command line startup arguments for your JVM. If you are deploying on an application server, refer to the application server documentation for the specifics of specifying Java 2 security policies. The DFC example policy template file is named `dfc.example.java.policy`, and is found in the `dfc.jar` file installed in the `DOCUMENTUM_SHARED` directory (by default, in Windows systems, `C:\Program Files\Documentum\Shared`).

If you deploy DFC on WebSphere Application Server 7.x, you must set the `dfc.security.enforce_existing_policy` JVM runtime system property before policy configuration.

Name	Description	Type	Required or Not	Default
dfc.security.enforce_existing_policy	<p>Specifies whether to enforce the existing security policy settings or not.</p> <p>true: DFC enforces the existing security policy settings. And thus, policy configuration is skipped.</p> <p>false: DFC does not enforces the existing security policy settings. And thus, you can initialize policy configuration to update the security settings.</p>	boolean	Yes	false

DFC Programming Basics

This chapter describes some common patterns of DFC application code. It contains the following sections.

- [Client/Server model, page 23](#)
- [Packages, page 23](#)
- [Interfaces, page 24](#)
- [Processing a repository object, page 25](#)
- [Failover enabled methods, page 27](#)

Client/Server model

The Documentum architecture generally follows the client/server model. DFC-based programs are client programs, even if they run on the same machine as a Documentum server. DFC encapsulates its client functionality in the IDfClient interface, which serves as the entry point for DFC code. IDfClient handles basic details of connecting to Documentum servers.

You obtain an IDfClient object by calling the static method DfClientX.getLocalClient().

An IDfSession object represents a connection with the Documentum server and provides services related to that session. DFC programmers create new Documentum objects or obtain references to existing Documentum objects through the methods of IDfSession.

To get a session, first create an IDfSessionManager by calling IDfClient.newSessionManager(). Next, get the session from the session manager using the procedure described in the sections and session managers below. For more information about sessions, refer to [Chapter 3, Sessions and Session Managers](#).

Packages

DFC comprises a number of packages, that is, sets of related classes and interfaces.

- The names of DFC Java classes begin with Df (for example, DfCollectionX).
- Names of interfaces begin with IDf (for example, IDfSessionManager).

Interfaces expose DFC's public methods. Each interface contains a set of related methods. The Javadocs describe each package and its purpose.

Note:

- The `com.documentum.operations` package and the `IDfSysObject` interface in the `com.documentum.fc.client` package have some methods for the same basic tasks (for example, checkin, checkout). In these cases, the `IDfSysObject` methods are mostly for internal use and for supporting legacy applications. The methods in the operations package perform the corresponding tasks at a higher level. For example, they keep track of client-side files and implement Content Server XML functionality.
- The `IDfClientX` is the correct interface for accessing factory methods (all of its `getXxx` methods, except for those dealing with the DFC version or trace levels).

The DFC interfaces form a hierarchy; some derive methods and constants from others. Use the Tree link from the home page of the DFC online reference (see [DFC online reference documentation, page 16](#)) to examine the interface hierarchy. Click any interface to go to its definition.

Each interface inherits the methods and constants of the interfaces above it in the hierarchy. For example, `IDfPersistentObject` has a `save` method. `IDfSysObject` is below `IDfPersistentObject` in the hierarchy, so it inherits the `save` method. You can call the `save` method of an object of type `IDfSysObject`.

Interfaces

Because DFC is large and complex, and because its underlying implementation is subject to change, you should use DFC's public interfaces.

Tip: DFC provides factory methods to instantiate objects that implement specified DFC interfaces. If you bypass these methods to instantiate implementation classes directly, your programs may fail to work properly, because the factory methods sometimes do more than simply instantiate the default implementation class. For most DFC programming, the only implementation classes you should instantiate directly are `DfClientX` and the exception classes (`DfException` and its subclasses).

DFC does not generally support direct access to, replacement of, or extension of its implementation classes. The principal exception to these rules is the *Business Object Framework* (BOF). For more information about BOF, refer to [Chapter 7, Using the Business Object Framework \(BOF\)](#).

IDfPersistentObject

An `IDfPersistentObject` corresponds to a persistent object in a repository. With DFC you usually don't create objects directly. Instead, you obtain objects by calling factory methods that have `IDfPersistentObject` as their return type.



Caution: If the return value of a factory method has type `IDfPersistentObject`, you can cast it to an appropriate interface (for example, `IDfDocument` if the returned object implements

that interface). Do not cast it to an implementation class (for example, DfDocument). Doing so produces a ClassCastException.

Processing a repository object

The following general procedure may help to clarify the DFC approach. This example is not meant to give complete details, but to provide an overview of a typical transaction. More detail on the process is provided in subsequent chapters.

To process a repository object:

1. Obtain an IDfClientX object. For example, execute the following Java code:

```
IDfClientX cx = new DfClientX();
```

2. Obtain an IDfClient object by calling the getLocalClient method of the IDfClientX object. For example, execute the following Java code:

```
IDfClient c = cx.getLocalClient();
```

The IDfClient object must reside in the same process as the Documentum client library, DMCL.

3. Obtain a session manager by calling the newSessionManager method of the IDfClient object. For example, execute the following Java code:

```
IDfSessionManager sm = c.newSessionManager();
```

4. Use the session manager to obtain a session with the repository, that is, a reference to an object that implements the IDfSession interface. For example, execute the following Java code:

```
IDfSession s = sm.getSession();
```

Refer to [Sharable and Private Sessions, page 30](#) for information about the difference between the getSession and newSession methods of IDfSessionManager.

5. If you do not have a reference to the Documentum object, call an IDfSession method (for example, newObject or getObjectByQualification) to create an object or to obtain a reference to an existing object.
6. Use routines of the operations package to manipulate the object, that is, to check it out, check it in, and so forth. For simplicity, the example below does not use the operations package. (Refer to [Chapter 6, Working with Document Operations](#) for examples that use operations).
7. Release the session.

Example 2-1. Processing a repository object

The following fragment from a Java program that uses DFC illustrates this procedure.

```
IDfClientX cx = new DfClientX();           //Step 1

IDfClient client = cx.getLocalClient();     //Step 2

IDfSessionManager sMgr = client.newSessionManager(); //Step 3
IDfLoginInfo loginInfo = clientX.getLoginInfo();
loginInfo.setUser( "Mary" );
loginInfo.setPassword( "ganDalF" );
loginInfo.setDomain( "" );
sMgr.setIdentity( strRepositoryName, loginInfo );
```

```
IDfSession session = sMgr.getSession( strRepoName );    //Step 4

try {
    IDfDocument document =
        (IDfDocument) session.newObject( "dm_document" ); //Step 5

    document.setObjectName( "Report on Wizards" );    //Step 6
    document.save();

}
finally {
    sMgr.release( session );                          //Step 7
}
```

Steps 1 through 4 obtain an IDfSession object, which encapsulates a session for this application program with the specified repository.

The following example shows how to obtain an IDfClient object in Visual Basic:

```
Dim myclient As IDfClient          'Steps 1-2
Dim myclientx As new DfClientX
Set myclient = myclientX.getLocalClient
```

Step 5 creates an IDfDocument object. The return type of the newObject method is IDfPersistentObject. You must cast the returned object to IDfDocument in order to use methods that are specific to documents.

Step 6 of the example code sets the document object name and saves it.

Note that the return type of the newObject method is IDfPersistentObject. The program explicitly casts the return value to IDfDocument, then uses the object's save method, which IDfDocument inherits from IDfPersistentObject. This is an example of interface inheritance, which is an important part of DFC programming. The interfaces that correspond to repository types mimic the repository type hierarchy.

Step 7 releases the session, that is, places it back under the control of the session manager, sMgr. The session manager will most likely return the same session the next time the application calls sMgr.getSession.

Most DFC methods report errors by throwing a DfException object. Java code like that in the above example normally appears within a try/catch/finally block, with an error handler in the catch block.

Tip: When writing code that calls DFC, it is a best practice to include a *finally* block to ensure that you release storage and sessions.

Failover enabled methods

Content Server installations can be set up and configured for high availability. One part of high availability is that the Java Method Server (JMS) is set up to failover to another JMS. If you are writing custom methods, and you wish them to be failover enabled, here are some minimum requirements:

- Logic that modifies repository objects should be done in a transaction.
 - If failover occurs, failed or incomplete methods will be rerun, from the beginning, on the backup JVM.
 - Use a transaction to enable roll back of any object changes in the event of a failure.
- Do not write code that assumes file, directory paths, or other resources exist if they only exist on one host.
- Test failover of custom method code by configuring multiple Content Server/JMS and shutting down one while the method is running.

Sessions and Session Managers

This chapter describes how to get, use, and release sessions, which enable your application to connect to a repository and access repository objects.

Note: If you are programming in the WDK environment, be sure to refer to Managing Sessions in *Web Development Kit Development Guide* for information on session management techniques and methods specific to WDK.

This chapter contains the following major sections:

- [Sessions, page 29](#)
- [Session Managers, page 30](#)
- [Getting session managers and sessions, page 30](#)
- [Objects disconnected from sessions, page 34](#)
- [Related sessions \(subconnections\), page 34](#)
- [Original vs. object sessions, page 35](#)
- [Transactions, page 35](#)
- [Configuring sessions using IDfSessionManagerConfig, page 36](#)
- [Getting sessions using login tickets, page 36](#)
- [Principal authentication support, page 44](#)

Sessions

To do any work in a repository, you must first get a session on the repository. A session (IDfSession) maintains a connection to a repository, and gives access to objects in the repository for a specific logical user whose credentials are authenticated before establishing the connection. The IDfSession interface provides a large number of methods for examining and modifying the session itself, the repository, and its objects, as well as for using transactions (refer to IDfSession in the Javadocs for a complete reference).

Sharable and Private Sessions

A sharable session can be shared by threads or applications. A private session is a session that is not shared. To get a shared session, use `IDfSessionManager.getSession`. To get a private session, use `IDfSessionManager.newSession`.

Session Managers

A session manager (`IDfSessionManager`) manages sessions for a single user on one or more repositories. You create a session manager using the `DfClient.newSessionManager` factory method.

The session manager serves as a factory for generating new `IDfSession` objects using the `IDfSessionManager.newSession` method. Immediately after using the session to do work in the repository, the application should release the session using the `IDfSessionManager.release()` method in a *finally* clause. The session initially remains available to be reclaimed by the session manager instance that released it, and simultaneously, the connection associated with that session will be placed to the connection pool where it can be shared.

The `IDfSessionManager.getSession` method checks for an available shared session, and, if one is available, uses it instead of creating a new session. This makes for efficient use of content server connections, which are an extremely expensive resource, in a web programming environment where a large number of sessions are required.

Getting session managers and sessions

The following sections describe how to instantiate a session manager, set its identities, and use it to get and release sessions.

Instantiating a Session Manager

The following sample method instantiates and returns a session manager object without setting any identities.

Example 3-1. Instantiating a session manager without identities

```
import com.documentum.com.DfClientX;
import com.documentum.fc.client.IDfClient;
import com.documentum.fc.client.IDfSessionManager;
import com.documentum.fc.common.DfException;

/**
 * Creates a IDfSessionManager with no prepopulated identities
 */
public static IDfSessionManager getSessionManager() throws Exception
{
    // Create a client object using a factory method in DfClientX.

    DfClientX clientx = new DfClientX();
```

```

IDfClient client = clientx.getLocalClient();

// Call a factory method to create the session manager.

IDfSessionManager sessionMgr = client.newSessionManager();
return sessionMgr;
}

```

In this example, you directly instantiate a DfClientX object. Based on that object, you are able to use factory methods to create the other objects required to interact with the repository.

Setting Session Manager Identities

To set a session manager identity, encapsulate a set of user credentials in an IDfLoginInfo object and pass this with the repository name to the IDfSessionManager.setIdentity method. In simple cases, where the session manager will be limited to providing sessions for a single repository, or where the login credentials for the user is the same in all repositories, you can set a single identity to IDfLoginInfo.ALL_DOCBASES (= *). This causes the session manager to map any repository name for which there is no specific identity defined to a default set of login credentials.

```

/**
 * Creates a simplest-case IDfSessionManager
 * The user in this case is assumed to have the same login
 * credentials in any available repository
 */
public static IDfSessionManager getSessionManager
    (String userName, String password) throws Exception
{
    // create a client object using a factory method in DfClientX

    DfClientX clientx = new DfClientX();
    IDfClient client = clientx.getLocalClient();

    // call a factory method to create the session manager

    IDfSessionManager sessionMgr = client.newSessionManager();

    // create an IDfLoginInfo object and set its fields
    IDfLoginInfo loginInfo = clientx.getLoginInfo();
    loginInfo.setUser(userName);
    loginInfo.setPassword(password);

    // set single identity for all docbases
    sessionMgr.setIdentity(IDfSessionManager.ALL_DOCBASES, loginInfo);
    return sessionMgr;
}

```

If the session manager has multiple identities, you can add these lazily, as sessions are requested. The following method adds an identity to a session manager, stored in the session manager referred to by the Java instance variable sessionMgr. If there is already an identity set for the repository name, setIdentity will throw a DfServiceException. To allow your method to overwrite existing identities, you can check for the identity (using hasIdentity) and clear it (using clearIdentity) before calling setIdentity.

```

public void addIdentity
    (String repository, String userName, String password) throws DfServiceException
{
    // create an IDfLoginInfo object and set its fields

```

```
IDfLoginInfo loginInfo = this.clientx.getLoginInfo();
loginInfo.setUser(userName);
loginInfo.setPassword(password);

if (sessionMgr.hasIdentity(repository))
{
    sessionMgr.clearIdentity(repository);
}
sessionMgr.setIdentity(repository, loginInfo);
}
```

Note that `setIdentity` does not validate the repository name nor authenticate the user credentials. This normally is not done until the application requests a session using the `getSession` or `newSession` method; however, you can authenticate the credentials stored in the identity without requesting a session using the `IDfSessionManager.authenticate` method. The `authenticate` method, like `getSession` and `newSession`, uses an identity stored in the session manager object, and throws an exception if the user does not have access to the requested repository.

Getting and releasing sessions

To get and assume ownership of a managed session, call `IDfSessionManager.getSession` (to get a sharable session) or `newSession` (to get a private session), passing a repository name as a `String`. If there is no identity set for the repository, `getSession` throws a `DfIdentityException`. If a user name and password stored in the identity fail to authenticate, `getSession` will throw a `DfAuthenticationException`.

To release a session, call `IDfSessionManager.release()`, passing it the session reference. You should release a session as soon as the work for which it is immediately required is complete, and you should release the session in a `finally` block to ensure that it gets released in the event of an error or exception. This discipline will help you avoid problems with session leaks, in which sessions are created and remain open. It is safer and much more efficient to release sessions as soon as they are no longer actively being used and get new sessions as needed than to store sessions for later use.

It is important to note that this pattern for releasing sessions applies only when the session was obtained using a factory method of `IDfSessionManager`.

```
public static void demoSessionGetRelease
(String repository, String userName, String password) throws DfException
{
    IDfSession mySession = null;

    // Get a session manager with a single identity.

    DfClientX clientx = new DfClientX();
    IDfClient client = clientx.getLocalClient();
    mySessMgr = client.newSessionManager();
    IDfLoginInfo logininfo = clientx.getLoginInfo();
    loginInfo.setUser(userName);
    loginInfo.setPassword(password);
    mySessMgr.setIdentity(repository, loginInfo);

    // Get a session using a factory method of session manager.

    IDfSession mySession = mySessMgr.getSession(repository);
    try
    {
        /*
        * Insert code that performs tasks in the repository.
        */
    }
}
```



```

*/
}
finally
{
    mySessMgr.release(mySession);
}
}

```

Some legacy applications might still be getting sessions from IDfClient directly, which is now discouraged in favor of the method above. In those cases, you should call IDfSession.disconnect().

When you can and cannot release a managed session

You can only release a managed session that was obtained using a factory method of the session manager; that is, IDfSessionManager.getSession or IDfSessionManager.newSession. Getting a session in this way implies ownership, and confers responsibility for releasing the session.

If you get a reference to an existing session, which might for example be stored as a data member of a typed object, no ownership is implied, and you cannot release the session. This would be the case if you obtained the session using IDfTypedObject.getSession.

The following snippet demonstrates these two cases:

```

// session is owned
IDfSession session = sessionManager.getSession("myRepository");
IDfSysObject object = session.getObject(objectId);
mySbo.doSomething(object);
sessionManager.release(session);

public class MySbo
{
    private void doSomething( IDfSysObject object )
    {
        // session is not owned
        IDfSession session = object.getSession();
        IDfQuery query = new DfClientX().getQuery();
        query.setDQL("select r_object_id from dm_relation where ...");
        IDfCollection co = query.execute(session, IDfQuery.DF_READ_QUERY );
        ...
    }
}

```

Note that the session instantiated using the session manager factory method is released as soon as the calling code has finished using it. The session obtained by object.getSession, which is in a sense only borrowed, is not released, and in fact cannot be released in this method.

Sessions can be released only once

Once a session is released, you cannot release or disconnect it again using the same session reference. The following code will throw a runtime exception:

```

IDfSession session = sessionManager.getSession("myRepository");
sessionManager.release(session);
sessionManager.release(session); // throws runtime exception

```

Sessions cannot be used once released

Once you have released a session, you cannot use the session reference again to do anything with the session (such as getting an object).

```
IdfSession session = sessionManager.getSession("myRepository");
sessionManager.release(session);
session.getObject(objectId); // throws runtime exception
```

Objects disconnected from sessions

There may be cases when it is preferable to release a session, but store the object(s) retrieved through the session for later use. DFC handles this for you transparently. If you attempt to do something that requires a session with an object after the session with which the object was obtained has been released, DFC will automatically reopen a session for the object. You can get a reference to this session by calling `object.getSession`.

Users of versions of DFC earlier than release 6, should note that you no longer need to call `setSessionManager` to explicitly disconnect objects, nor is there a need to use `beginClientControl/endClientControl` to temporarily turn off management of a session.

Related sessions (subconnections)

It is sometimes necessary to get a session to another repository based on an existing session. For example, you may be writing a method that takes an object and opens a session on another repository using the user credentials that were used to obtain the object. There are two recommended approaches to solving this problem.

The first is to get a session using the session manager associated with the original session. This new session is a peer of the original session, and your code is responsible for releasing both sessions.

```
IdfSession peerSession =
    session.getSessionManager().getSession(repository2Name);
try
{
    doSomethingUsingRepository2(peerSession);
}
finally
{
    session.getSessionManager().releaseSession(peerSession);
}
```

A second approach is to use a related session (subconnection), obtained by calling `IdfSession.getRelatedSession`. The lifetime of the related session will be dependent on the lifetime of the original session; you cannot explicitly release it.

```
IdfSession relatedSession = session.getRelatedSession(repository2Name);
doSomethingUsingRepository2(relatedSession);
```

Both of these techniques allow you to make use of identities stored in the session manager.

Users of earlier versions of DFC should note that using `setDocbaseScope` for creating subconnections is no longer recommended. Use one of the preceding techniques instead.

Original vs. object sessions

The original session is the session used by an application to obtain an object.

The object session is a session to the object's repository that DFC used internally to get the object.

Generally, the original session and object session are the same.

`IDfTypedObject.getOriginalSession()` returns the original session.

`IDfTypedObject.getObjectSession()` returns the object session.

`IDfTypedObject.getSession()` is provided for compatibility purposes and returns the original session. It is equivalent to `IDfTypedObject.getObjectSession()`.

Transactions

DFC supports transactions at the session manager level and at the session level. A transaction at the session manager level includes operations on any sessions obtained by a thread using `IDfSessionManager.newSession()` or `IDfSessionManager.getSession` after the transaction is started (See `IDfSessionManager.beginTransaction()` in the DFC Javadoc) and before it completes the transaction (see `IDfSessionManager.commitTransaction()` and `IDfSessionManager.abortTransaction()`).

A transaction at the session level includes operations on the session that occur after the transaction begins (see `IDfSession.beginTransaction()`) and occur before it completes (see `IDfSession.commitTrans()` and `IDfSession.abortTrans()`). Previous versions of DFC did not support calling `beginTransaction()` on a session obtained from a session manager. This restriction has been removed. The code below shows how a TBO can use a session-level transaction.

```
public class MyTBO
{
    protected void doSave() throws DfException
    {
        boolean txStartedHere = false;
        if ( !getObjectSession().isTransactionActive() )
        {
            getObjectSession().beginTransaction();
            txStartedHere = true;
        }
        try
        {
            doSomething();           // Do something that requires transactions
            if ( txStartedHere )
                getObjectSession().commitTrans();
        }
        finally
        {
            if ( txStartedHere && getObjectSession().isTransactionActive() )
                getObjectSession().abortTrans();
        }
    }
}
```

Configuring sessions using IDfSessionManagerConfig

You can configure common session settings stored in the session configuration objects using the `IDfSessionManagerConfig` interface. You can get an object of this type using the `IDfSessionManager.getConfig` method. You can set attributes related to locale, time zone, dynamic groups, and application codes. Generally the settings that you apply will be applied to future sessions, not to sessions that exist at the time the settings are applied. The following snippet demonstrates the use `IDfSessionManagerConfig`:

```
// get session manager
IDfClient client = new DfClientX().getLocalClient();
IDfSessionManager sm = client.newSessionManager();

// configure session manager
// settings common among future sessions
sm.getConfig().setLocale("en_US");
sm.getConfig().addApplicationCode("finance");

// continue setting up session manager
IDfLoginInfo li = new DfClientX().getLoginInfo();
li.setUser(user);
li.setPassword(password);
sm.setIdentity("bank", li);

// now you can get sessions and continue processing
IDfSession session = sessionManager.getSession("myRepository");
```

Getting sessions using login tickets

A login ticket is a token string that you can pass in place of a password to obtain a repository session. You generate a login ticket using methods of an `IDfSession` object. There are two main use cases to consider.

The first use is to provide additional sessions for a user who already has an authenticated session. This technique is typically employed when a web application that already has a Documentum session needs to build a link to a WDK-based application. This enables a user who is already authenticated to link to the second application without going through an additional login screen. The ticket in this case is typically embedded in a URL. For documentation of this technique, see *Ticketed Login* in the *Web Development Kit Development Guide*.

The second use is to allow a session authenticated for a superuser to grant other users access to the repository. This strategy can be used in a workflow method, in which the method has to perform a task on the part of a user without requesting the user password. It is also used in JEE principal authentication support, which allows a single login to the Web server and the Content Server. For information on how to use this technique in WDK applications, see *J2EE Principal Authentication* in *Web Development Kit Development Guide*. For information on building support for principal authentication in custom (non-WDK) web applications, see [Principal authentication support, page 44](#).

For further information on login ticket behavior and server configuration dependencies, see *Login Tickets* in *Content Server Fundamentals*. For information on related server configuration settings see *Configuring Login Tickets* in *Content Server Administrators Guide*.

Methods for getting login tickets

The `IDfSession` interface defines three methods for generating login tickets: `getLoginTicket`, `getLoginTicketForUser`, and `getLoginTicketEx`.

- The `getLoginTicket` method does not require that the session used to obtain the ticket be a superuser, and it generates a ticket that can be used only by the same user who owned the session that requested the ticket.
- The `getLoginTicketEx` requires that the session used to obtain a ticket be a superuser, and it can be used to generate a ticket for any user. It has parameters that can be used to supplement or override server settings: specifically the scope of the ticket (whether its use is restricted to a specific server or repository, or whether it can be used globally); the time in minutes that the ticket will remain valid after it is created; whether the ticket is for a single use; and if so, the name of the server that it is restricted to.
- The `getLoginTicketForUser` requires that the session used to obtain a ticket belong to a superuser, and it can be used to generate a ticket for any user. It uses default server settings for ticket scope, ticket expiration time, and whether the ticket is for single use.

Generating a login ticket using a superuser account

The `IDfSession.getLoginTicketEx` method allows you to obtain tickets that can be used to log in any user using a superuser account. The following sample assumes that you have already instantiated a session manager (`IDfSessionManager adminSessionMgr`), and set an identity in `adminSessionMgr` for a repository using credentials of a superuser account on that repository. The sample then obtains a session for the repository from `adminSessionMgr`, and returns a login ticket. For information on instantiating session managers and setting identities, see [Getting session managers and sessions](#), page 30.

```
/**
 * Obtains a login ticket for userName from a superuser session
 *
 */
public String dispenseTicket(String repository, String userName)
    throws DfException
{
    // This assumes we already have a session manager (IDfSessionManager)
    // with an identity set for repository with superuser credentials.
    session = adminSessionMgr.getSession(repository);
    try
    {
        String ticket = session.getLoginTicketForUser(userName);
        return ticket;
    }
    finally
    {
        adminSessionMgr.release(session);
    }
}
```

To get a session for the user using this login ticket, you pass the ticket in place of the user's password when setting the identity for the user's session manager. The following sample assumes that you have already instantiated a session manager for the user.

```
public IDfSession getSessionWithTicket
    (String repository, String userName) throws DfException
{
    // get a ticket using the preceding sample method

    String ticket = dispenseTicket(repository, userName);

    // set an identity in the user session manager using the
    // ticket in place of the password

    DfClientX clientx = new DfClientX();
    IDfLoginInfo loginInfo = clientx.getLoginInfo();
    loginInfo.setUser(userName);
    loginInfo.setPassword(ticket);

    // if an identity for this repository already exists, replace it silently

    if (userSessionMgr.hasIdentity(repository))
    {
        userSessionMgr.clearIdentity(repository);
    }
    userSessionMgr.setIdentity(repository, loginInfo);

    // get the session and return it

    IDfSession session = userSessionMgr.getSession(repository);
    return session;
}
```

User authentication

You can use either Kerberos SSO or Java Enterprise Edition (JEE) principal authentication for user authentication:

- [Kerberos SSO, page 38](#)
- [Principal authentication support, page 44](#)

Kerberos SSO

These topics are included:

- [Overview, page 39](#)
- [Procedure to enable Kerberos SSO, page 39](#)
- [Configuring the DFC application's service principal name and *.keytab file, page 39](#)
- [Configuring the DFC application for Kerberos, page 41](#)

Overview

EMC Documentum supports Kerberos secure Single-Sign-On (SSO) using Microsoft Active Server Domain Services for Kerberos Key Distribution Center (KDC) services in the following ways:

- In a single domain.
- In one-way and two-way trusts between multiple domains in the same forest only; that is, cross-forest trusts are not supported.

Procedure to enable Kerberos SSO

Make sure that you have configured the following components:

- (Required for cross-domain support only) One-way or two-way trusts between all applicable domains in the same forest
- Kerberos SSO on Content Server

Note: For more information, see the *Documentum Content Server Administration and Configuration Guide*.

1. Register the DFC application's service principal name (SPN) in the Active Directory and generate a keytab file. See [Configuring the DFC application's service principal name and *.keytab file, page 39](#).
2. Enable the DFC application for Kerberos. See [Configuring the DFC application for Kerberos, page 41](#).

Configuring the DFC application's service principal name and *.keytab file

To enable authentication of the DFC application on the Kerberos Key Distribution Center (KDC), register the DFC application's service principal name (SPN) on the Active Server KDC using the Microsoft `ktpass` utility. A Kerberos SPN uniquely identifies a service that uses Kerberos authentication. In this case, the service is your DFC application. Executing the `ktpass` utility also generates a `*.keytab` file. The `*.keytab` file contains name/value pairs consisting of an SPN and a long-term key derived from a password. Both the DFC application and the KDC must be able to access the `*.keytab` file. You copy the `*.keytab` file to the DFC application machine (the machine where the Kerberos service ticket (ST) is validated) and specify the location of the `*.keytab` file in the JAAS configuration.

Note: Although the `*.keytab` file is usually used on non-Windows machines, DFC leverages the `*.keytab` file to improve network performance by eliminating Kerberos authentication communication between Windows machines and the KDC.

In some cases, you can register the SPNs of more than one DFC application to the same account. For example, in load-balanced environments support for Kerberos can be achieved by joining all load-balanced nodes into a single account and assigning a single SPN to the cluster. If access to the service is required through a different SPN (for example, based on the service host IP address rather than the load balancer name), then this SPN can also be registered with the same account. The

following procedure describes the main steps for registering an SPN using a one-to-one mapping between the DFC application's SPN and user account, or a many-to-one mapping in which multiple SPNs are registered to one user account.

To configure the SPN and keytab file (main steps):

1. Create a new user (or use an existing one) for the DFC application in the Active Directory.
Note: Make sure to enable delegation trust for the service accounts who create the SPNs.
2. Map the DFC application's SPN to a user and generate the *.keytab file. See [Mapping the SPN to a user name](#), page 40.

Mapping the SPN to a user name

The recommended SPN format for DFC applications is:

```
DFC/<host>:<port>@<REALM>
```

EMC recommends using a host name rather than an IP address as the host string. For example, myhost.mydomain.com. realm is the name of the Kerberos realm, which is defined in the Kerberos configuration file (see [Setting up the Kerberos configuration file](#), page 41).

To map the SPN to a user name:

Note:

- By default, Windows Server 2008 R2 SP1 does not support DES-related ciphers (for example, DES-CBC-MD5). For more information about DES-related ciphers on Windows Server 2008, see [http://technet.microsoft.com/en-us/library/dd560670\(v=WS.10\).aspx](http://technet.microsoft.com/en-us/library/dd560670(v=WS.10).aspx).
- For the ktpass utility syntax, see <http://technet.microsoft.com/en-us/library/cc753771%28v=WS.10%29.aspx>.

1. Perform one of the following tasks:

- To map the SPN to a user name using a one-to-one mapping, execute the ktpass utility as follows:

Note: For a one-to-one mapping, do not map the same SPN to more than one user account.

```
ktpass /pass <password> -out <keytab_file> -princ <SPN>  
-crypto <crypto_type> +DumpSalt -ptype KRB5_NT_PRINCIPAL +desOnly  
/mapOp set /mapUser <user_name> /target <domain_controller>
```

- To map multiple SPNs to a user name using many-to-one mapping, perform the following steps:

- a. Execute the ktpass utility as follows:

```
ktpass /pass <password> -out <keytab_file> -princ <SPN>  
-crypto <crypto_type> +DumpSalt -ptype KRB5_NT_PRINCIPAL +desOnly  
/mapOp set /mapUser <user_name> /target <domain_controller>
```

Remember the salt string and the key version number (vno) because you need to use them in [Step c](#).

- b. To map the next SPN to the same user account, execute the setspn utility as follows:

```
setspn -A <SPN> <user_name>
```


- c. Execute `ktpass` utility for the second SPN without setting with the same user as follows:

Note: Use the `salt` and key version number (`kvno`) that were displayed as the output in [Step a](#).

```
ktpass /pass <password> -out <keytab_file> -princ <SPN>
-crypto <crypto_type> +DumpSalt -ptype KRB5_NT_PRINCIPAL +desOnly
/mapOp set +RawSalt <salt> -in <keytab_file> -kvno <vno>
```

- d. Repeat [Step b](#) and [Step c](#) for each additional SPN.

Configuring the DFC application for Kerberos

To enable Kerberos support for the DFC application, perform these tasks:

- [Setting up the Kerberos configuration file, page 41](#)
- [Setting up JAAS configuration, page 42](#)
- [Performance best practices, page 43](#)

Setting up the Kerberos configuration file

The Kerberos configuration file (typically named `krb5.ini`) specifies Kerberos settings such as the KDC address and default realm name. The realm name includes the KDC and administration server addresses.

```
[libdefaults]
    default_realm = <REALM>
    forwardable = true
    ticket_lifetime = 24h
    clocks skew = 72000
    default_tkt_enctypes =
    default_tgs_enctypes =

[realms]
    <REALM> = {
        kdc = <kdc_server_ip>
        admin_server = <admin_server_ip>
    }

[domain_realm]
    <domain> = <REALM>

[logging]
    default = c:\kdc.log
    kdc = c:\kdc.log

[appdefaults]
    autologin = true
    forward = true
    forwardable = true
    encrypt = true
```

<code><kdc_server_ip></code>	The IP address of the KDC server.
<code><admin_server_ip></code>	The IP address of the Administration server.

<code><domain></code>	The domain in which the DFC application's SPN exists.
<code><REALM></code>	The realm name. For example: SRV01.COM

Setting up JAAS configuration

The Java Authentication and Authorization Service (JAAS) configuration file entry contains JAAS specific settings such as the `<LoginContext>` name (which is also the name of the configuration entry), settings for the Kerberos login module, the DFC application's SPN, and the location of the `*.keytab` file. Copy the generated `*.keytab` file to the location specified in the JAAS configuration on the DFC application's machine.

The location and format of the JAAS configuration settings might be different for each application. You can specify the configuration file setting in a JVM command-line parameter; for example:

```
-Djava.security.auth.login.config=
```

In a `jaas.conf` file (for example, located in `C:/jaas.conf`), specify the following format for the JAAS configuration:

```
<loginContext>
{
    <LoginModule> required
    debug=true
    principal="<SPN>"
    realm="<REALM>"
    refreshKrb5Config=true
    noTGT=true
    useKeyTab=true
    storeKey=true
    doNotPrompt=true
    useTicketCache=false
    isInitiator=false
    keyTab="<dfcuser_keytab_path>";
};
```

where:

<code><loginContext></code>	<p>Corresponds to the DFC application's SPN. You replace separator characters with hyphen characters and omit the @REALM segment in the SPN. For example, the following <code>LoginContext</code> is derived from the corresponding SPN:</p> <ul style="list-style-type: none"> • <code>LoginContext</code>: DFC-myhost-mydomain-com-8080 • <code>SPN</code>: DFC/myhost.mydomain.com:8080@SRV01.COM
<code><LoginModule></code>	<p>Specify the Kerberos login module to be used to perform user authentication:</p> <ul style="list-style-type: none"> • For single-domain support only:

	<pre>com.sun.security.auth.module. Krb5LoginModule</pre> <ul style="list-style-type: none"> For both multi- and single-domain support: <pre>com.dstc.security.kerberos.jaas. KerberosLoginModule</pre> <p>Note: This module is the Quest KerberosLoginModule.</p>
<SPN>	<p>The DFC application's SPN.</p> <p>For example, for single-domain support:</p> <pre>DFC/myhost.mydomain.com:8080@SRV01. COM</pre> <p>For multi-domain support, instead of appending the domain name to the SPN, use the <code>realm</code> property to specify the domain name.</p>
<REALM>	<p>(Multi-domain support only) The realm name.</p> <p>For example: SRV01.COM</p>
<dfsuser_keytab_path>	<p>The path to the DFC user account's *.keytab file. For example: c:\dfcuser.keytab</p>

Performance best practices

To make DFC, multi-domain Kerberos perform better, apply the best practices described in this section.

Note: The best practices and/or test results are derived or obtained after testing the product in the EMC testing environment. Every effort is made to simulate common customer usage scenarios during performance testing, but actual performance results will vary due to differences in hardware and software configurations, data, and other variables.

Known Performance Issue with QUEST's TCP Protocol

Performance issues are noticed while using QUEST libraries for Kerberos multiple or single domain authentication. However, the same performance issues are not noticed with SUN libraries for Kerberos single domain authentication. This large time lag or performance issue with QUEST libraries is due to the reason that SUN libraries use UDP as a default protocol to communicate with KDC, whereas QUEST libraries use TCP as default protocol.

QUEST was initially supporting UDP as a default protocol like SUN and noticed packet dropouts in their user environments. Hence, QUEST moved to TCP protocol from UDP protocol for better reliability. EMC recommends to use the QUEST's default supporting protocol that is TCP.

The performance issue with QUEST libraries (using TCP protocol) can only happen in the networks where the packet size of Kerberos request (AS-REQ/TGS-REQ) is less than Ethernet packet size

(1472 bytes). If the packet size of Kerberos request is large enough (a full Ethernet packet or larger) then the performance degrade does not happen.

If you have well behaved reliable network, looking for better performance, and want to avoid performance overhead occurs due to packet size of Kerberos request when TCP protocol is used by QUEST libraries, you may want to use UDP protocol to gain performance benefit. Based on the value provided to the QUEST supported property `jcsi.kerberos.maxpacketsize`, the QUEST libraries choose the protocol TCP or UDP for its Kerberos communication. If the packet size of Kerberos request is more than the value provided to this property, QUEST uses TCP protocol over UDP. However, it is advised to set this property value to 1460 (close to Ethernet packet size) to reduce the packet dropout scenarios. If you still notice the packet dropouts in your environment, use TCP protocol.

Principal authentication support

Java Enterprise Edition (JEE) principal authentication allows you to use the authentication mechanism of a JEE application server to authenticate users of a web application. However, a separate mechanism is required to obtain a session for the user on the repository, once the user is authenticated by the application server. WDK provides built-in support for this, so the information presented in this section is useful primarily for developers of custom web applications not derived from WDK. For more information on using principal authentication in WDK-based applications, see J2EE Principal Authentication in *Web Development Kit Development Guide*.

Implementing principal support in your custom application

To implement principal support, you must provide your own implementation of the `IDfPrincipalSupport` interface and, optionally, of the `IDfTrustManager` interface. The sole method to implement in `IDfPrincipalSupport` is `getSession`, which takes as arguments the repository and principal names, and returns an `IDfSession`. The purpose of this method is to obtain a session for the principal (that is, the user) on the specified repository using a login ticket generated using the login credentials of a trusted authenticator (who is a superuser).

```
public IDfSession getSession  
    (String docbaseName, String principalName) throws DfPrincipalException
```

The intent of the `IDfTrustManager` interface is to obtain login credentials for the trusted authenticator from a secure data source and return them in an `IDfLoginInfo` object to be used internally by the `IDfPrincipalSupport` object. This is done in the `IDfTrustManager.getTrustCredential` method:

```
public IDfLoginInfo getTrustCredential(String docbaseName)
```

To obtain sessions using DFC principal support, use the `IDfClient.setPrincipalSupport` method, passing it an instance of your `IDfPrincipalSupport` implementation class. This changes the behavior of any session manager generated from the `IDfClient`, so that it will delegate the task of obtaining the session to the `IDfPrincipalSupport.getSession` method. You then set the principal name in the session manager using the `setPrincipalName` method. Thereafter, when a session is requested from this session manager instance, it will use the principal support object `getSession` method to obtain a session for the principal on a specified repository.

The following sample demonstrates the basic principal support mechanism.

```
public static void demoPrincipalIdentity(String repository, String principalName)
```

```

    throws Exception
{
    // This assumes that your trust manager implementation
    // gets and stores the superuser credentials when it is constructed.
    IDfTrustManager trustManager = new YourTrustManager();

    // Initialize the client.
    IDfClientX clientX = new DfClientX();
    IDfClient localClient = clientX.getLocalClient();

    // Set principal support on the client. This delegates the task of
    // getting sessions to the IDfPrincipalSupport object.
    IDfPrincipalSupport principalSupport =
        new YourPrincipalSupport(localClient, trustManager);
    localClient.setPrincipalSupport(principalSupport);

    // Create session manager and set principal.
    IDfSessionManager sessMgr = localClient.newSessionManager();
    sessMgr.setPrincipalName(principalName());

    // Session is obtained for the principal, not for the authenticator.
    IDfSession session = sessMgr.getSession(repository);

    try
    {
        System.out.println("Got session: " + session.getSessionId());
        System.out.println("Username: " + session.getLoginInfo().getUser());
    }
    // Release the session in a finally clause.
    finally
    {
        sessMgr.release(session);
    }
}

```

Default classes for demonstrating principal support implementation

DFC includes a default implementation of `IDfPrincipalSupport`, as well as of the supporting interface `IDfTrustManager`, which together demonstrate a design pattern that you can use in building your own principal support implementation. Direct use of these classes is not supported, because they do not provide security for the trusted authenticator password, and because they may change in future releases.

The default implementation of `IDfPrincipalSupport`, `DfDefaultPrincipalSupport`, gets the session by generating a login ticket for the principal using the credentials of a trusted authenticator. The task of obtaining the credentials for the authenticator is managed by another object, of type `IDfTrustManager`. The `IDfTrustManager.getTrustCredential` method returns an `IDfLoginInfo` object containing the credentials of the trusted authenticator, who must be a superuser. `DfDefaultPrincipalSupport`, which is passed an instance of the `IDfTrustManager` implementation class, obtains the trust credentials from the trust manager and uses them to request a login ticket for the principal.

The default implementation of `IDfTrustManager`, `DfDefaultTrustManager`, has overloaded constructors that get the authenticator credentials from either a properties file, or from a Java Properties object passed directly to the constructor. These `DfDefaultTrustManager` constructors do not provide any security for the authenticator password, so it is critical that you do provide your

own implementation of `IDfTrustManager`, and obtain the credentials in a manner that meets your application's security requirements.

Note: Source code for these classes (which should be used as templates only) is available on the developer network: <http://developer.emc.com/developer/samplecode.htm>.

Maintaining state in a session manager

You can cause the session manager to maintain the state of a repository object. The `setSessionManager` method of `IDfTypedObject` accomplishes this. This method copies the state of the object from the session to the session manager, so that disconnecting the session no longer makes references to the object invalid.



Caution: Use `setSessionManager` sparingly. It is an expensive operation. If you find yourself using it often, try to find a more efficient algorithm.

However, using `setSessionManager` is preferable to using the begin/end client control mechanism (refer to the Javadocs for details) to prevent the session manager from disconnecting a session.

Creating a Test Application

The primary use of DFC is to add business logic to your applications. The presentation layer is built using the Web Development Kit, most often via customization of the Webtop interface. Building a custom UI on top of DFC requires a great deal of work, and will largely recreate effort that has already been done for you.

While you should not create a completely new interface for your users, it can be helpful to have a small application that you can use to work with the API directly and see the results. With that in mind, here is a rudimentary interface class that will enable you to add and test behavior using the Operation API.

These examples were created with Oracle JDeveloper, and feature its idiosyncratic ways of building the UI. You can use any IDE you prefer, using this example as a guideline.

To the extent possible, the sample classes have been created as self-contained entities. However, a best practice is to instantiate a Session Manager when you first open an application and continue to use it until the application stops. Therefore, the logic for creating a Session Manager has been incorporated as a method of the `DfcTutorialFrame` class.

This chapter contains the following sections:

- [The `DfcTutorialFrame` class, page 47](#)
- [The `DfcTutorialApplication` class, page 52](#)
- [The `TutorialGetDirectory` class, page 53](#)
- [The `TutorialSessionManager` class, page 54](#)
- [Running the tutorial application, page 55](#)

Note: If you belong to the EMC Developer Network, you can download these files from the following website:

<https://community.emc.com/community/edn/documentum>

You can search for the document *DFC Tutorial Source Files for 6.5 and 6.6* or use the tags `dfc` and `samplecode`.

The `DfcTutorialFrame` class

`DfcTutorialFrame` displays the controls used to set arguments, execute commands, and view results. In this example, it displays buttons used to navigate a directory in the repository.

Example 4-1. DfcBaseTutorialFrame.java

```
package com.emc.tutorial;

import com.documentum.fc.client.IDfCollection;
import com.documentum.fc.client.IDfSession;
import com.documentum.fc.client.IDfSysObject;
import com.documentum.fc.client.IDfTypedObject;

import com.documentum.fc.common.IDfId;

import java.awt.Dimension;

import java.awt.List;
import java.awt.Rectangle;

import java.awt.event.ActionEvent;
import java.awt.event.ActionListener;

import java.util.StringTokenizer;
import java.util.Vector;

import javax.swing.ActionMap;
import javax.swing.JButton;
import javax.swing.JFrame;
import javax.swing.JLabel;
import javax.swing.JList;
import javax.swing.JScrollPane;
import javax.swing.JTextArea;
import javax.swing.JTextField;
import javax.swing.SwingConstants;

public class DfcTutorialFrame extends JFrame {
    private JTextField jTextField_repositoryName = new JTextField();
    private JLabel jLabel_repositoryName = new JLabel();
    private JLabel jLabel_userName = new JLabel();
    private JTextField jTextField_userName = new JTextField();
    private JTextField jTextField_password = new JTextField();
    private JLabel jLabel_password = new JLabel();
    private JTextField jTextField_cwd = new JTextField();
    private JLabel jLabel_cwd = new JLabel();
    private JButton jButton_getDirectory = new JButton();
    private Vector m_fileIDs = new Vector();
    private List list_id = new List();
    private JLabel jLabel_messages = new JLabel();
    private JScrollPane jScrollPane_results = new JScrollPane();
    private JTextArea jTextArea_results = new JTextArea();
    private JLabel jLabel_results = new JLabel();
    private JLabel jLabel_directoryContents = new JLabel();
    private Vector m_breadcrumbs = new Vector();
    private JButton jButton_upDirectory = new JButton();

    public DfcTutorialFrame() {
        try {
            jbInit();
        } catch (Exception e) {
            e.printStackTrace();
        }
    }

    private void jbInit() throws Exception {
        this.getContentPane().setLayout( null );
        this.setSize(new Dimension(679, 479));
        this.setTitle( "DFC Tutorial Frame" );
    }
}
```



```

jTextField_repositoryName.setBounds(new Rectangle(15, 20, 100, 25));
jTextField_repositoryName.setText("MyRepository");
jLabel_repositoryName.setText("Repository Name");
jLabel_repositoryName.setBounds(new Rectangle(15, 5, 100, 15));
jLabel_repositoryName.setHorizontalAlignment(SwingConstants.CENTER);
jLabel_userName.setText("User Name");
jLabel_userName.setBounds(new Rectangle(135, 5, 95, 15));
jLabel_userName.setHorizontalAlignment(SwingConstants.CENTER);
jTextField_userName.setBounds(new Rectangle(135, 20, 95, 25));
jTextField_userName.setText("MyLogin");
jTextField_password.setBounds(new Rectangle(245, 20, 100, 25));
jTextField_password.setText("MyPassword");
jLabel_password.setText("Password");
jLabel_password.setBounds(new Rectangle(245, 5, 100, 15));
jLabel_password.setHorizontalAlignment(SwingConstants.CENTER);
jTextField_cwd.setBounds(new Rectangle(15, 70, 330, 25));
jTextField_cwd.setText("/MyLogin");
jTextField_cwd.addActionListener(new ActionListener() {
    public void actionPerformed(ActionEvent e) {
        jTextField_cwd_actionPerformed(e);
    }
});
jLabel_cwd.setText("Current Working Directory");
jLabel_cwd.setBounds(new Rectangle(15, 50, 130, 20));
jLabel_cwd.setHorizontalAlignment(SwingConstants.CENTER);
jButton_getDirectory.setText("Get Directory");
jButton_getDirectory.setBounds(new Rectangle(15, 230, 100, 20));
jButton_getDirectory.addActionListener(new ActionListener() {
    public void actionPerformed(ActionEvent e) {
        jButton_getDirectory_actionPerformed(e);
    }
});
jLabel_messages.setText("Messages appear here.");
jLabel_messages.setBounds(new Rectangle(5, 430, 495, 15));
jScrollPane_results.setBounds(new Rectangle(15, 275, 330, 140));
jLabel_results.setText("Results:");
jLabel_results.setBounds(new Rectangle(15, 260, 135, 15));
jLabel_directoryContents.setText("Directory contents:");
jLabel_directoryContents.setBounds(new Rectangle(15, 100, 140, 15));
jButton_upDirectory.setText("Up One Level");
jButton_upDirectory.setBounds(new Rectangle(245, 230, 100, 20));
jButton_upDirectory.addActionListener(new ActionListener() {
    public void actionPerformed(ActionEvent e) {
        jButton_upDirectory_actionPerformed(e);
    }
});
list_id.setBounds(new Rectangle(15, 115, 330, 110));
list_id.setSize(new Dimension(330, 110));
list_id.addActionListener(new ActionListener() {
    public void actionPerformed(ActionEvent e) {
        list_id_actionPerformed(e);
    }
});
jScrollPane_results.getViewport().add(jTextArea_results, null);
this.getContentPane().add(jButton_upDirectory, null);
this.getContentPane().add(jLabel_directoryContents, null);
this.getContentPane().add(jLabel_results, null);
this.getContentPane().add(jScrollPane_results, null);
this.getContentPane().add(list_id, null);
this.getContentPane().add(jLabel_messages, null);
this.getContentPane().add(jButton_getDirectory, null);
this.getContentPane().add(jLabel_cwd, null);
this.getContentPane().add(jTextField_cwd, null);
this.getContentPane().add(jLabel_password, null);
this.getContentPane().add(jTextField_password, null);

```

```
this.getContentPane().add(jTextField_userName, null);
this.getContentPane().add(jLabel_userName, null);
this.getContentPane().add(jLabel_repositoryName, null);
this.getContentPane().add(jTextField_repositoryName, null);
}

private void getDirectory() {
    // Get the arguments and assign variable values.

    String repository = jTextField_repositoryName.getText();
    String userName = jTextField_userName.getText();
    String password = jTextField_password.getText();
    String directory = jTextField_cwd.getText();
    TutorialGetDirectory tgd = new TutorialGetDirectory();

    // Empty the file IDs member variable.
    m_fileIDs.clear();

    // Empty the file list display.
    list_id.removeAll();

    IDfCollection folderList = tgd.getDirectory(
        repository,
        userName,
        password,
        directory
    );
    try {
        // Cycle through the collection getting the object ID and adding it to the
        // m_listIDs Vector. Get the object name and add it to the file list control.
        while (folderList.next())
        {
            IDfTypedObject doc = folderList.getTypedObject();
            list_id.add(doc.getString("object_name"));
            m_fileIDs.addElement(doc.getString("r_object_id"));
        }

        // Display a success message for 1 or more objects, or a different message if
        // the command succeeds but there are no files in the directory. Enable the
        // document manipulation buttons if there are items available in the list box.
        if (list_id.getItemCount() > 0) {
            jLabel_messages.setText("Directory query complete.");
        }
        else {
            jLabel_messages.setText("Directory query complete. No items found.");
        }
    }

    // Handle any exceptions.
    catch (Exception ex) {
        jLabel_messages.setText("Exception has been thrown: " + ex);
        ex.printStackTrace();
        list_id.removeAll();
    }
}

private void jButton_getDirectory_actionPerformed(ActionEvent e) {
    String repository = jTextField_repositoryName.getText();
    String userName = jTextField_userName.getText();
    String password = jTextField_password.getText();
    String directory = jTextField_cwd.getText();

    IDfSession mySession = null;
    TutorialSessionManager mySessMgr = null;
    try {
```

```

mySessMgr = new TutorialSessionManager
    (repository, userName, password);
mySession = mySessMgr.getSession();

if(list_id.getSelectedItem()!=null) {

    String objectIdString =
        m_fileIDs.elementAt(list_id.getSelectedIndex()).toString();
    IDfId idObj = mySession.getIdByQualification(
        "dm_sysobject where r_object_id='" + objectIdString + "'");
    IDfSysObject sysObj = (IDfSysObject) mySession.getObject(idObj);
    if(sysObj.getTypeName().equals("dm_folder")) {
        jTextField_cwd.setText(directory + "/" + list_id.getSelectedItem());
        getDirectory();
        StringTokenizer st = new StringTokenizer(jTextField_cwd.getText(), "/");
        m_breadcrumbs.removeAllElements();
        while(st.hasMoreTokens()) m_breadcrumbs.add(st.nextToken());
        StringBuffer newDirectory = new StringBuffer("");
        for (int i=0; i<m_breadcrumbs.elementAt(i));
        }
        jTextField_cwd.setText(newDirectory.toString());
    }
    else {
        getDirectory();
    }
}
else {
    getDirectory();
}
}
catch (Exception ex) {
    jLabel_messages.setText("Exception has been thrown: " + ex);
    ex.printStackTrace();
}
finally {
    mySessMgr.releaseSession(mySession);
}
}

private void changeDirectory() {
    String repository = jTextField_repositoryName.getText();
    String userName = jTextField_userName.getText();
    String password = jTextField_password.getText();
    String directory = jTextField_cwd.getText();

    IDfSession mySession = null;
    TutorialSessionManager mySessMgr = null;
    try {
        mySessMgr = new TutorialSessionManager
            (repository, userName, password);
        mySession = mySessMgr.getSession();
        String objectIdString =
            m_fileIDs.elementAt(list_id.getSelectedIndex()).toString();

        IDfId idObj = mySession.getIdByQualification(
            "dm_sysobject where r_object_id='" + objectIdString + "'");

        // Instantiate an object from the ID.

        IDfSysObject sysObj = (IDfSysObject) mySession.getObject(idObj);
        if(sysObj.getTypeName().equals("dm_folder")) {

            jTextField_cwd.setText(directory + "/" + list_id.getSelectedItem());
            getDirectory();

```

```

        StringTokenizer st = new StringTokenizer(jTextField_cwd.getText(), "/");
        m_breadcrumbs.removeAllElements();
        while(st.hasMoreTokens()) m_breadcrumbs.add(st.nextToken());
        StringBuffer newDirectory = new StringBuffer("");
        for (int i=0; i<m_breadcrumbs.size(); i++) {
            newDirectory.append("/") + m_breadcrumbs.elementAt(i);
        }
        jTextField_cwd.setText(newDirectory.toString());
    }
}
catch (Exception ex) {
    jLabel_messages.setText("Exception has been thrown: " + ex);
    ex.printStackTrace();
}
finally {
    mySessMgr.releaseSession(mySession);
}
}

private void jButton_changeDirectory_actionPerformed(ActionEvent e) {
    changeDirectory();
}

private void jTextField_cwd_actionPerformed(ActionEvent e) {
    getDirectory();
}

private void list_id_actionPerformed(ActionEvent e) {
    changeDirectory();
}

private void jButton_upDirectory_actionPerformed(ActionEvent e) {
    StringBuffer newDirectory = new StringBuffer("");
    newDirectory.append("/") + m_breadcrumbs.elementAt(0);
    if (m_breadcrumbs.size()>1) {
        m_breadcrumbs.removeElementAt(m_breadcrumbs.size()-1);
        for (int i=1; i<m_breadcrumbs.size(); i++) {
            newDirectory.append("/") + m_breadcrumbs.elementAt(i);
        }
        jTextField_cwd.setText(newDirectory.toString());
        getDirectory();
    }
}
}
}

```

The DfcTutorialApplication class

This is the runnable class with the main() method, used to start and stop the application.

Example 4-2. DfcTutorialApplication.java

```

package com.emc.tutorial;

import java.awt.Dimension;
import java.awt.Toolkit;

import javax.swing.JFrame;
import javax.swing.UIManager;

public class DfcTutorialApplication {
    public DfcTutorialApplication() {
        JFrame frame = new DfcTutorialFrame();
    }
}

```

```

        Dimension screenSize = Toolkit.getDefaultToolkit().getScreenSize();
        Dimension frameSize = frame.getSize();
        if (frameSize.height > screenSize.height) {
            frameSize.height = screenSize.height;
        }
        if (frameSize.width > screenSize.width) {
            frameSize.width = screenSize.width;
        }
        frame.setLocation( ( screenSize.width - frameSize.width ) / 2,
            ( screenSize.height - frameSize.height ) / 2 );
        frame.setDefaultCloseOperation( JFrame.EXIT_ON_CLOSE );
        frame.setVisible(true);
    }

    public static void main(String[] args) {
        try {
            UIManager.setLookAndFeel(UIManager.getSystemLookAndFeelClassName());
        } catch (Exception e) {
            e.printStackTrace();
        }
        new DfcTutorialApplication();
    }
}

```

The TutorialGetDirectory class

This class returns the directory contents.

Example 4-3. TutorialGetDirectory.java

```

package com.emc.tutorial;

import com.documentum.fc.client.IDfCollection;
import com.documentum.fc.client.IDfFolder;
import com.documentum.fc.client.IDfSession;

public class TutorialGetDirectory {
    public TutorialGetDirectory() {
    }
    public IDfCollection getDirectory(
        String repository,
        String userName,
        String password,
        String directory
    ) {

        IDfSession mySession = null;
        TutorialSessionManager mySessMgr = null;
        IDfCollection fileList = null;

        try {
            mySessMgr = new TutorialSessionManager(repository, userName, password);
            mySession = mySessMgr.getSession();
            IDfFolder myFolder = mySession.getFolderByPath(directory);
            fileList = myFolder.getContents(null);
        }
        catch (Exception ex) {
            ex.printStackTrace();
        }
        finally {

```

```
        mySessMgr.releaseSession(mySession);
    }
    return fileList;
}
}
```

The TutorialSessionManager class

This class creates the session manager to the repository.

Example 4-4. TutorialSessionManager.java

```
package com.emc.tutorial;

import com.documentum.com.DfClientX;

import com.documentum.fc.client.IDfClient;
import com.documentum.fc.client.IDfSession;
import com.documentum.fc.client.IDfSessionManager;

import com.documentum.fc.common.DfException;
import com.documentum.fc.common.IDfLoginInfo;

public class TutorialSessionManager {

    // Member variables used to encapsulate session and user information.
    private IDfSessionManager m_sessionMgr;
    private String m_repository;
    private String m_userName;
    private String m_password;

    public TutorialSessionManager(String rep, String user, String pword) {
        try {

            // Populate member variables.
            m_repository = rep;
            m_userName = user;
            m_password = pword;

            // Call the local createSessionManager method.
            m_sessionMgr = createSessionManager();
        }
        catch (Exception e) {
            System.out.println("An exception has been thrown: " + e);
        }
    }

    private IDfSessionManager createSessionManager() throws Exception {

        // The only class we instantiate directly is DfClientX.
        DfClientX clientx = new DfClientX();

        // Most objects are created using factory methods in interfaces.
        // Create a client based on the DfClientX object.
        IDfClient client = clientx.getLocalClient();

        // Create a session manager based on the local client.
        IDfSessionManager sMgr = client.newSessionManager();

        // Set the user information in the login information variable.
```

```

        IDfLoginInfo loginInfo = clientx.getLoginInfo();
        loginInfo.setUser(m_userName);
        loginInfo.setPassword(m_password);

// Set the identity of the session manager object based on the repository
// name and login information.
        sMgr.setIdentity(m_repository, loginInfo);

// Return the populated session manager to the calling class. The session
// manager object now has the required information to connect to the
// repository, but is not actively connected.
        return sMgr;
    }

// Request an active connection to the repository.
    public IDfSession getSession() throws DfException {
        return m_sessionMgr.getSession(m_repository);
    }

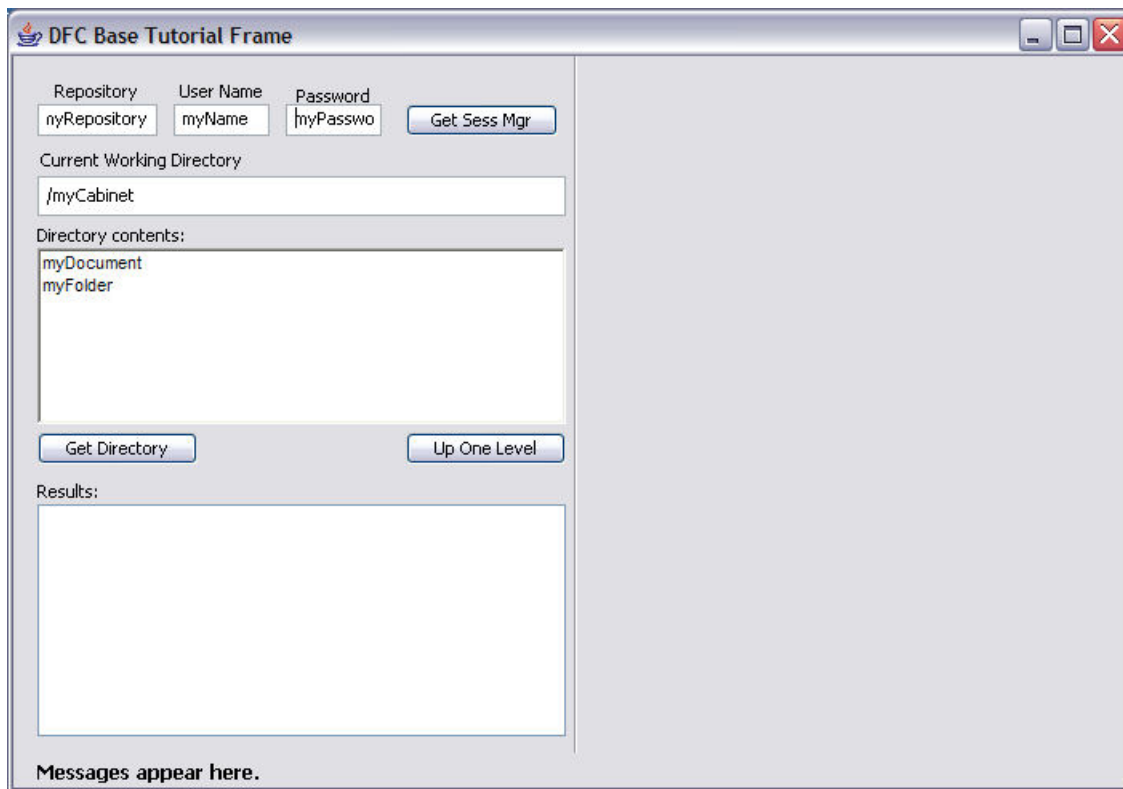
// Release an active connection to the repository for reuse.
    public void releaseSession(IDfSession session) {
        m_sessionMgr.release(session);
    }
}

```

Running the tutorial application

Compile the three base classes and run `DfcBaseTutorialApplication.class`. The application interface appears.

To list the contents of a directory, enter values in the *Repository Name*, *User Name*, *Password*, and *Current Working Directory* fields. Click the Get Directory button to update the file listing. For convenience, you can modify the `DfcBaseTutorialFrame` class to launch the application with these fields already populated with your own settings.

Figure 1. The DFC base tutorial frame

This example demonstrated how you can create an application that connects to the content server and retrieves information. Many of the remaining examples in this manual provide the code for a button handler that you can add to this test application to call classes that demonstrate implementation of individual behaviors using DFC.

Working with Objects

This chapter describes how to use DFC calls to work with objects stored in the Documentum repository.

When you are working with the metadata associated with a repository object, using DFC methods directly can be more efficient and quicker to deploy. When you are working with document content, a better choice is to work with document operations, covered in [Working with Document Operations](#).

This chapter contains the following main sections:

- [Understanding repository objects, page 57](#)
- [Creating a cabinet, page 58](#)
- [Creating a folder, page 60](#)
- [Creating a document object, page 62](#)
- [Accessing attributes, page 64](#)
- [Setting attributes, page 68](#)
- [Removing an attribute value, page 76](#)
- [Getting object content, page 79](#)
- [Destroying an object, page 80](#)
- [Real-time content deletion, page 82](#)

Understanding repository objects

A repository object has two basic components: metadata and content. Metadata describe the object (format, type, etc.) and store other pertinent information such as the object name, owner, creation date, and modification date. As a developer, you are primarily concerned with manipulating document metadata. Metadata are stored in database tables and associated with content objects stored in a local file system.

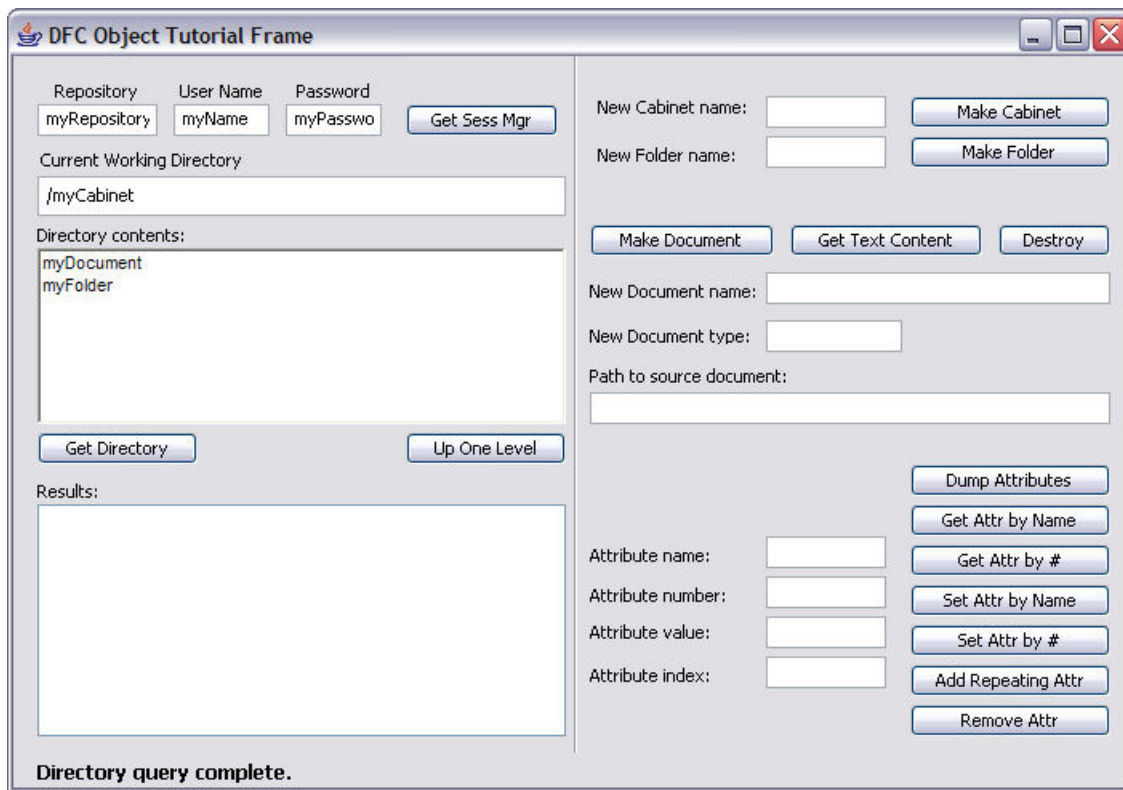
Object content is the information of most interest to end users. If the object is a text document, the content is the text. If the object is a graphic, the content is the binary information used to create the picture.

Using DFC methods, developers can examine and manipulate the metadata associated with repository objects.

The DFC Object Tutorial Frame

The examples in this chapter start with the DFC Base Tutorial Frame and add buttons and fields to demonstrate object manipulation techniques. You can implement any or all of these self-contained examples.

Figure 2. The DFC object tutorial frame



Creating a cabinet

Before you begin manipulating repository objects, you might want to create your own cabinet to work in. A cabinet is a top level container object used to hold folders. The cabinet is, in fact, a type of folder, and is created using the interface `IDfFolder`. Setting its object type to “dm_cabinet” gives it additional features, including the ability to exist as a top-level object.

You can create a test button for creating a cabinet in the DFC Object Tutorial Frame.

To add the Create Cabinet button to the DFC Base Tutorial Frame

1. Create a `TextField` control named `textField_cabinetName`.
2. Create a `Button` control named `button_makeCabinet`.
3. Add a button handler method for [Make Cabinet](#)
4. Create the class [TutorialMakeCabinet](#)

The source code for the button handler and TutorialMakeCabinet class follow.

Example 5-1. Make Cabinet button handler method

```
private void jButton_makeCabinet_actionPerformed(ActionEvent e)
{
    String repositoryName = jTextField_repositoryName.getText();
    String cabinetName = jTextField_cabinetName.getText();

    if (!cabinetName.equals(""))
    {
        TutorialMakeCabinet tmc = new TutorialMakeCabinet();
        if (tmc.makeCabinet(
            m_sessionManager,
            repositoryName,
            cabinetName)
        )
        {
            jTextField_cwd.setText("/") + cabinetName);
            getDirectory();
            jLabel_messages.setText(
                "Created cabinet " + cabinetName + "."
            );
        }
        else
        {
            jLabel_messages.setText("Cabinet creation failed.");
        }
    }
    else
    {
        jLabel_messages.setText("Enter a unique cabinet name " +
            "to create a new cabinet.");
    }
}
```

Example 5-2. The TutorialMakeCabinet class

```
package com.emc.tutorial;

import com.documentum.fc.client.IDfFolder;
import com.documentum.fc.client.IDfSession;
import com.documentum.fc.client.IDfSessionManager;
import com.documentum.fc.client.IDfSysObject;

public class TutorialMakeCabinet
{
    public TutorialMakeCabinet()
    {
    }

    public Boolean makeCabinet(
        IDfSessionManager sessionManager,
        String repositoryName,
        String cabinetName
    )
    {
        IDfSession mySession = null;
        try
        {
            mySession = sessionManager.getSession(repositoryName);
```

```
// check to see if the cabinet already exists
IDfFolder myCabinet = mySession.getFolderByPath("/") + cabinetName);
if (myCabinet == null)
{
    IDfSysObject newCabinet =
        (IDfFolder) mySession.newObject(DM_CABINET);
    newCabinet.setObjectName(cabinetName);
    newCabinet.save();
    return true;
}
else
{
    return false;
}
}
catch (Exception ex)
{
    ex.printStackTrace();
    return false;
}
finally
{
    sessionManager.release(mySession);
}
}
public static final String DM_CABINET = "dm_cabinet";
}
```

Creating a folder

Creating a folder is similar to creating a cabinet. The essential differences are that you will create a *dm_folder* object and identify the parent cabinet or folder in which you'll create it.

To add a Make Folder button to the DFC Base Tutorial Frame

1. Create a JTextField control named jTextField_folderName.
2. Create a JButton control named jButton_makeFolder.
3. Add a handler method for the [Make Folder](#) button.
4. Create the class [TutorialMakeFolder](#)

The code for the [Make Folder button handler](#) and the [TutorialMakeFolder](#) class follow.

Example 5-3. Make Folder button handler method

```
private void jButton_makeFolder_actionPerformed(ActionEvent e)
{
    String repositoryName = jTextField_repositoryName.getText();
    String folderName = jTextField_newFolder.getText();
    String parentName = jTextField_cwd.getText();

    if (!folderName.equals("") & !parentName.equals(""))
    {
        TutorialMakeFolder tmf = new TutorialMakeFolder();
        if (
            tmf.makeFolder(
                m_sessionManager,
```

```

        repositoryName,
        folderName,
        parentName)
    )
    {
        getDirectory();
        jLabel_messages.setText("Created folder " + folderName +
                                ".");
    }

    else
    {
        jLabel_messages.setText("Folder creation failed.");
    }
}
else
{
    jLabel_messages.setText("Enter a folder name and current working" +
                            " directory to create a new folder.");
}
}
}

```

Example 5-4. The TutorialMakeFolder class

```

package com.emc.tutorial;

import com.documentum.fc.client.IDfFolder;
import com.documentum.fc.client.IDfSession;
import com.documentum.fc.client.IDfSessionManager;
import com.documentum.fc.client.IDfSysObject;

public class TutorialMakeFolder
{
    public TutorialMakeFolder()
    {
    }

    public Boolean makeFolder(
        IDfSessionManager sessionManager,
        String repositoryName,
        String folderName,
        String parentName
    )
    {
        IDfSession mySession = null;
        try
        {
            mySession = sessionManager.getSession(repositoryName);
            IDfSysObject newFolder =
                (IDfFolder) mySession.newObject(DM_FOLDER);
            IDfFolder aFolder =
                mySession.getFolderByPath(parentName + "/" + folderName);
            if (aFolder == null)
            {
                newFolder.setObjectName(folderName);
                newFolder.link(parentName);
                newFolder.save();
                return true;
            }
            else
            {
                return false;
            }
        }
    }
}

```

```
    }  
    catch (Exception ex)  
    {  
        ex.printStackTrace();  
        return false;  
    }  
    finally  
    {  
        sessionManager.release(mySession);  
    }  
}  
public static final String DM_FOLDER = "dm_folder";  
}
```

Creating a document object

A document object represents both the content of a document and the metadata that describe the document. In most cases, you create a document by importing an existing document from a local source to the repository.

In the example [TutorialMakeDocument](#), we create a document setting just the minimal required information: the name of the document, the document type, the source of the content, and the parent folder in which to create the new document object.

To add a Make Document button to the Dfc BaseTutorial Frame

1. Create a JTextField control named jTextField_documentName.
2. Create a JTextField control named jTextField_srcDocPath.
3. Create a JTextField control named jTextField_docType.
4. Create a JButton control named jButton_makeDocument.
5. Add a [button handler method for Make Document](#).
6. Create the class [TutorialMakeDocument](#)

The following is an example of the Make Document button handler.

Example 5-5. The Make Document button handler

```
private void jButton_makeDocument_actionPerformed(ActionEvent e)  
{  
    String repository = jTextField_repositoryName.getText();  
    String documentName = jTextField_documentName.getText();  
    String folderName = jTextField_cwd.getText();  
    String srcPath = jTextField_sourceDocumentPath.getText();  
    String docType = jTextField_documentType.getText();  
  
    TutorialMakeDocument tmd = new TutorialMakeDocument();  
    if (tmd.makeDocument(  
        m_sessionManager,  
        repository,  
        documentName,  
        docType,  
        srcPath,  
        folderName  
    )  
}
```

```

    )
    {
        jLabel_messages.setText("Created document " + documentName + ".");
        getDirectory();
    }
    else
    {
        jLabel_messages.setText("Document creation failed.");
    }
}

```

The following example creates a document object in the repository.

Example 5-6. The TutorialMakeDocument class

```

package com.emc.tutorial;

import com.documentum.fc.client.IDfSession;
import com.documentum.fc.client.IDfSessionManager;
import com.documentum.fc.client.IDfSysObject;

public class TutorialMakeDocument
{
    public TutorialMakeDocument()
    {
    }

    public Boolean makeDocument(
        IDfSessionManager sessionManager,
        String repositoryName,
        String documentName,
        String documentType,
        String sourcePath,
        String parentName)
    {
        IDfSession mySession = null;
        try
        {
            mySession = sessionManager.getSession(repositoryName);
            IDfSysObject newDoc =
                (IDfSysObject) mySession.newObject(DM_DOCUMENT);
            newDoc.setObjectName(documentName);
            newDoc.setContentType(documentType);
            newDoc.setFile(sourcePath);
            newDoc.link(parentName);
            newDoc.save();
            return true;
        }
        catch (Exception ex)
        {
            ex.printStackTrace();
            return false;
        }
        finally
        {
            sessionManager.release(mySession);
        }
    }

    public static final String DM_DOCUMENT = "dm_document";
}

```

Accessing attributes

Having created a document or folder, you can examine its attributes by performing a “dump” of all attributes or by specifying a particular attribute you want to examine.

Dumping Attributes

The most convenient way to begin working with attributes is to “dump” all of the object attributes to get a list of their names and values.

To add a Dump Attributes button to the Dfc Base Tutorial Frame

1. Create a JButton control named `jButton_dumpAttributes`.
2. Add a [button handler method for Dump Attributes](#).
3. Create and deploy the class [TutorialDumpAttributes](#)

Example 5-7. The Dump Attributes button handler

```
private void jButton_dumpAttributes_actionPerformed(ActionEvent e)
{
    String repositoryName = jTextField_repositoryName.getText();

    if (list_id.getSelectedIndex() == -1)
    {
        jLabel_messages.setText(
            "Select an item in the file list to dump attributes."
        );
    }
    else
    {
        String objectIdString =
            m_fileIDs.elementAt(list_id.getSelectedIndex()).toString();
        TutorialDumpAttributes tda = new TutorialDumpAttributes();
        jTextArea_results.setText(
            tda.dumpAttributes(
                m_sessionManager,
                repositoryName,
                objectIdString
            );
        jLabel_messages.setText("Query complete.");
    }
}
```

Example 5-8. TutorialDumpAttributes class

```
package com.emc.tutorial;

import com.documentum.fc.client.IDfContainment;
import com.documentum.fc.client.IDfSession;
import com.documentum.fc.client.IDfSessionManager;
import com.documentum.fc.client.IDfSysObject;
import com.documentum.fc.common.DfId;
import com.documentum.fc.common.IDfId;

public class TutorialDumpAttributes
```



```

{
    public TutorialDumpAttributes()
    {
    }

    public String dumpAttributes
    (
        IDfSessionManager sessionManager,
        String repositoryName,
        String objectIdString
    )
    {
        IDfSession mySession = null;
        StringBuffer attributes = new StringBuffer("");
        try
        {
            mySession = sessionManager.getSession(repositoryName);

            IDfId idObj = mySession.getIdByQualification(
                "dm_sysobject where r_object_id='" + objectIdString + "'"
            );
            IDfSysObject sysObj = (IDfSysObject) mySession.getObject(idObj);
            attributes.append(sysObj.dump());
        }
        catch (Exception ex)
        {
            ex.printStackTrace();
            attributes.append("\nException: " + ex.toString());
        }
        finally
        {
            sessionManager.release(mySession);
        }
        return attributes.toString();
    }
}

```

Getting a single attribute by name

If you know the name of the attribute you want to access (or you just looked it up by dumping all of the attributes), you can request the specific attribute directly.

To add a Get Attribute By Name button to the DFC Base Tutorial Frame

1. Create a JTextField control named jTextField_attributeName
2. Create a JButton control named jButton_getAttributeByName
3. Add a handler method for the [Get Attribute By Name](#) button.
4. Create and deploy the class [TutorialGetAttributeByName](#)

Example 5-9. Handler for the Get Attribute By Name button

```

private void jButton_getAttributeByName_actionPerformed(ActionEvent e)
{
    String repositoryName = jTextField_repositoryName.getText();

    if (list_id.getSelectedIndex() == -1 |

```

```
        jTextField_attributeName.getText().equals(""))
    {
        jLabel_messages.setText(
            "Select an item and enter an attribute name to get" +
            " attribute information."
        );
    }
    else
    {
        String objectIdString =
            m_fileIDs.elementAt(list_id.getSelectedIndex()).toString();
        String attributeName = jTextField_attributeName.getText();
        TutorialGetAttributeByName tgabn =
            new TutorialGetAttributeByName();
        jTextArea_results.setText(
            tgabn.getAttribute(
                m_sessionManager,
                repositoryName,
                objectIdString,
                attributeName
            )
        );
        jLabel_messages.setText("Query complete.");
    }
}
```

Example 5-10. The TutorialGetAttributeByName class

```
package com.emc.tutorial;

import com.documentum.fc.client.IDfSession;
import com.documentum.fc.client.IDfSessionManager;
import com.documentum.fc.client.IDfSysObject;
import com.documentum.fc.common.IDfId;

public class TutorialGetAttributeByName
{
    public TutorialGetAttributeByName()
    {
    }

    public String getAttribute(
        IDfSessionManager sessionManager,
        String repositoryName,
        String objectIdString,
        String attributeName
    )
    {
        IDfSession mySession = null;
        StringBuffer attribute = new StringBuffer("");
        try
        {
            mySession = sessionManager.getSession(repositoryName);
            IDfId idObj =
                mySession.getIdByQualification(
                    "dm_sysobject where r_object_id='" + objectIdString + "'"
                );
            IDfSysObject sysObj = (IDfSysObject) mySession.getObject(idObj);
            attribute.append(attributeName);
            attribute.append(": ");
            attribute.append(sysObj.getValue(attributeName).toString());
        }
        catch (Exception ex)
        {
        }
    }
}
```

```

        {
            ex.printStackTrace();
            attribute = new StringBuffer("");
            attribute.append(ex.toString());
        }
// Always, always, always release the session when finished using it.
finally
{
    sessionManager.release(mySession);
}
return attribute.toString();
}
}

```

Getting a single attribute by number

There may be times where you capture the attribute's index number rather than its name. The `IDfSystemObject.getAttr()` method requires the attribute name as the argument. You can use the attribute's index number to get the attribute name, then continue on to get the attribute.

To add a Get Attribute By Number button to the DFC Base Tutorial Frame

1. Create a `TextField` control named `jTextField_attributeNumber`
2. Create a `Button` control named `jButton_getAttributeByNumber`
3. Add a handler method for the [Get Attribute By Number](#) button.
4. Create and deploy the class [TutorialGetAttributeByNumber](#)

Example 5-11. Handler for the Get Attribute By Number button

```

private void jButton_getAttributeByNumber_actionPerformed(ActionEvent e)
{
    String repositoryName = jTextField_repositoryName.getText();

    if (list_id.getSelectedIndex() == -1 |
        jTextField_attributeNum.getText().equals(""))
    {
        jLabel_messages.setText(
            "Select an item and enter an attribute number to get" +
            "attribute information."
        );
    }
    else
    {
        String objectIdString =
            m_fileIDs.elementAt(list_id.getSelectedIndex()).toString();
        int theIndex = Integer.parseInt(jTextField_attributeNum.getText());

        TutorialGetAttributeByNumber tga =
            new TutorialGetAttributeByNumber();

        jTextArea_results.setText(
            tga.getAttribute(
                m_sessionManager,
                repositoryName,
                objectIdString,
                theIndex
            )
        );
    }
}

```

```
        )
    );
    jLabel_messages.setText("Query complete.");
}
}
```

Example 5-12. The TutorialGetAttributeByNumber class

```
package com.emc.tutorial;

import com.documentum.fc.client.IDfSession;
import com.documentum.fc.client.IDfSessionManager;
import com.documentum.fc.client.IDfSysObject;
import com.documentum.fc.common.IDfId;

public class TutorialGetAttributeByNumber {
    public TutorialGetAttributeByNumber() {
    }
    public String getAttribute (
        IDfSessionManager sessionManager,
        String repositoryName,
        String objectIdString,
        int theIndex
    )
    {
        IDfSession mySession = null;
        StringBuffer attribute = new StringBuffer("");
        try {
            mySession = sessionManager.getSession(repositoryName);
            IDfId idObj = mySession.getIdByQualification(
                "dm_sysobject where r_object_id='" + objectIdString + "'");
            IDfSysObject sysObj = (IDfSysObject) mySession.getObject(idObj);
            attribute.append(sysObj.getAttr(theIndex).getName());
            attribute.append(": ");
            attribute.append(sysObj.getValueAt(theIndex).toString());
        }
        catch (Exception ex) {
            ex.printStackTrace();
            attribute = new StringBuffer("");
            attribute.append(ex.toString());
        }
        finally {
            sessionManager.release(mySession);
        }
        return attribute.toString();
    }
}
```

Setting attributes

Attributes are set according to their datatype. For example, you set a String value by calling the following method.

```
IDfTypedObject.setString(String attributeName, String value)
```

The available datatypes are DFC versions of Boolean, integer, String, time, DfID, and undefined. In practice, you will not work with the DfID because all object IDs are set internally and are not mutable by custom applications.

Time values (including date values) are stored as numeric values, and are entered or displayed using a pattern mask. For example, in these examples, we use the pattern `IDfTime.DF_TIME_PATTERN_2`, `mm/dd/yyyy`. The complete list of time formats can be found in the Javadoc entry for `IDfTime`.

Setting a single attribute

You can set attributes directly by type. Most often, you will have a specific control that will set a specific data type. Alternatively, this example queries for the data type of the attribute name the user supplies, then uses a switch statement to set the value accordingly.

To add a Set Attribute By Name button to the DFC Base Tutorial Frame

1. If you have not done so already, create a `TextField` control named `textField_attributeName`.
2. Create a `TextField` control named `textField_attributeValue`
3. Create a `Button` control named `button_setAttributeByName`.
4. Add a [button handler method for Set Attribute By Name](#).
5. Create and deploy the class [TutorialSetAttributeByName](#)

Example 5-13. Handler for the Set Attribute By Name button

```
private void jButton_setAttributeByName_actionPerformed(ActionEvent e)
{
    String repositoryName = textField_repositoryName.getText();
    String attributeName = textField_attributeName.getText();
    String attributeValue = textField_attributeValue.getText();

    if (list_id.getSelectedIndex() == -1 |
        textField_attributeName.getText().equals(""))
    {
        jLabel_messages.setText(
            "Select an item, enter an attribute name and value to set" +
            " attribute information.");
    }
    else
    {
        String objectIdString =
            m_fileIDs.elementAt(list_id.getSelectedIndex()).toString();
        TutorialSetAttributeByName tsabn =
            new TutorialSetAttributeByName();
        JTextArea_results.setText(
            tsabn.setAttributeByName(
                m_sessionManager,
                repositoryName,
                objectIdString,
                attributeName,
                attributeValue)
            );
        jLabel_messages.setText("Set attribute complete.");
    }
}
```

Example 5-14. The TutorialSetAttributeByName class

```
package com.emc.tutorial;
```

```
import com.documentum.fc.client.IDfSession;
import com.documentum.fc.client.IDfSessionManager;
import com.documentum.fc.client.IDfSysObject;
import com.documentum.fc.client.IDfType;
import com.documentum.fc.common.DfId;
import com.documentum.fc.common.DfTime;
import com.documentum.fc.common.IDfId;
import com.documentum.fc.common.IDfTime;

public class TutorialSetAttributeByName
{
    public TutorialSetAttributeByName()
    {
    }

    public String setAttributeByName(
        IDfSessionManager sessionManager,
        String repositoryName,
        String objectIdString,
        String attributeName,
        String attributeValue)
    {
        IDfSession mySession = null;
        try
        {
            mySession = sessionManager.getSession(repositoryName);
            IDfId idObj =
                mySession.getIdByQualification(
                    "dm_sysobject where r_object_id='" + objectIdString + "'"
                );

            IDfSysObject sysObj = (IDfSysObject) mySession.getObject(idObj);

            int attributeDatatype = sysObj.getAttrDataType(attributeName);
            StringBuffer results = new StringBuffer("");
            results.append("Previous value: " +
                sysObj.getValue(attributeName).toString());
            switch (attributeDatatype)
            {
                case IDfType.DF_BOOLEAN:
                    if (attributeValue.equals("F") |
                        attributeValue.equals("f") |
                        attributeValue.equals("0") |
                        attributeValue.equals("false") |
                        attributeValue.equals("FALSE"))
                        sysObj.setBoolean(attributeName, false);
                    if (attributeValue.equals("T") |
                        attributeValue.equals("t") |
                        attributeValue.equals("1") |
                        attributeValue.equals("true") |
                        attributeValue.equals("TRUE"))
                        sysObj.setBoolean(attributeName, true);
                    break;

                case IDfType.DF_INTEGER:
                    sysObj.setInt(attributeName,
                        Integer.parseInt(attributeValue));
                    break;

                case IDfType.DF_STRING:
                    sysObj.setString(attributeName, attributeValue);
                    break;

                // This case should not arise - no user-settable IDs
            }
        }
        catch (Exception e)
        {
            return e.getMessage();
        }
        return results.toString();
    }
}
```

```

        case IDfType.DF_ID:
            IDfId newId = new DfId(attributeValue);
            sysObj.setId(attributeName, newId);
            break;

        case IDfType.DF_TIME:
            DfTime newTime =
                new DfTime(attributeValue, IDfTime.DF_TIME_PATTERN2);
            sysObj.setTime(attributeName, newTime);
            break;

        case IDfType.DF_UNDEFINED:
            sysObj.setString(attributeName, attributeValue);
            break;

    }
    if (sysObj.fetch(null))
    {
        results = new StringBuffer("Object is no longer current.");
    }
    else
    {
        sysObj.save();
        results.append("\nNew value: " + attributeValue);
    }
    return results.toString();
}
catch (Exception ex)
{
    ex.printStackTrace();
    return "Set attribute command failed.";
}
finally
{
    sessionManager.release(mySession);
}
}
}

```

Setting an attribute by number

The `IDfSysObject.setAttr()` methods require an attribute name. If your application captures the attribute's index number, you can use it to get the name of the attribute, then continue on to set the value.

To add a Set Attribute By Number button to the DFC Base Tutorial Frame

1. If you have not done so already, create a `TextField` control named `textField_attributeNumber`.
2. If you have not done so already, create a `TextField` control named `textField_attributeValue`.
3. Create a `Button` control named `button_setAttributeByNumber`.
4. Add a [button handler method for Set Attribute By Number](#).
5. Create and deploy the class [TutorialSetAttributeByNumber](#).

Example 5-15. Handler for the Set Attribute By Number button

```
private void jButton_setAttributeByNumber_actionPerformed(ActionEvent e)
{
    String repositoryName = jTextField_repositoryName.getText();
    String attributeNumber = jTextField_attributeNum.getText();
    String attributeValue = jTextField_attributeValue.getText();

    if (list_id.getSelectedIndex() == -1 |
        jTextField_attributeNum.getText().equals(""))
    {
        jLabel_messages.setText(
            "Select an item, enter an attribute name and value to set" +
            " attribute information."
        );
    }
    else
    {
        String objectIdString =
            m_fileIDs.elementAt(list_id.getSelectedIndex()).toString();
        TutorialSetAttributeByNumber tsabn =
            new TutorialSetAttributeByNumber();
        jTextArea_results.setText(
            tsabn.setAttributeByNumber(
                m_sessionManager,
                repositoryName,
                objectIdString,
                attributeNumber,
                attributeValue
            )
        );
        jLabel_messages.setText("Set attribute complete.");
    }
}
```

Example 5-16. The TutorialSetAttributeByNumber class

```
package com.emc.tutorial;

import com.documentum.fc.client.IDfSession;
import com.documentum.fc.client.IDfSessionManager;
import com.documentum.fc.client.IDfSysObject;
import com.documentum.fc.client.IDfType;
import com.documentum.fc.common.DfId;
import com.documentum.fc.common.DfTime;
import com.documentum.fc.common.IDfId;
import com.documentum.fc.common.IDfTime;

public class TutorialSetAttributeByNumber
{
    public TutorialSetAttributeByNumber()
    {
    }

    public String setAttributeByNumber
    (
        IDfSessionManager sessionManager,
        String repositoryName,
        String objectIdString,
        String attributeNumber,
        String attributeValue
    )
    {
        IDfSession mySession = null;
```



```

try
{
    mySession = sessionManager.getSession(repositoryName);
    IDfId idObj =
        mySession.getIdByQualification(
            "dm_sysobject where r_object_id='" +
                objectIdString + "'");

    IDfSysObject sysObj = (IDfSysObject) mySession.getObject(idObj);
    int attributeInt = Integer.parseInt(attributeNumber);
    String attributeName = sysObj.getAttr(attributeInt).getName();
    int attributeDatatype = sysObj.getAttrDataType(attributeName);

    StringBuffer results = new StringBuffer("");
    results.append("Previous value: " +
        sysObj.getValue(attributeName).toString());
    switch (attributeDatatype)
    {
        case IDfType.DF_BOOLEAN:
            if (attributeValue.equals("F") |
                attributeValue.equals("f") |
                attributeValue.equals("0") |
                attributeValue.equals("false") |
                attributeValue.equals("FALSE"))
                sysObj.setBoolean(attributeName, false);
            if (attributeValue.equals("T") |
                attributeValue.equals("t") |
                attributeValue.equals("1") |
                attributeValue.equals("true") |
                attributeValue.equals("TRUE"))
                sysObj.setBoolean(attributeName, true);
            break;

        case IDfType.DF_INTEGER:
            sysObj.setInt(attributeName,
                Integer.parseInt(attributeValue));
            break;

        case IDfType.DF_STRING:
            sysObj.setString(attributeName, attributeValue);
            break;

        // This case should not arise - no user-settable IDs
        case IDfType.DF_ID:
            IDfId newId = new DfId(attributeValue);
            sysObj.setId(attributeName, newId);
            break;

        case IDfType.DF_TIME:
            DfTime newTime =
                new DfTime(attributeValue, IDfTime.DF_TIME_PATTERN2);
            sysObj.setTime(attributeName, newTime);
            break;

        case IDfType.DF_UNDEFINED:
            sysObj.setString(attributeName, attributeValue);
            break;
    }
    if (sysObj.fetch(null))
    {
        results = new StringBuffer("Object is no longer current.");
    }
    else

```

```
        {
            sysObj.save();
            results.append("\nNew value: " + attributeValue);
        }
        return results.toString();
    }
}
catch (Exception ex)
{
    ex.printStackTrace();
    return "Set attribute command failed.";
}
finally
{
    sessionManager.release(mySession);
}
}
```

Appending a repeating attribute

If you use the regular `set[Datatype]` method for setting a repeating attribute, the first value (at the zero index) will be set with the value you provide. You can use an `append[Datatype]` method to add a value to a repeating attribute.

To add an Append Repeating Attribute button to the DFC Base Tutorial Frame

1. If you have not already done so, create a `TextField` control named `jTextField_attributeName`.
2. If you have not already done so, create a `TextField` control named `jTextField_attributeValue`.
3. Create a `Button` control named `jButton_appendRepeatingAttribute`
4. Add a [button handler method for Append Repeating Attribute](#).
5. Create and deploy the class [TutorialAppendRepeatingAttribute](#)

Example 5-17. Handler for the Append Repeating Attribute button

```
private void jButton_appendRepeatingAttribute_actionPerformed(ActionEvent e)
{
    String repositoryName = jTextField_repositoryName.getText();
    String attributeName = jTextField_attributeName.getText();
    String attributeValue = jTextField_attributeValue.getText();

    if (list_id.getSelectedIndex() == -1 |
        jTextField_attributeName.getText().equals(""))
    {
        jLabel_messages.setText(
            "Select an item, enter an attribute name, index, and value " +
            "to set repeating attribute information."
        );
    }
    else
    {
        String objectIdString =
            m_fileIDs.elementAt(list_id.getSelectedIndex()).toString();
        TutorialAppendRepeatingAttribute tara =
            new TutorialAppendRepeatingAttribute();
        tara.appendRepeatingAttribute(
```

```

        m_sessionManager,
        repositoryName,
        objectIdString,
        attributeName,
        attributeValue
    );

    jLabel_messages.setText("Set attribute complete.");
}
}

```

Example 5-18. The TutorialAppendRepeatingAttribute class

```

package com.emc.tutorial;

import com.documentum.fc.client.IDfSession;
import com.documentum.fc.client.IDfSessionManager;
import com.documentum.fc.client.IDfSysObject;
import com.documentum.fc.client.IDfType;
import com.documentum.fc.common.DfId;
import com.documentum.fc.common.DfTime;
import com.documentum.fc.common.IDfId;
import com.documentum.fc.common.IDfTime;

public class TutorialAppendRepeatingAttribute
{
    public TutorialAppendRepeatingAttribute()
    {
    }

    public String appendRepeatingAttribute
    (
        IDfSessionManager sessionManager,
        String repositoryName,
        String objectIdString,
        String attributeName,
        String attributeValue)
    {
        IDfSession mySession = null;
        try
        {
            mySession = sessionManager.getSession(repositoryName);
            IDfId idObj =
                mySession.getIdByQualification(
                    "dm_sysobject where r_object_id='" + objectIdString + "'"
                );

            IDfSysObject sysObj = (IDfSysObject) mySession.getObject(idObj);

            int attributeDatatype = sysObj.getAttrDataType(attributeName);
            StringBuffer results = new StringBuffer("");
            results.append("Previous value: " +
                sysObj.getValue(attributeName).toString());

            switch (attributeDatatype)
            {
                case IDfType.DF_BOOLEAN:
                    if (attributeValue.equals("F") |
                        attributeValue.equals("f") |
                        attributeValue.equals("0") |
                        attributeValue.equals("false") |
                        attributeValue.equals("FALSE"))
                        sysObj.setBoolean(attributeName, false);

```

```
        if (attributeValue.equals("T") |
            attributeValue.equals("t") |
            attributeValue.equals("1") |
            attributeValue.equals("true") |
            attributeValue.equals("TRUE"))
            sysObj.appendBoolean(attributeName, true);
        results.append("\nNew value: " + attributeValue);
        break;

    case IDfType.DF_INTEGER:
        sysObj.appendInt(attributeName,
                        Integer.parseInt(attributeValue));
        break;

    case IDfType.DF_STRING:
        sysObj.appendString(attributeName, attributeValue);
        break;

    // This case should not arise - no user-settable IDs
    case IDfType.DF_ID:
        IDfId newId = new DfId(attributeValue);
        sysObj.appendId(attributeName, newId);
        break;

    case IDfType.DF_TIME:
        DfTime newTime =
            new DfTime(attributeValue, IDfTime.DF_TIME_PATTERN2);
        if (newTime.isValid())
            sysObj.appendTime(attributeName, newTime);
        else
            results = new StringBuffer("");
        break;

    case IDfType.DF_UNDEFINED:
        sysObj.appendString(attributeName, attributeValue);
        break;
    }
    sysObj.save();
    return results.toString();
}
catch (Exception ex)
{
    ex.printStackTrace();
    return "Set attribute command failed.";
}
finally
{
    sessionManager.release(mySession);
}
}
```

Removing an attribute value

You can also remove attribute values from system objects. If the attribute value is a single value, you can set the value of the variable to *null*. For a repeating attribute, you can either delete an individual value by providing the index of the value, or delete all values for the attribute. The example below uses simple tests to decide which of these options to choose: your own application will most likely

need to use more robust decision-making logic, but the API calls will be the same once you've determined which attributes are to be removed.

To add a Remove Attribute button to the DFC Base Tutorial Frame

1. If you haven't already done so, create a JTextField control named jTextField_attributeName.
2. Create a JTextField control named jTextField_attributeIndex.
3. Create a JButton control named jButton_removeAttribute
4. Add a [button handler method for Remove Attribute](#).
5. Create and deploy the class [TutorialRemoveAttribute](#)

Example 5-19. Handler for the Remove Attribute button

```
private void jButton_removeAttribute_actionPerformed(ActionEvent e)
{
    String repositoryName = jTextField_repositoryName.getText();
    String attributeName = jTextField_attributeName.getText();
    String objectIdString =
        m_fileIDs.elementAt(list_id.getSelectedIndex()).toString();
    String attributeIndex = jTextField_attributeIndex.getText();
    TutorialRemoveAttribute tra = new TutorialRemoveAttribute();
    if (
        tra.removeAttribute(
            m_sessionManager,
            repositoryName,
            objectIdString,
            attributeName,
            attributeIndex
        )
    )
    {
        jLabel_messages.setText("Attribute removed.");
    }
    else
    {
        jLabel_messages.setText("Attribute removal operation failed.");
    }
}
```

Example 5-20. The TutorialRemoveAttribute class

```
package com.emc.tutorial;

import com.documentum.fc.client.IDfSession;
import com.documentum.fc.client.IDfSessionManager;
import com.documentum.fc.client.IDfSysObject;
import com.documentum.fc.common.IDfId;

public class TutorialRemoveAttribute
{
    public TutorialRemoveAttribute()
    {
    }

    public Boolean removeAttribute
    (
        IDfSessionManager sessionManager,
        String repositoryName,
        String objectIdString,
```

```
String attributeName,
String attributeIndex
)
{
    IDfSession mySession = null;

    TutorialSessionManager mySessMgr = null;
    try
    {
        mySession = sessionManager.getSession(repositoryName);
        IDfId idObj =
            mySession.getIdByQualification(
                "dm_sysobject where r_object_id='" + objectIdString + "'"
            );
        IDfSysObject sysObj = (IDfSysObject) mySession.getObject(idObj);
        int indexInt = -1;
        try
        {
            indexInt = Integer.parseInt(attributeIndex);
        }
        catch (Exception ex)
        {
            ex.printStackTrace();
        }
        if (sysObj.isAttrRepeating(attributeName))
        {
            if (indexInt > -1)
            {
                sysObj.remove(attributeName, indexInt);
                sysObj.save();
            }
            else
            {
                sysObj.removeAll(attributeName);
                sysObj.save();
            }
        }
        else
        {
            TutorialSetAttributeByName tsabn =
                new TutorialSetAttributeByName();
            tsabn.setAttributeByName(
                sessionManager,
                repositoryName,
                objectIdString,
                attributeName,
                null
            );
        }
        return true;
    }
    catch (Exception ex)
    {
        ex.printStackTrace();
        return false;
    }
    finally
    {
        sessionManager.release(mySession);
    }
}
```

Getting object content

The `IDfSysObject.getContent()` command lets you get the contents of a document as a `ByteArrayInputStream`. If you are working with the content of a file, particularly if you intend to save changes, you should use the operation interfaces, which wrap the behavior required to safely check out, update, and check in document content. There may be times, though, that you want to read the content of a file without manipulating it. In those cases, you can use the `getContent()` method.

To add a Get Content button to the DFC Base Tutorial Frame

1. Create a `JButton` control named `jButton_getDocumentContent`
2. Add a [button handler method for Get Document Content](#)..
3. Create and deploy the class [TutorialGetTextContent](#)

Example 5-21. Handler for the Get Document Content button

```
private void jButton_getDocumentContent_actionPerformed(ActionEvent e)
{
    String repositoryName = jTextField_repositoryName.getText();
    String objectIdString =
        m_fileIDs.elementAt(list_id.getSelectedIndex()).toString();

    TutorialGetTextContent tgtc = new TutorialGetTextContent();

    jTextArea_results.setText(
        tgtc.getContent(
            m_sessionManager,
            repositoryName,
            objectIdString
        )
    );
}
```

Example 5-22. The TutorialGetTextContent class

```
package com.emc.tutorial;

import com.documentum.fc.client.IDfSession;
import com.documentum.fc.client.IDfSessionManager;
import com.documentum.fc.client.IDfSysObject;
import com.documentum.fc.common.IDfId;

import java.io.BufferedReader;
import java.io.ByteArrayInputStream;
import java.io.InputStreamReader;

public class TutorialGetTextContent
{
    public TutorialGetTextContent()
    {
    }

    public String getContent(
        IDfSessionManager sessionManager,
        String repositoryName,
        String objectIdString
    )
    {
    }
}
```

```
    IDfSession mySession = null;
    StringBuffer sb = new StringBuffer("");

    try
    {
        mySession = sessionManager.getSession(repositoryName);

        // Get the object ID based on the object ID string.
        IDfId idObj =
            mySession.getIdByQualification(
                "dm_sysobject where r_object_id='" + objectIdString + "'"
            );

        // Instantiate an object from the ID.

        IDfSysObject sysObj = (IDfSysObject) mySession.getObject(idObj);
        if (sysObj.getContentType().equals("crtxt") |
            sysObj.getContentType().equals("text"))
        {
            ByteArrayInputStream buf = sysObj.getContent();
            int i = 0;
            InputStreamReader readInput =
                new InputStreamReader(buf, "UTF-8");
            BufferedReader br = new BufferedReader(readInput);
            while (br.ready())
            {
                sb.append(br.readLine());
                sb.append("\n");
            }
            return sb.toString();
        }
        else
        {
            return "Use getContent to view text documents.";
        }
    }

    // Handle any exceptions.
    catch (Exception ex)
    {
        ex.printStackTrace();
        return "Exception has been thrown: " + ex;
    }

    // Always, always, release the session in the "finally" clause.
    finally
    {
        sessionManager.release(mySession);
    }
}
```

Destroying an object

You can use the `IDfSystemObject.destroyAllVersions()` method to permanently remove an object from the database. If you use the `IDfPersistentObject.destroy()` method, you will destroy only the specific system object corresponding to the `r_object_id` you provide. In this example, we use the `destroyAllVersions()` method, which destroys not only the system object with the corresponding ID but all iterations of the object.

If you attempt to destroy a directory that has children, the method will return an error.

To add a Destroy button to the DFC Base Tutorial Frame

1. Create a JButton control named jButton_destroy
2. Add a handler method for the [Destroy](#) button.
3. Create and deploy the class [TutorialDestroyObject](#).

Example 5-23. Handler for the Destroy button

```
private void jButton_destroy_actionPerformed(ActionEvent e)
{
    String repositoryName = jTextField_repositoryName.getText();
    String objectIdString =
        m_fileIDs.elementAt(list_id.getSelectedIndex()).toString();

    TutorialDestroyObject tdo = new TutorialDestroyObject();
    if (tdo.destroyObject(
        m_sessionManager,
        repositoryName,
        objectIdString
    )
    )
    {
        getDirectory();
        jLabel_messages.setText("Destroyed object '" + objectIdString +
                                "'.");
    }
    else
    {
        jLabel_messages.setText("Destroy command failed.");
    }
}
```

Example 5-24. The TutorialDestroyObject class

```
package com.emc.tutorial;

import com.documentum.fc.client.IDfSession;
import com.documentum.fc.client.IDfSessionManager;
import com.documentum.fc.client.IDfSysObject;
import com.documentum.fc.common.IDfId;

public class TutorialDestroyObject
{
    public TutorialDestroyObject()
    {
    }

    public Boolean destroyObject(
        IDfSessionManager sessionManager,
        String repositoryName,
        String objectIdString
    )
    {
        IDfSession mySession = null;
        try
        {
            mySession = sessionManager.getSession(repositoryName);
            IDfId idObj =
                mySession.getIdByQualification(
```

```
        "dm_sysobject where r_object_id='" + objectIdString + "'"
    );
    IDfSysObject sysObj = (IDfSysObject) mySession.getObject(idObj);
    sysObj.destroyAllVersions();
    return true;
}
catch (Exception ex)
{
    ex.printStackTrace();
    return false;
}
finally
{
    sessionManager.release(mySession);
}
}
```

Real-time content deletion

For a conventional object deletion, when you delete an object, the contents of the object are deleted asynchronously. A real-time content deletion removes the contents immediately with the object.

Starting with DFC 7.0, you can use the `setRealtimeContentDel (bool enable)` method of `DfSession` to control whether or not to enable the real-time content deletion feature for a DFC session. This method throws `DfException` with the `DFC_REALTIME_CONT_DEL_UN SUPPORT` error if Content Server (an earlier version, such as 6.7) does not support this feature.

You can use the following queries to search for the contents that are not deleted successfully:

- For real-time content deletion failures:

```
Select r_object_id from dmr_content where content_state = 2
```

- For conventional (asynchronous) content deletion failures

```
Select r_object_id from dmr_content where content_state = 1
```

- For all deletion failures:

```
Select r_object_id from dmr_content where content_state <> 0
```

Working with Document Operations

This chapter describes the way to use DFC to perform the most common operations on documents. Most information about documents also applies to the broader category of repository objects represented by the `IDfSysObject` interface.

The chapter contains the following main sections:

- [Understanding documents, page 83](#)
- [Understanding operations, page 85](#)
- [Types of operation, page 86](#)
- [Basic steps for manipulating documents, page 87](#)
- [Operations for manipulating documents, page 90](#)
- [Handling document manipulation errors, page 122](#)
- [Operations and transactions, page 124](#)

Understanding documents

The *Documentum Content Server Fundamentals* manual explains Documentum facilities for managing documents. This section provides a concise summary of what you need to know to understand the remainder of this chapter.

Documentum maintains a repository of objects that it classifies according to a type hierarchy. For this discussion, *SysObjects* are at the top of the hierarchy. A *document* is a specific kind of `SysObject`. Its primary purpose is to help you manage content.

Documentum maintains more than one version of a document. A *version tree* is an original document and all of its versions. Every version of the document has a unique object ID, but every version has the same *chronicle ID*, namely, the object ID of the original document.

Virtual documents

A *virtual document* is a container document that includes one or more objects, called *components*, organized in a tree structure. A component can be another virtual document or a simple document. A

virtual document can have any number of components, nested to any level. Documentum imposes no limit on the depth of nesting in a virtual document.

Documentum uses two sets of terminology for virtual documents. In the first set, a virtual document that contains a component is called the component's *parent*, and the component is called the virtual document's *child*. Children, or children of children to any depth, are called *descendants*.

Note: Internal variables, Javadoc comments, and registry keys sometimes use the alternate spelling *descendent*.

The second set of terminology derives from graph theory, even though a virtual document forms a tree, and not an arbitrary graph. The virtual document and each of its descendants is called a *node*. The directed relationship between a parent node and a child node is called an *edge*.

In both sets of terminology, the original virtual document is sometimes called the *root*.

You can associate a particular version of a component with the virtual document (this is called *early binding*) or you can associate the component's entire version tree with the virtual document. The latter allows you to select which version to include at the time you construct the document (this is called *late binding*).

Documentum provides a flexible set of rules for controlling the way it assembles documents. An *assembly* is a snapshot of a virtual document. It consists of the set of specific component versions that result from assembling the virtual document according to a set of binding rules. To preserve it, you must attach it to a SysObject: usually either the root of the virtual document or a SysObject created to hold the assembly. A SysObject can have at most one attached assembly.

You can version a virtual document and manage its versions just as you do for a simple document. Deleting a virtual document version also removes any containment objects or assembly objects associated with that version.

When you copy a virtual document, the server can make a copy of each component, or it can create an internal reference or pointer to the source component. It maintains information in the containment object about which of these possibilities to choose. One option is to require the copy operation to specify the choice.

Whether it copies a component or creates a reference, Documentum creates a new containment object corresponding to that component.

Note: DFC allows you to process the root of a virtual document as an ordinary document. For example, suppose that *doc* is an object of type IDfDocument and also happens to be the root of a virtual document. If you tell DFC to check out *doc*, it does not check out any of the descendants. If you want DFC to check out the descendants along with the root document, you must first execute an instruction like

```
IDfVirtualDocument vDoc =  
    doc.asVirtualDocument(CURRENT, false)
```

If you tell DFC to check out vDoc, it processes the current version of doc and each of its descendants. The DFC Javadocs explain the parameters of the asVirtualDocument method.

Documentum represents the nodes of virtual documents by *containment objects* and the nodes of assemblies by *assembly objects*. An assembly object refers to the SysObject to which the assembly is attached, and to the virtual document from which the assembly came.

If an object appears more than once as a node in a virtual document or assembly, each node has a separate associated containment object or assembly object. No object can appear as a descendant of itself in a virtual document.

XML Documents

Documentum's XML support has many features. Information about those subjects appears in *Documentum Content Server Fundamentals* and in the *XML Application Development Guide*.

Using XML support requires you to provide a controlling XML application. When you import an XML document, DFC examines the controlling application's configuration file and applies any chunking rules that you specify there.

If the application's configuration file specifies chunking rules, DFC creates a virtual document from the chunks it creates. It imports other documents that the XML document refers to as entity references or links, and makes them components of the virtual document. It uses attributes of the containment object associated with a component to remember whether it came from an entity or a link and to maintain other necessary information. Assembly objects have the same XML-related attributes as containment objects do.

Understanding operations

Operations are used to manipulate documents in Documentum. Operations provide interfaces and a processing environment to ensure that Documentum can handle a variety of documents and collections of documents in a standard way. You obtain an operation of the appropriate kind, place one or more documents into it, and execute the operation.

All of the examples in this chapter pertain only to documents, but operations can be used to work with objects of type `IDfSysObject`, not just the subtype `IDfDocument`.

For example, to check out a document, take the following steps:

1. Obtain a checkout operation.
2. Add the document to the operation.
3. Execute the operation.

DFC carries out the behind-the-scenes tasks associated with checking out a document. For a virtual document, for example, DFC adds all of its components to the operation and ensures that links between them are still valid when it stores the documents into the checkout directory on the file system. It corrects filename conflicts, and it keeps a local record of which documents it checks out. This is only a partial description of what DFC does when you check out a document. Because of the number and complexity of the underlying tasks, DFC wraps seemingly elementary document-manipulation tasks in operations.

Note: A `SysObject` in the operations framework is processed only once. If you add the same object twice to an operation, it only performs the operation steps once per object. For example, if you add a document object multiple times in order to retrieve multiple content pages with one operation, you will only retrieve one page of content. In order to retrieve multiple content pages (for example, renditions), you must use a separate operation for each operation.

An IDfClientX object provides factory methods for creating operations. Once you have an IDfClientX object (say cX) and a SysObject (say doc) representing the document, the code for the checkout looks like this:

```
// Obtain a checkout operation
IDfCheckoutOperation checkout =cX.getCheckoutOperation();

// Add the document to the checkout operation
checkout.add(doc);    //This might fail and return a null

// Check the document out
checkout.execute();    //This might produce errors without
//throwing an exception
```

In your own applications, you would add code to handle a null returned by the add method or errors produced by the execute method.

Types of operation

DFC provides operation types and corresponding nodes (to be explained in subsequent sections) for many tasks you perform on documents or, where appropriate, files or folders. The following table summarizes these.

Table 1. DFC operation types and nodes

Task	Operation Type	Operation Node Type
Import into a repository	IDfImportOperation	IDfImportNode
Export from a repository	IDfExportOperation	IDfExportNode
Check into a repository	IDfCheckinOperation	IDfCheckinNode
Check out of a repository	IDfCheckoutOperation	IDfCheckoutNode
Cancel a checkout	IDfCancelCheckoutOperation	IDfCancelCheckoutNode
Delete from a repository	IDfDeleteOperation	IDfDeleteNode
Copy from one repository location to another	IDfCopyOperation	IDfCopyNode
Move from one repository location to another	IDfMoveOperation	IDfMoveNode
Validate an XML document against a DTD or Schema	IDfValidationOperation	IDfValidationNode
Transform an XML document using XSLT	IDfXMLTransformOperation	IDfXMLTransformNode
Pre-cache objects in a BOCS repository	IDfPredictiveCachingOperation	IDfPredictiveCachingNode

Basic steps for manipulating documents

This section describes the basic steps common to using the facilities of the operations package to manipulate documents. It sets forth the basic steps, then discusses the steps in greater detail.

Steps for manipulating documents

This section describes the basic steps common to document manipulation operations. [Details of manipulating documents, page 88](#) provides more detailed information about these steps.

To perform a document-manipulation task:

1. Use the appropriate factory method of IDfClientX to obtain an operation of the type appropriate to the document-manipulation task.
 For example, if you want to check documents into a repository, start by calling `getCheckinOperation`.
Note: Each operation has a type (for example, `IDfCheckinOperation`) that inherits most of its methods (in particular, its `add` and `execute` methods) from `IDfOperation`.
2. Set parameters to control the way the operation performs the task.
 Each operation type has `setXxx` methods for setting its parameters.
 The operation behaves in predefined (default) ways if you do not set optional parameters. Some parameters (the session for an import operation, for example) are mandatory.
3. Add documents to the operation:
 - a. Use its inherited `add` method to place a document, file, or folder into the operation.
 The `add` method returns the newly created node, or a null if it fails (refer to [Handling document manipulation errors, page 122](#)).
 - b. Set parameters to change the way the operation handles this item and its descendants.
 Each type of operation node has methods for setting parameters that are important for that type of node. These are generally the same as the methods for the corresponding type of operation. If you do not set parameters, the operation handles this item according to the `setXxx` methods.
 - c. Repeat the previous two substeps for all items you add to the operation.
4. Invoke the operation's inherited `execute` method to perform the task.
 Note that this step may add and process additional nodes. For example, if part of the execution entails scanning an XML document for links, DFC may add the linked documents to the operation.
 The `execute` method returns a boolean value to indicate its success (true) or failure (false). See [Handling document manipulation errors, page 122](#)) for more information.
5. Process the results.
 - a. Handle errors.
 If it detects errors, the `execute` method returns the boolean value *false*. You can use the operation's inherited `getErrors` method to obtain a list of failures.

For details of how to process errors, see [Processing the results, page 89](#).

- b. Perform tasks specific to the operation.

For example, after an import operation, you may want to take note of all of the new objects that the operation created in the repository. You might want to display or modify their properties.

Details of manipulating documents

This section discusses some issues and background for the steps of the general procedure in [Steps for manipulating documents, page 87](#).

Obtaining the operation

Each operation factory method of IDfClientX instantiates an operation object of the corresponding type. For example, getImportOperation factory method instantiates an IDfImportOperation object.

Setting parameters for the operation

Different operations accept different parameters to control the way they carry out their tasks. Some parameters are optional, some mandatory.

Note: You must use the setSession method of IDfImportOperation or IDfXMLTransformOperation to set a repository session before adding nodes to either of these types of operation.

Adding documents to the operation

An operation contains a structure of nodes and descendants. When you obtain the operation, it has no nodes. When you use the operation's add method to include documents in the operation, it creates new root nodes. The add method returns the node as an IDfOperationNode object. You must cast it to the appropriate operation node type to use any methods the type does not inherit from IDfOperationNode (see [Working with nodes, page 89](#)).

Note: If the add method cannot create a node for the specified document, it returns a null argument. Be sure to test for this case, because it does not usually throw an exception.

DFC might include additional nodes in the operation. For example, if you add a repository folder, DFC adds nodes for the documents linked to that folder, as children of the folder's node in the operation.

Each node can have zero or more child nodes. If you add a virtual document, the add method creates as many descendant nodes as necessary to create an image of the virtual document's structure within the operation.

You can add objects from more than one repository to an operation.

You can use a variety of methods to obtain and step through all nodes of the operation (see [Working with nodes, page 89](#)). You might want to set parameters on individual nodes differently from the way you set them on the operation.

Executing the Operation

The operations package processes the objects in an operation as a group, possibly invoking many DFC calls for each object. Operations encapsulate Documentum client conventions for registering, naming, and managing local content files.

DFC executes the operation in a predefined set of steps, applying each step to all of the documents in the operation before proceeding to the next step. It processes each document in an operation only once, even if the document appears at more than one node.

Once DFC has executed a step of the operation on all of the documents in the operation, it cannot execute that step again. If you want to perform the same task again, you must construct a new operation to do so.

Normally, you use the operation's `execute` method and let DFC proceed through the execution steps. DFC provides a limited ability for you to execute an operation in steps, so that you can perform special processing between steps. Documentum does not recommend this approach, because the number and identity of steps in an operation may change with future versions of DFC. If you have a programming hurdle that you cannot get over without using steps, work with Documentum Technical Support or Consulting to design a solution.

Processing the results

If DFC encounters an error while processing one node in an operation, it continues to process the other nodes. For example, if one object in a checkout operation is locked, the operation checks out the others. Only fatal conditions cause an operation to throw an exception. DFC catches other exceptions internally and converts them into `IDfOperationError` objects. The `getErrors` method returns an `IDfList` object containing those errors, or a null if there are no errors. The calling program can examine the errors, and decide whether to undo the operation, or to accept the results for those objects that did not generate errors.

Once you have checked the errors you may want to examine and further process the results of the operation. The next section, [Working with nodes, page 89](#), shows how to access the objects and results associated with the nodes of the operation.

Working with nodes

This section shows how to access the objects and results associated with the nodes of an operation.

Note: Each operation node type (for example, `IDfCheckinNode`) inherits most of its methods from `IDfOperationNode`.

The `getChildren` method of an `IDfOperationNode` object returns the first level of nodes under the given node. You can use this method recursively to step through all of the descendant nodes.

Alternatively, you can use the operation's `getNodes` method to obtain a flat list of descendant nodes, that is, an `IDfList` object containing of all of its descendant nodes without the structure.

These methods return nodes as objects of type `IDfOperationNode`, not as the specific node type (for example, `IDfCheckinNode`).

The `getId` method of an `IDfOperationNode` object returns a unique identifier for the node, *not the object ID of the corresponding document*. `IDfOperationNode` does not have a method for obtaining the object ID of the corresponding object. Each operation node type (for example, `IDfCheckinNode`) has its own `getObjectID` method. You must cast the `IDfOperationNode` object to a node of the specific type before obtaining the object ID.

Operations for manipulating documents

This section provides sample code and discusses specific details of the following kinds of document manipulation operations:

- [Checking out, page 92](#)
- [Checking in, page 96](#)
- [Cancelling checkout, page 99](#)
- [Importing, page 103](#)
- [Exporting, page 106](#)
- [Copying, page 109](#)
- [Moving, page 112](#)
- [Deleting, page 114](#)
- [Predictive caching, page 117](#)
- [Validating an XML document against a DTD or schema, page 118](#)
- [Performing an XSL transformation of an XML document, page 119](#)

The examples use the terms *file* and *directory* to refer to entities on the file system and the terms *document* and *folder* to repository entities represented by DFC objects of type `IDfDocument` and `IDfFolder`.

These examples assume the use of a simple UI, similar to the one described in [Chapter 4, Creating a Test Application](#), which enables the user to pass a Session Manager and additional information such as a document ID or directory path to give the essential information required by the operation. The figure shows the base tutorial extended with the operations discussed in this chapter.

Figure 3. The DFC operations tutorial frame

The screenshot shows a window titled "DFC Operations Tutorial Frame" with a standard Windows-style title bar (minimize, maximize, close buttons). The window is divided into two main panels. The left panel contains a form with three input fields: "Repository" (containing "nyRepository"), "User Name" (containing "myName"), and "Password" (containing "myPasswo"). To the right of these fields is a button labeled "Get Sess Mgr". Below these fields is a text box labeled "Current Working Directory" containing "/myCabinet". Underneath is a list box labeled "Directory contents:" containing two items: "myDocument" and "myFolder". Below the list box are two buttons: "Get Directory" and "Up One Level". At the bottom of the left panel is a large empty text box labeled "Results:". At the very bottom of the left panel, the text "Messages appear here." is displayed. The right panel contains several groups of buttons and text boxes. At the top are "Check Out" and "Check In" buttons, with a "Cancel Checkout" button below them. Below these is a text box labeled "Destination Directory in Repository:" followed by a "Copy" button and a "Move" button. To the right of the "Move" button is a button labeled "Set Copy/Move Destination". Below these is a "Delete" button and a checkbox labeled "Delete Current Version Only" which is checked. At the bottom of the right panel are two more sections: "Local Directory:" with an "Export" button and a text box, and "Path to source document:" with an "Import" button and a text box.

DFC Operations Tutorial Frame

Repository: nyRepository User Name: myName Password: myPasswo Get Sess Mgr

Current Working Directory: /myCabinet

Directory contents:

- myDocument
- myFolder

Get Directory Up One Level

Results:

Messages appear here.

Check Out Check In

Cancel Checkout

Destination Directory in Repository:

Copy Move Set Copy/Move Destination

Delete ☒ Delete Current Version Only

Local Directory:

Export

Path to source document:

Import

Checking out

The execute method of an IDfCheckoutOperation object checks out the documents in the operation. The checkout operation:

- Locks the document
- Copies the document to your local disk
- Always creates registry entries to enable DFC to manage the files it creates on the file system

Example 6-1. Handler for the Check Out button

```
private void jButton_checkOut_actionPerformed(ActionEvent e)
{
    String repositoryName = jTextField_repositoryName.getText();
    String docId =
        m_fileIDs.elementAt(list_id.getSelectedIndex()).toString();
    TutorialCheckOut tco = new TutorialCheckOut();
    jLabel_messages.setText(
        tco.checkoutExample(
            m_sessionManager,
            repositoryName,
            docId
        )
    );
    initDirectory();
    getDirectory();
}
```

Example 6-2. TutorialCheckout.java

```
package com.emc.tutorial;

import com.documentum.com.DfClientX;
import com.documentum.com.IDfClientX;
import com.documentum.fc.client.IDfDocument;
import com.documentum.fc.client.IDfSession;
import com.documentum.fc.client.IDfSessionManager;
import com.documentum.fc.client.IDfSysObject;
import com.documentum.fc.common.DfId;
import com.documentum.fc.common.IDfId;
import com.documentum.operations.IDfCheckoutNode;
import com.documentum.operations.IDfCheckoutOperation;

public class TutorialCheckOut
{
    public TutorialCheckOut()
    {
    }

    public String checkoutExample
    (
        IDfSessionManager sessionManager,
        String repositoryName,
        String docId
    )
    {
        StringBuffer result = new StringBuffer("");
        IDfSession mySession = null;

        try
```

```

{
    mySession = sessionManager.getSession(repositoryName);

    // Get the object ID based on the object ID string.
    IDfId idObj =
        mySession.getIdByQualification(
            "dm_sysobject where r_object_id='" + docId + "'"
        );

    // Instantiate an object from the ID.

    IDfSysObject sysObj = (IDfSysObject) mySession.getObject(idObj);

    // Instantiate a client.
    IDfClientX clientx = new DfClientX();

    // Use the factory method to create a checkout operation object.
    IDfCheckoutOperation coOp = clientx.getCheckoutOperation();

    // Set the location where the local copy of the checked out file
    // is stored.
    coOp.setDestinationDirectory("C:\\");

    // Get the document instance using the document ID.
    IDfDocument doc =
        (IDfDocument) mySession.getObject(new DfId(docId));

    // Create the checkout node by adding the document to the checkout
    // operation.
    IDfCheckoutNode coNode = (IDfCheckoutNode) coOp.add(doc);

    // Verify that the node exists.
    if (coNode == null)
    {
        result.append("coNode is null");
    }

    // Execute the checkout operation. Return the result.
    if (coOp.execute())
    {
        result.append("Successfully checked out file ID: " + docId);
    }
    else
    {
        result.append("Checkout failed.");
    }
    return result.toString();
}
catch (Exception ex)
{
    ex.printStackTrace();
    return "Exception has been thrown: " + ex;
}
finally
{
    sessionManager.release(mySession);
}
}
}

```

Special considerations for checkout operations

Follow the steps in [Steps for manipulating documents, page 87](#).

If any node corresponds to a document that is already checked out, the system does not check it out again. DFC does not treat this as an error. If you cancel the checkout, however, DFC cancels the checkout of the previously checked out node as well.

DFC applies XML processing to XML documents. If necessary, it modifies the resulting files to ensure that it has enough information to check in the documents properly.

You can use many of the same methods for setting up checkout operations and processing results that you use for export operations.

Checking out a virtual document

If the operation's add method receives a virtual document as an argument, it also adds all of the document's descendants (determined by applying the applicable binding rules), creating a separate node for each.

Example 6-3. The TutorialCheckoutVdm class

```
package dfctestenvironment;

import com.documentum.com.DfClientX;
import com.documentum.com.IDfClientX;
import com.documentum.fc.client.IDfDocument;
import com.documentum.fc.client.IDfSession;
import com.documentum.fc.client.IDfVirtualDocument;
import com.documentum.fc.common.DfId;
import com.documentum.operations.IDfCheckoutNode;
import com.documentum.operations.IDfCheckoutOperation;

public class TutorialCheckoutVdm {
    public TutorialCheckoutVdm() {}

    public String checkoutExample(
        IDfSessionManager sessionManager,
        String repositoryName,
        String docId)
    {
        try {
            String result = "";

            // Instantiate a session.
            IDfSession mySession = sessionManager.getSession (repositoryName);
            // Instantiate a client.
            IDfClientX clientx = new DfClientX();

            // Use the factory method to create a checkout operation object.
            IDfCheckoutOperation coOp = clientx.getCheckoutOperation();

            // Set the location where the local copy of the checked out file is stored.
            coOp.setDestinationDirectory("C:\\");

            // Get the document instance using the document ID.
            IDfDocument doc = (IDfDocument) mySession.getObject(new DfId(docId));
```

```
// Create an empty checkout node object.
    IDfCheckoutNode coNode;

// If the doc is a virtual document, instantiate it as a virtual document
// object and add it to the checkout operation. Otherwise, add the document
// object to the checkout operation.
    if (doc.isVirtualDocument()){
        IDfVirtualDocument vDoc = doc.asVirtualDocument( "CURRENT",false);
        coNode = (IDfCheckoutNode) coOp.add(vDoc);
    }
    else {
        coNode = (IDfCheckoutNode) coOp.add(doc);
    }

// Verify that the node exists.
    if (coNode == null) {
        result = ("coNode is null");
    }

// Execute the checkout operation. Return the result.
    if ( coOp.execute()) {
        result = "Successfully checked out file ID: " + docId;
    }
    else {
        result = ("Checkout failed.");
    }
    return result;
}
catch (Exception ex) {
    ex.printStackTrace();
    return "Exception has been thrown: " + ex;
}
finally {
    sessionManager.release(mySession);
}
}
```

Checking in

The execute method of an IDfCheckinOperation object checks documents into the repository. It creates new objects as required, transfers the content to the repository, and removes local files if appropriate. It checks in existing objects that any of the nodes refer to (for example, through XML links).

Example 6-4. Handler for the Check In button

Check in a document as the next major version (for example, version 1.2 would become version 2.0). The default increment is NEXT_MINOR (for example, version 1.2 would become version 1.3).

```
private void jButton_checkIn_actionPerformed(ActionEvent e)
{
    String repositoryName = jTextField_repositoryName.getText();
    String docId =
        m_fileIDs.elementAt(list_id.getSelectedIndex()).toString();
    TutorialCheckIn tci = new TutorialCheckIn();
    jLabel_messages.setText(
        tci.checkinExample(
            m_sessionManager,
            repositoryName,
            docId
        )
    );
    initDirectory();
    getDirectory();
}
```

Example 6-5. The TutorialCheckIn class

```
package com.emc.tutorial;

import com.documentum.com.DfClientX;
import com.documentum.com.IDfClientX;
import com.documentum.fc.client.IDfDocument;
import com.documentum.fc.client.IDfSession;
import com.documentum.fc.client.IDfSessionManager;
import com.documentum.fc.client.IDfSysObject;
import com.documentum.fc.common.DfId;
import com.documentum.fc.common.IDfId;
import com.documentum.operations.IDfCheckinNode;
import com.documentum.operations.IDfCheckinOperation;

public class TutorialCheckIn
{
    public TutorialCheckIn()
    {
    }

    public String checkinExample(
        IDfSessionManager sessionManager,
        String repositoryName,
        String docId
    )
    {
        IDfSession mySession = null;
        try
        {
            mySession = sessionManager.getSession(repositoryName);

            // Get the object ID based on the object ID string.
```



```

        IDfId idObj =
            mySession.getIdByQualification(
                "dm_sysobject where r_object_id='" + docId + "'"
            );

// Instantiate an object from the ID.

        IDfSysObject sysObj = (IDfSysObject) mySession.getObject(idObj);

// Instantiate a client.
        IDfClientX clientx = new DfClientX();

// Use the factory method to create an IDfCheckinOperation instance.
        IDfCheckinOperation cio = clientx.getCheckinOperation();

// Set the version increment. In this case, the next major version
// ( version + 1)
        cio.setCheckinVersion(IDfCheckinOperation.NEXT_MAJOR);

// When updating to the next major version, you need to explicitly
// set the version label for the new object to "CURRENT".
        cio.setVersionLabels("CURRENT");

// Create a document object that represents the document being
// checked in.
        IDfDocument doc =
            (IDfDocument) mySession.getObject(new DfId(docId));

// Create a checkin node, adding it to the checkin operation.
        IDfCheckinNode node = (IDfCheckinNode) cio.add(doc);

// Execute the checkin operation and return the result.
        if (!cio.execute())
        {
            return "Checkin failed.";
        }

// After the item is created, you can get it immediately using the
// getNewObjectId method.

        IDfId newId = node.getNewObjectId();
        return "Checkin succeeded - new object ID is: " + newId;
    }
    catch (Exception ex)
    {
        ex.printStackTrace();
        return "Checkin failed.";
    }
    finally
    {
        sessionManager.release(mySession);
    }
}
}

```

Special considerations for checkin operations

Follow the steps in [Steps for manipulating documents, page 87](#).

Setting up the operation

To check in a document, you pass an object of type `IDfSysObject` or `IDfVirtualDocument`, *not the file on the local file system*, to the operation's add method. In the local client file registry, DFC records the path and filename of the local file that represents the content of an object. If you move or rename the file, DFC loses track of it and reports an error when you try to check it in.

Setting the content file, as in `IDfCheckinNode.setFilePath`, overrides DFC's saved information.

If you specify a document that is not checked out, DFC does not check it in. DFC does not treat this as an error.

You can specify checkin version, symbolic label, or alternate content file, and you can direct DFC to preserve the local file.

If between checkout and checkin you remove a link between documents, DFC adds the orphaned document to the checkin operation as a root node, but the relationship between the documents no longer exists in the repository.

Processing the checked in documents

Executing a checkin operation normally results in the creation of new objects in the repository. If `opCheckin` is the `IDfCheckinOperation` object, you can obtain a complete list of the new objects by calling

```
IDfList list = opCheckin.getNewObjects();
```

The list contains the object IDs of the newly created `SysObjects`.

In addition, the `IDfCheckinNode` objects associated with the operation are still available after you execute the operation (see [Working with nodes, page 89](#)). You can use their methods to find out many other facts about the new `SysObjects` associated with those nodes.

Behavior Change after Failed Save or Checkin

In Documentum 6.5 and later, if a save or checkin of an object fails, DFC reverts the object automatically, before returning the object to the application issuing the save or checkin. Reverting the object removes any changes made to the object before the attempted save or checkin. The application must then reapply any changes made before the save or checkin before reattempting the save or checkin operation. This behavior is different from how earlier DMCL releases handled failed saves or checkins. The DMCL in prior releases simply marked the object with an error state and returned it to the application.

Cancelling checkout

The execute method of an IDfCancelCheckoutOperation object cancels the checkout of documents by releasing locks, deleting local files if appropriate, and removing registry entries.

If the operation's add method receives a virtual document as an argument, it also adds all of the document's descendants (determined by applying the applicable binding rules), creating a separate operation node for each.

Example 6-6. Handler for the Cancel Checkout button

```
private void jButton_cancelCheckOut_actionPerformed(ActionEvent e)
{
    String repositoryName = jTextField_repositoryName.getText();
    String docId =
        m_fileIDs.elementAt(list_id.getSelectedIndex()).toString();
    TutorialCancelCheckOut tcco = new TutorialCancelCheckOut();
    JLabel_messages.setText(
        tcco.cancelCheckOutExample(
            m_sessionManager,
            repositoryName,
            docId
        )
    );
}
```

Example 6-7. TutorialCancelCheckout.java

Cancel checkout of a document.

```
package com.emc.tutorial;

import com.documentum.com.DfClientX;
import com.documentum.com.IDfClientX;
import com.documentum.fc.client.IDfDocument;
import com.documentum.fc.client.IDfSession;
import com.documentum.fc.client.IDfSessionManager;
import com.documentum.fc.client.IDfSysObject;
import com.documentum.fc.common.DfException;
import com.documentum.fc.common.DfId;
import com.documentum.fc.common.IDfId;
import com.documentum.operations.IDfCancelCheckoutNode;
import com.documentum.operations.IDfCancelCheckoutOperation;

public class TutorialCancelCheckOut
{
    public TutorialCancelCheckOut()
    {
    }

    public String cancelCheckOutExample(
        IDfSessionManager sessionManager,
        String repositoryName,
        String docId
    )
    {
        IDfSession mySession = null;

        try
        {
            mySession = sessionManager.getSession(repositoryName);
```

```
// Get the object ID based on the object ID string.
    IDfId idObj =
        mySession.getIdByQualification(
            "dm_sysobject where r_object_id='" + docId + "'"
        );

// Instantiate an object from the ID.

    IDfSysObject sysObj = (IDfSysObject) mySession.getObject(idObj);

// Get a new client instance.
    IDfClientX clientx = new DfClientX();

// Use the factory method to create a checkout operation object.
    IDfCancelCheckoutOperation cco =
        clientx.getCancelCheckoutOperation();

// Instantiate the document object from the ID string.
    IDfDocument doc =
        (IDfDocument) mySession.getObject(new DfId(docId));

// Indicate whether to keep the local file.
    cco.setKeepLocalFile(true);

// Create an empty cancel checkout node.
    IDfCancelCheckoutNode node;

// Populate the cancel checkout node and add it to the cancel checkout
// operation.
    node = (IDfCancelCheckoutNode) cco.add(doc);

// Check to see if the node is null - this will not throw an error.
    if (node == null)
    {
        return "Node is null";
    }

    // Execute the operation and return the result.
    if (!cco.execute())
    {
        return "Operation failed";
    }
    return "Successfully cancelled checkout of file ID: " + docId;
}

// Handle any exceptions.
catch (Exception ex)
{
    ex.printStackTrace();
    return "Exception has been thrown: " + ex;
}

// Always, always, release the session in the "finally" clause.
finally
{
    sessionManager.release(mySession);
}
}
```

Special considerations for cancel checkout operations

Follow the steps in [Steps for manipulating documents, page 87](#).

If a document in the cancel checkout operation is not checked out, DFC does not process it. DFC does not treat this as an error.

Cancel checkout for virtual document

If the operation's add method receives a virtual document as an argument, it also adds all of the document's descendants (determined by applying the applicable binding rules), creating a separate operation node for each.

Example 6-8. The TutorialCancelCheckoutVdm class

```
package dfctestenvironment;

import com.documentum.com.DfClientX;
import com.documentum.com.IDfClientX;
import com.documentum.fc.client.IDfDocument;
import com.documentum.fc.client.IDfSession;
import com.documentum.fc.client.IDfVirtualDocument;
import com.documentum.fc.common.DfException;
import com.documentum.fc.common.DfId;
import com.documentum.operations.IDfCancelCheckoutNode;
import com.documentum.operations.IDfCancelCheckoutOperation;

public class TutorialCancelCheckoutVdm {
    public TutorialCancelCheckoutVdm() {
    }
    public String cancelCheckoutExample(
        IDfSession mySession,
        String docId) throws DfException
    {
        try
        {
            // Get a new client instance.
            IDfClientX clientx = new DfClientX();

            // Use the factory method to create a checkout operation object.
            IDfCancelCheckoutOperation cco =
                clientx.getCancelCheckoutOperation();

            // Instantiate the document object from the ID string.
            IDfDocument doc =
                (IDfDocument) mySession.getObject(new DfId(docId));

            // Indicate whether to keep the local file.
            cco.setKeepLocalFile(true);

            // Create an empty cancel checkout node.
            IDfCancelCheckoutNode node;

            // If it is a virtual document, instantiate it as a virtual document and add
            // the virtual document to the operation. Otherwise, add the doc to the
            // operation.
            if (doc.isVirtualDocument()) {
                IDfVirtualDocument vdoc = doc.asVirtualDocument("CURRENT", false);
```

```
        node = (IDfCancelCheckoutNode)cco.add(vdoc);
    }
    else
    {
        node = (IDfCancelCheckoutNode)cco.add(doc);
    }

    // Check to see if the node is null - this will not throw an error.
    if (node==null) {return "Node is null";}

    // Execute the operation and return the result.
    if (!cco.execute()){
        return "Operation failed";
    }
    return "Successfully cancelled checkout of file ID: " + docId;
}
catch (Exception e){
    e.printStackTrace();
    return "Exception thrown.";
}
}
```

Importing

The execute method of an IDfImportOperation object imports files and directories into the repository. It creates objects as required, transfers the content to the repository, and removes local files if appropriate. If any of the nodes of the operation refer to existing local files (for example, through XML or OLE links), it imports those into the repository too.

Example 6-9. Handler for the Import button

```
private void jButton_import_actionPerformed(ActionEvent e)
{
    String repositoryName = jTextField_repositoryName.getText();
    String sourceFilePath = jTextField_importFilePath.getText();
    String destinationDirectory = jTextField_cwd.getText();
    TutorialImport ti = new TutorialImport();
    jLabel_messages.setText(
        ti.importExample(
            m_sessionManager,
            repositoryName,
            destinationDirectory,
            sourceFilePath
        )
    );
    initDirectory();
    getDirectory();
}
```

Example 6-10. TutorialImport.java

```
package com.emc.tutorial;

import com.documentum.com.DfClientX;
import com.documentum.com.IDfClientX;
import com.documentum.fc.client.IDfFolder;
import com.documentum.fc.client.IDfSession;
import com.documentum.fc.client.IDfSessionManager;
import com.documentum.fc.common.DfException;
import com.documentum.operations.IDfFile;
import com.documentum.operations.IDfImportNode;
import com.documentum.operations.IDfImportOperation;

public class TutorialImport
{
    public TutorialImport()
    {
    }

    public String importExample(
        IDfSessionManager sessionManager,
        String repositoryName,
        String destinationDirectory,
        String sourceFilePath
    )
    {
        IDfSession mySession = null;

        try
        {
            mySession = sessionManager.getSession(repositoryName);

            // Create a new client instance.
```

```
        IDfClientX clientx = new DfClientX();

// Use the factory method to create an IDfImportOperation instance.
        IDfImportOperation opi = clientx.getImportOperation();

// You must explicitly set the session for an import operation.
        opi.setSession(mySession);

// Create an instance of the target folder.
        IDfFolder folder = mySession.getFolderByPath(destinationDirectory);

// Create a file instance for the local file.
        IDfFile file = clientx.getFile(sourceFilePath);
        if (!file.exists())
            return ("File does not exist.");

// Set the destination folder.
        opi.setDestinationFolderId(folder.getObjectId());

// Create the import node, adding the file to the import operation.
        IDfImportNode node = (IDfImportNode) opi.add(file);
        if (node == null)
            return ("Node is null.");

// Execute the import operation and return the results.
        if (opi.execute())
        {
            String resultString =
                ("Item" + opi.getNewObjects().toString() +
                 " imported successfully.");
            return resultString;
        }
        else
        {
            return ("Error during import operation.");
        }
    }

// Handle any exceptions.
    }
    catch (Exception ex)
    {
        ex.printStackTrace();
        return "Exception has been thrown: " + ex;
    }
}

// Always, always, release the session in the "finally" clause.
finally
{
    {
        sessionManager.release(mySession);
    }
}
}
```

Special Considerations for Import Operations

Follow the steps in [Steps for manipulating documents, page 87](#).

Setting up the operation

Use the object's `setSession` method to specify a repository session and the object's `setDestinationFolderId` method to specify the repository cabinet or folder into which the operation should import documents.

You *must* set the session before adding files to the operation.

You can set the destination folder, either on the operation or on each node. The node setting overrides the operation setting. If you set neither, DFC uses its default destination folder.

You can add an `IDfFile` object or specify a file system path. You can also specify whether to keep the file on the file system (the default choice) or delete it after the operation is successful.

If you add a file system directory to the operation, DFC imports all files in that directory and proceeds recursively to add each subdirectory to the operation. The resulting repository folder hierarchy mirrors the file system directory hierarchy.

You can also control version labels, object names, object types and formats of the imported objects.

If you are importing a document with OLE links, all of the linked files can be imported for you automatically and stored as nodes in a virtual document. See [Microsoft Object Linking and Embedding \(OLE\)](#), page 20 for information on configuring your system to accommodate OLE links.

XML processing

You can import XML files without doing XML processing. If `nodeImport` is an `IDfImportNode` object, you can turn off XML processing on the node and all its descendants by calling

```
nodeImport.setXMLApplicationName("Ignore");
```

Turning off this kind of processing can shorten the time it takes DFC to perform the operation.

Processing the imported documents

Executing an import operation results in the creation of new objects in the repository. If `opImport` is the `IDfImportOperation` object, you can obtain a complete list of the new objects by calling

```
IDfList list = opImport.getNewObjects();
```

The list contains the object IDs of the newly created `SysObjects`.

In addition, the `IDfImportNode` objects associated with the operation are still available after you execute the operation (see [Working with nodes](#), page 89). You can use their methods to find out many other facts about the new `SysObjects` associated with those nodes. For example, you can find out object IDs, object names, version labels, file paths, and formats.

Exporting

The execute method of an IDfExportOperation object creates copies of documents on the local file system. If the operation's add method receives a virtual document as an argument, it also adds all of the document's descendants (determined by applying the applicable binding rules), creating a separate node for each.

Example 6-11. Handler for the Export button

```
private void jButton_export_actionPerformed(ActionEvent e)
{
    String repositoryName = jTextField_repositoryName.getText();
    String docId =
        m_fileIDs.elementAt(list_id.getSelectedIndex()).toString();
    String targetLocalDirectory = jTextField_localDirectory.getText();
    TutorialExport te = new TutorialExport();
    jLabel_messages.setText(
        te.exportExample(
            m_sessionManager,
            repositoryName,
            docId,
            targetLocalDirectory)
    );
}
```

Example 6-12. TutorialExport.java

This example does not create registry information about the resulting file.

```
package com.emc.tutorial;

import com.documentum.com.DfClientX;
import com.documentum.com.IDfClientX;
import com.documentum.fc.client.IDfDocument;
import com.documentum.fc.client.IDfFormat;
import com.documentum.fc.client.IDfSession;
import com.documentum.fc.client.IDfSessionManager;
import com.documentum.fc.client.IDfSysObject;
import com.documentum.fc.common.DfId;
import com.documentum.fc.common.IDfId;
import com.documentum.operations.IDfExportNode;
import com.documentum.operations.IDfExportOperation;

public class TutorialExport
{
    public TutorialExport()
    {
    }

    public String exportExample(
        IDfSessionManager sessionManager,
        String repositoryName,
        String docId,
        String targetLocalDirectory
    )
    {
        IDfSession mySession = null;
        StringBuffer sb = new StringBuffer("");

        try
        {
            mySession = sessionManager.getSession(repositoryName);
```

```

// Get the object ID based on the object ID string.
    IDfId idObj =
        mySession.getIdByQualification(
            "dm_sysobject where r_object_id='" + docId + "'"
        );

// Instantiate an object from the ID.

    IDfSysObject sysObj = (IDfSysObject) mySession.getObject(idObj);

// Create a new client instance.
    IDfClientX clientx = new DfClientX();

// Use the factory method to create an IDfExportOperation instance.
    IDfExportOperation eo = clientx.getExportOperation();

// Create a document object that represents the document being exported.
    IDfDocument doc =
        (IDfDocument) mySession.getObject(new DfId(docId));

// Create an export node, adding the document to the export operation object.
    IDfExportNode node = (IDfExportNode) eo.add(doc);

// Get the document's format.
    IDfFormat format = doc.getFormat();

// If necessary, append a path separator to the targetLocalDirectory value.
    if (targetLocalDirectory.lastIndexOf("/") !=
        targetLocalDirectory.length() - 1 &&
        targetLocalDirectory.lastIndexOf("\\") !=
        targetLocalDirectory.length() - 1)
    {
        targetLocalDirectory += "/";
    }

// Set the full file path on the local system.
    node.setFilePath(targetLocalDirectory + doc.getObjectName() + "." +
        format.getDOSExtension());

// Execute and return results
    if (eo.execute())
    {
        return "Export operation successful." + "\n" + sb.toString();
    }
    else
    {
        return "Export operation failed.";
    }
}

// Handle any exceptions.
catch (Exception ex)
{
    ex.printStackTrace();
    return "Exception has been thrown: " + ex;
}

// Always, always, release the session in the "finally" clause.
finally
{
    sessionManager.release(mySession);
}
}

```

Special considerations for export operations

Follow the steps in [Steps for manipulating documents, page 87](#).

By default, an export operation creates files on the local system and makes no provision for Documentum to manage them. You can tell DFC to create registry entries for the files by invoking the `setRecordInRegistry` method of an object of type either `IDfExportOperation` or `IDfExportNode`, using the parameters described in the Javadocs.

If any node corresponds to a checked out document, DFC copies the latest repository version to the local file system. DFC does not treat this as an error.

You can find out where on the file system the export operation creates files. Use the `getDefaultDestinationDirectory` and `getDestinationDirectory` methods of `IDfExportOperation` objects and the `getFilePath` method of `IDfExportNode` objects to do this.

Exporting the contents of a folder requires adding each document individually to the operation.

If you are exporting a virtual document with OLE links, the document can be reassembled for you automatically. See [Microsoft Object Linking and Embedding \(OLE\), page 20](#) for information on configuring your system to accommodate OLE links.

Copying

The execute method of an IDfCopyOperation object copies the *current versions* of documents or folders from one repository location to another.

If the operation's add method receives a virtual document as an argument, it also adds all of the document's descendants (determined by applying the applicable binding rules), creating a separate node of the operation for each.

If the add method receives a folder (unless you override this default behavior), it also adds all documents and folders linked to that folder. This continues recursively until the entire hierarchy of documents and subfolders under the original folder is part of the operation. The execute method replicates this hierarchy at the target location.

Example 6-13. Handler for the Copy button

```
private void jButton_copy_actionPerformed(ActionEvent e)
{
    String repositoryName = jTextField_repositoryName.getText();
    String docId =
        m_fileIDs.elementAt(list_id.getSelectedIndex()).toString();

    // Add a field to the application that allows you to enter a path
    // for the location of copy of the document.
    String destinationDirectory = jTextField_destinationDirectory.getText();
    TutorialCopy tc = new TutorialCopy();
    jLabel_messages.setText(
        tc.copyExample(
            m_sessionManager,
            repositoryName,
            docId,
            destinationDirectory
        )
    );
}
```

Example 6-14. TutorialCopy.java

```
package com.emc.tutorial;

import com.documentum.com.DfClientX;
import com.documentum.com.IDfClientX;
import com.documentum.fc.client.IDfDocument;
import com.documentum.fc.client.IDfFolder;
import com.documentum.fc.client.IDfSession;
import com.documentum.fc.client.IDfSessionManager;
import com.documentum.fc.client.IDfSysObject;
import com.documentum.fc.common.DfId;
import com.documentum.fc.common.IDfId;
import com.documentum.operations.IDfCopyNode;
import com.documentum.operations.IDfCopyOperation;

public class TutorialCopy
{
    public TutorialCopy()
    {
    }

    public String copyExample(
        IDfSessionManager sessionManager,
```

```
        String repositoryName,
        String docId,
        String destination
    )
    {
        IDfSession mySession = null;
        StringBuffer sb = new StringBuffer("");

        try
        {
            mySession = sessionManager.getSession(repositoryName);

            // Get the object ID based on the object ID string.
            IDfId idObj = mySession.getIdByQualification(
                "dm_sysobject where r_object_id='" + docId + "'"
            );

            // Instantiate an object from the ID.

            IDfSysObject sysObj = (IDfSysObject) mySession.getObject(idObj);

            // Create a new client instance.
            IDfClientX clientx = new DfClientX();

            // Use the factory method to create an IDfCopyOperation instance.
            IDfCopyOperation co = clientx.getCopyOperation();

            // Remove the path separator if it exists.
            if (destination.lastIndexOf("/") == destination.length()-1 ||
                destination.lastIndexOf("\\") == destination.length()-1)
            {
                destination = destination.substring(0,destination.length()-1);
            }

            // Create an instance for the destination directory.
            IDfFolder destinationDirectory =
                mySession.getFolderByPath(destination);

            // Set the destination directory by ID.
            co.setDestinationFolderId(destinationDirectory.getObjectId());

            // Create a document object that represents the document being copied.
            IDfDocument doc =
                (IDfDocument) mySession.getObject(new DfId(docId));

            // Create a copy node, adding the document to the copy operation object.
            IDfCopyNode node = (IDfCopyNode) co.add(doc);

            // Execute and return results
            if (co.execute())
            {
                return "Copy operation successful.";
            }
            else
            {
                return "Copy operation failed.";
            }
        }
    }

    // Handle any exceptions.
    catch (Exception ex)
    {
        ex.printStackTrace();
        return "Exception has been thrown: " + ex;
    }
}
```

```
// Always, always, release the session in the "finally" clause.
    finally
    {
        sessionManager.release(mySession);
    }
}
```

Special considerations for copy operations

Follow the steps in [Steps for manipulating documents, page 87](#).

You must set the destination folder, either on the operation or on each of its nodes.

You can use the `setDeepFolders` method of the operation object (node objects do not have this method) to override the default behavior of recursively adding folder contents to the operation.

Certain settings of the attributes of `dm_relation` and `dm_relation_type` objects associated with an object may cause DFC to add related objects to the copy operation. Refer to the *Server Object Reference* manual for details.

Moving

The execute method of an IDfMoveOperation object moves the *current versions* of documents or folders from one repository location to another by unlinking them from the source location and linking them to the destination. Versions other than the current version remain linked to the original location.

If the operation's add method receives a virtual document as an argument, it also adds all of the document's descendants (determined by applying the applicable binding rules), creating a separate node for each.

If the add method receives a folder (unless you override this default behavior), it adds all documents and folders linked to that folder. This continues recursively until the entire hierarchy of documents and subfolders under the original folder is part of the operation. The execute method links this hierarchy to the target location.

Example 6-15. TutorialMove.java

Move a document.

```
package dfctutorialenvironment;
import com.documentum.com.DfClientX;
import com.documentum.com.IDfClientX;
import com.documentum.fc.client.IDfDocument;
import com.documentum.fc.client.IDfFolder;
import com.documentum.fc.client.IDfSession;
import com.documentum.fc.common.DfException;
import com.documentum.fc.common.DfId;
import com.documentum.fc.common.IDfId;
import com.documentum.operations.IDfMoveNode;
import com.documentum.operations.IDfMoveOperation;

public class TutorialMove {
    public TutorialMove() {
    }
    public String moveExample (
        IDfSession mySession,
        String docId,
        String destination
    ) throws DfException {

        // Create a new client instance.
        IDfClientX clientx = new DfClientX();

        // Use the factory method to create an IDfCopyOperation instance.
        IDfMoveOperation mo = clientx.getMoveOperation();

        // Create an instance for the destination directory.
        IDfFolder destinationDirectory = mySession.getFolderByPath(destination);

        // Set the destination directory by ID.
        mo.setDestinationFolderId(destinationDirectory.getObjectId());

        // Create a document object that represents the document being copied.
        IDfDocument doc = (IDfDocument) mySession.getObject(new DfId(docId));

        // Create a move node, adding the document to the move operation object.
        IDfMoveNode node = (IDfMoveNode)mo.add(doc);

        // Execute and return results
        if (mo.execute()) {
            return "Move operation successful.";
        }
    }
}
```



```
        else {  
            return "Move operation failed.";  
        }  
    }  
}
```

Special considerations for move operations

Follow the steps in [Steps for manipulating documents, page 87](#). Options for moving are essentially the same as for copying.

If the operation entails moving a checked out document, DFC leaves the document unmodified and reports an error.

Deleting

The execute method of an IDfDeleteOperation object removes documents and folders from the repository.

If the operation's add method receives a virtual document as an argument, it also adds all of the document's descendants (determined by applying the applicable binding rules), creating a separate node for each. You can use the enableDeepDeleteVirtualDocumentsInFolders method of IDfDeleteOperation to override this behavior.

Example 6-16. Handler for the Delete button

```
private void jButton_delete_actionPerformed(ActionEvent e)
{
    String repositoryName = jTextField_repositoryName.getText();
    String docId =
        m_fileIDs.elementAt(list_id.getSelectedIndex()).toString();

    // Create a checkbox control that lets you set a Boolean value
    // for whether or not the application should destroy all versions
    // or just the current version (current version only = true).
    Boolean currentVersionOnly = jCheckBox_deleteCurrentOnly.isSelected();
    TutorialDelete td = new TutorialDelete();
    jLabel_messages.setText(
        td.deleteExample(
            m_sessionManager,
            repositoryName,
            docId,
            currentVersionOnly
        )
    );
    initDirectory();
    getDirectory();
}
```

Example 6-17. TutorialDelete.java

Delete a document.

```
package com.emc.tutorial;

import com.documentum.com.DfClientX;
import com.documentum.com.IDfClientX;
import com.documentum.fc.client.IDfDocument;
import com.documentum.fc.client.IDfSession;
import com.documentum.fc.client.IDfSessionManager;
import com.documentum.fc.client.IDfSysObject;
import com.documentum.fc.client.IDfVirtualDocument;
import com.documentum.fc.client.IDfVirtualDocumentNode;
import com.documentum.fc.common.DfId;
import com.documentum.fc.common.IDfId;
import com.documentum.operations.IDfDeleteNode;
import com.documentum.operations.IDfDeleteOperation;

public class TutorialDelete
{
    public TutorialDelete()
    {
    }

    public String deleteExample(
        IDfSessionManager sessionManager,
```

```

        String repositoryName,
        String docId,
        Boolean currentVersionOnly
    )
    {
        IDfSession mySession = null;

        try
        {
            mySession = sessionManager.getSession(repositoryName);

            // Get the object ID based on the object ID string.
            IDfId idObj = mySession.getIdByQualification(
                "dm_sysobject where r_object_id='" + docId + "'");

            // Instantiate an object from the ID.
            IDfSysObject sysObj = (IDfSysObject) mySession.getObject(idObj);

            // Create a new client instance.
            IDfClientX clientx = new DfClientX();

            // Use the factory method to create an IDfDeleteOperation instance.
            IDfDeleteOperation delo = clientx.getDeleteOperation();

            // Create a document object that represents the document being copied.
            IDfDocument doc =
                (IDfDocument) mySession.getObject(new DfId(docId));

            // Set the deletion policy. You must do this prior to adding nodes to the
            // Delete operation.
            if (currentVersionOnly)
            {
                // Default value is SELECTED_VERSIONS
                // delo.setVersionDeletionPolicy(
                //     IDfDeleteOperation.SELECTED_VERSIONS
                // );
            }
            else
            {
                delo.setVersionDeletionPolicy(
                    IDfDeleteOperation.ALL_VERSIONS
                );
            }

            // Create a delete node using the factory method.
            IDfDeleteNode node = (IDfDeleteNode) delo.add(doc);

            if (node == null)
                return "Node is null.";

            // Execute the delete operation and return results.
            if (delo.execute())
            {
                return "Delete operation succeeded.";
            }
            else
            {
                return "Delete operation failed";
            }
        }
    }

    // Handle any exceptions.
    catch (Exception ex)
    {
        ex.printStackTrace();
    }
}

```

```
        return "Exception has been thrown: " + ex;
    }

    // Always, always, release the session in the "finally" clause.
    finally
    {
        sessionManager.release(mySession);
    }
}
```

Special considerations for delete operations

Follow the steps in [Steps for manipulating documents, page 87](#). If the operation entails deleting a checked out document, DFC leaves the document unmodified and reports an error.

Predictive caching

Predictive caching can help you to improve the user experience by sending system objects to Branch Office Caching Services servers before they are requested by users. For example, a company-wide report could be sent to all repository caches when it is added to the local repository rather than waiting for a user request on each server. Another use for this capability would be to cache an object in response to an advance in a workflow procedure, making the document readily available for the next user in the flow.

Example 6-18. Predictive caching operation for a single document

```
void preCache (IDfClientX clientx, IDfDocument doc, IDfList networkLocationIds) {
    IDfPredictiveCachingOperation pco =
        clientx.getPredictiveCachingOperation();
    // Add the document and cast the node to the appropriate type
    IDfPredictiveCachingNode node =
        (IDfPredictiveCachingNode)pco.add( doc );
    if( node == null ) { /* handle errors */ }

    // Set properties on the node
    node.setTimeToLive(IDfPredictiveCachingOperation.DAY);
    node.setNetworkLocationIds(networkLocationIds);
    node.setMinimumContentSize(1000);

    // Execute the operation
    if( !pco.execute() ) { /* handle errors */ }
}
```

Special considerations for predictive caching operations

Follow the steps in [Steps for manipulating documents, page 87](#).

`setTimeToLive` sets a time limit, in milliseconds, for BOCS. BOCS will attempt to pre-cache the content until the specified delay after successful execution of the operation.

`setNetworkLocationIds` sets the list of the network location identifiers to be used for content pre-caching. All BOCS servers for the specified network locations will attempt to pre-cache the content.

`setMinimumContentSize` is used to ensure that documents that are cached will provide a performance improvement. Smaller documents are transferred quickly enough that there is no detectable improvement in performance. Use this method to set the smallest content size, in bytes, that will be cached. Documents smaller than the minimum size will be skipped.

Validating an XML document against a DTD or schema

DFC uses a modified version of the Xerces XML parser, which it includes in dfc.jar. See the DFC release notes for details about the specific version and the Documentum modifications.

The execute method of an IDfValidationOperation object runs the parser in validation mode to determine whether or not documents are well formed and conform to the DTD or schema. If you do not specify it explicitly, DFC determines the XML application automatically and looks in the application's folder in the repository for the DTD or schema.

If any argument or the DTD or schema is in the repository, the execute method makes a temporary copy on the file system and performs the validation there.

If the parser detects errors, the operation's execute method returns a value of false, and you can use the operation's getErrors method to obtain the error information that the parser returns.

Example 6-19. Validating an XML document

Validate an XML document, using C:\Temp for a working directory.

```
void validateXMLDoc( IDfClientX clientx, // Factory for operations
    IDfDocument doc )                // Document to validate
throws DfException
{
    // Obtain validation operation
    IDfValidationOperation validate = clientx.getValidationOperation();
    validate.setDestinationDirectory( "C:/Temp" );

    // Convert the document to a node tree
    IDfVirtualDocument vDoc = doc.asVirtualDocument( "CURRENT", false );

    // Add the document to the operation
    IDfValidationNode node = (IDfValidationNode)validate.add( vDoc );
    if( node == null ) { /* handle errors */ }

    // Execute the operation
    if( !validate.execute() ) { /* handle errors */ }
}
```

Special considerations for validation operations

Follow the steps in [Steps for manipulating documents, page 87](#).

You can use the operation's setDestinationDirectory method to specify the file system directory to which the operation exports the XML files and DTDs or schemas that it passes to the parser.

Performing an XSL transformation of an XML document

The `execute` method of the `IDfXMLTransformOperation` interface uses the Xalan transformation engine and the specified XSLT stylesheet to transform an XML file or document. It places the output into a new document or attaches it to the original document as a rendition. It can also output an object of type `IDfFile` or any `java.io.Writer` or `java.io.OutputStream` stream.

Note: Refer to the DFC release notes for details about the specific version of the Xalan transformation engine that DFC requires.

Example 6-20. Transform to an HTML file

Transform the file `C:\Newsletter.xml` into an HTML file `C:\Newsletter.htm`, using an XSLT stylesheet from the repository.

```
void transformXML2HTMLUsingStylesheetObject(
    IDfClientX clientx,           // Factory for operations
    IDfSession session,          // Repository session (required)
    IDfDocument docStylesheet )  // XSL stylesheet in repository
throws DfException, IOException
{
    // Obtain transformation operation
    IDfXMLTransformOperation opTran = clientx.getXMLTransformOperation();

    // Set operation properties
    opTran.setSession( session );
    opTran.setTransformation( docStylesheet );
    FileOutputStream out = new FileOutputStream( "C:\\Newsletter.htm" );
    opTran.setDestination( out );

    // Add the XML file to the operation
    IDfXMLTransformNode node = (IDfXMLTransformNode)
        opTran.add( "C:\\Newsletter.xml" );
    if( node == null ) { /* handle errors */ }

    // Set node properties
    node.setOutputFormat( "html" );

    // Execute the operation
    if( !opTran.execute() ) { /* handle errors */ }
}
```

Example 6-21. Transform to an HTML file, import result into repository

Transform the file C:\Newsletter.xml into a new HTML document, using the XSLT stylesheet C:\Stylesheet.xsl, and import it into the repository.

```
void transformXML2HTMLUsingStylesheetFile(
    IDfClientX clientx,           // Factory for operations
    IDfSession session,          // Repository session (required)
    IDfId idDestFolder )         // Destination folder
throws DfException
{
    // Obtain transformation operation
    IDfXMLTransformOperation opTran = clientx.getXMLTransformOperation();

    /// Set transformation operation properties
    opTran.setSession( session );
    opTran.setTransformation( "C:\\Stylesheet.xsl" );

    // Obtain import operation
    IDfImportOperation opImp = clientx.getImportOperation();

    // Set import operation properties
    opImp.setSession( session );
    opImp.setDestinationFolderId( idDestFolder );

    // Specify the import operation as the destination of the
    // transformation operation. In effect, this adds the output
    // of the transformation operation to the import operation,
    // but it does not explicitly create an import node
    opTran.setDestination( opImp );

    // Add the XML file to the transform operation
    IDfXMLTransformNode nodeTran =
        (IDfXMLTransformNode)opTran.add( "C:\\Newsletter.xml" );
    if( nodeTran == null ) { /* handle errors */ }

    // Specify the output format
    // (NOTE: on the transformation node. There is no import node)
    nodeTran.setOutputFormat( "html" );

    // Execute the operation
    if( !opTran.execute() ) { /* handle errors */ }
}
```


Example 6-22. Transform an XML document into an HTML rendition

Transform an XML document into an HTML rendition, using an XSLT stylesheet from the repository.

```
void transformXML2HTMLRendition(
    IDfClientX clientx,          // Factory for operations
    IDfSession session,         // Repository session (required)
    IDfDocument docXml,         // Root of the XML document
    IDfDocument docStylesheet ) // XSL stylesheet in repository
throws DfException
{
    // Obtain transformation operation
    IDfXMLTransformOperation opTran = clientx.getXMLTransformOperation();

    opTran.setSession( session );
    opTran.setTransformation( docStylesheet );

    // Add XML document to the transformation operation
    IDfXMLTransformNode node = (IDfXMLTransformNode) opTran.add( docXml );
    if( node == null ) { /* handle errors */ }

    //Set HTML file for the rendition

    // Set format for the rendition
    node.setOutputFormat( "html" );

    // Execute the operation
    if( !opTran.execute() ) { /* handle errors */ }
}
```

DFC creates a rendition because the output format differs from the input format and you did not call `optran.setDestination` to specify an output directory.

Special considerations for XML transform operations

Follow the steps in [Steps for manipulating documents, page 87](#).

You must use the operations `setSession` method to specify a repository session. This operation requires a session, even if all of the files it operates on are on the file system.

The `add` method of an `IDfXMLTransformOperation` object accepts Java types as well as Documentum types. It allows you to specify the file to transform as an object of any of the following types:

- `IDfDocument`
- `IDfFile`
- `String` (for example `C:/PhoneInfo.xml`)
- `DOM` (`org.w3c.dom.Document`)
- `java.io.Reader`
- `URL`

Handling document manipulation errors

This section describes the ways that DFC reports errors that arise in the course of populating or executing operations.

The add Method Cannot Create a Node

The add method of any operation returns a null node if it cannot successfully add the document, file or folder that you pass it as an argument. Test for a null to detect and handle this failure. DFC does not report the reason for returning a null.

The execute Method Encounters Errors

The execute method of an operation throws DfException only in the case of a fatal error. Otherwise it creates a list of errors, which you can examine when the execute method returns or, if you use an operation monitor, as they occur.

Examining Errors After Execution

After you execute an operation, you can use its getErrors method to retrieve an IDfList object containing the errors. You must cast each to IDfOperationError to read its error message.

After detecting that the operation's execute method has returned errors, you can use the operation's abort method to undo as much of the operation as it can. You cannot undo XML validation or transform operations, nor can you restore deleted objects.

Example 6-23. Generate an Operation Exception

Generate a DfException to report the errors that occurred in the course of executing an operation.

```
public DfException generateOperationException(
    IDfOperation operation,
    IDfSession session,
    String msg )
throws DfException
{
    String message = msg;
    DfException e;
    try {

        // Initialize variables
        String strNodeName = "";
        String strNodeId = "";
        String strSucceeded = "";
        IDfId idNodesObj = null;
        IDfOperationError error = null;
        IDfOperationNode node = null;

        // Get the list of errors
```

```

IDfList errorList = operation.getErrors();

// Iterate through errors and build the error messages
for( int i = 0; i < errorList.getCount(); ++i ) {

    // Get next error
    error = (IDfOperationError)errorList.get( i );

    // Get the node at which the error happened
    node = error.getNode();

    // Use method described in another example
    idNodesObj = this.getObjectId( session, node );
    if( idNodesObj <> null ) {
        strNodeId = idNodesObj.getId();
        strNodeName = session.apiGet( "get", strNodeId + ",object_name" );
        message += "Node: [" + strNodeId + "], " + strNodeName + ", "
            + error.getMessage() + ", " + error.getException().toString();
    } // end for
} // end try
catch( Exception err )
{ message += err.toString(); }
finally {
    // Create a DfException to report the errors
    e = new DfException();
    e.setMessage( message );
} // end finally
return e;
}

```

Example 6-24. Obtain the Object ID of an Operation Node

Obtain the IDfId for the operation node.

```

IDfId getObjectId( IDfSession session, IDfOperationNode node ) {
try {
    return
    node instanceof IDfImportNode ? ((IDfImportNode)node).getObjectId()
    : node instanceof IDfExportNode ? ((IDfExportNode)node).getObjectId()
    : node instanceof IDfCheckoutNode ?
        ((IDfCheckoutNode)node).getObjectId()
    : node instanceof IDfCheckinNode ?
        ((IDfCheckinNode)node).getObjectId()
    : node instanceof IDfCancelCheckoutNode ?
        ((IDfCancelCheckoutNode)node).getObjectId()
    : node instanceof IDfDeleteNode ? ((IDfDeleteNode)node).getObjectId()
    : node instanceof IDfCopyNode ? ((IDfCopyNode)node).getObjectId()
    : node instanceof IDfMoveNode ? ((IDfMoveNode)node).getObjectId()
    : null;
} catch( Exception e ) { return null; }
}

```

Using an Operation Monitor to Examine Errors

You can monitor an operation for progress and errors. Create a class that implements the IDfOperationMonitor interface and register it by calling the setOperationMonitor method of IDfOperation. The operation periodically notifies the operation monitor of its progress or of errors that it encounters.

During execution, DFC calls the methods of the installed operation monitor to report progress or errors. You can display this information to an end user. In each case DFC expects a response that tells it whether or not to continue. You can make this decision in the program or ask an end user to decide.

Your operation monitor class must implement the following methods:

- `progressReport`

DFC supplies the percentage of completion of the operation and of its current step. DFC expects a response that tells it whether to continue or to abort the operation.

- `reportError`

DFC passes an object of type `IDfOperationError` representing the error it has encountered. It expects a response that tells it whether to continue or to abort the operation.

- `getYesNoAnswer`

This is the same as `reportError`, except that DFC gives you more choices. DFC passes an object of type `IDfOperationError` representing the error it has encountered. It expects a response of yes, no, or abort.

The Javadocs explain these methods and arguments in greater detail.

Operations and transactions

Operations do not use session transactions (see [Transactions, page 35](#)), because operations

- Support distributed operations involving multiple repositories.
- May potentially process vast numbers of objects.
- Manage non-database resources such as the system registry and the local file system.

You can undo most operations by calling an operation's `abort` method. The `abort` method is specific to each operation, but generally undoes repository actions and cleans up registry entries and local content files. Some operations (for example, `delete`) cannot be undone.

If you know an operation only contains objects from a single repository, and the number of objects being processed is small enough to ensure sufficient database resources, you can wrap the operation execution in a session transaction.

You can also include operations in session manager transactions. Session manager transactions can include operations on objects in different repositories, but you must still pay attention to database resources. Session manager transactions are not completely atomic, because they do not use a two-phase commit. For information about what session transactions can and cannot do, refer to [Transactions, page 35](#).

Using the Business Object Framework (BOF)

This chapter introduces the Business Object Framework (BOF). It contains the following major sections:

- [Overview of BOF, page 125](#)
- [BOF infrastructure, page 126](#)
- [Service-based Business Objects \(SBOs\), page 128](#)
- [Type-based Business Objects \(TBOs\), page 135](#)
- [Calling TBOs and SBOs, page 149](#)
- [Calling DFS, JAXB API, or other components from a BOF module, page 154](#)
- [Aspects, page 155](#)

Overview of BOF

BOF's main goals are to centralize and standardize the process of customizing Documentum functionality. BOF centralizes business logic within the framework. Using BOF, you can develop business logic that

- Always executes, regardless of the client program
- Can extend the implementation of core Documentum functionality
- Runs well in concert with an application server environment

In order to achieve this, the framework leaves customization of the user interface to the clients. BOF customizations embody business logic and are independent of considerations of presentation or format.

If you develop BOF customizations and want to access them from the .NET platform, you must take additional steps. We do not provide tools to assist you in this. You can, however, expose some custom functionality as web services. You can access web services from a variety of platforms (in particular, .NET). The *Web Services Framework Development Guide* provides information about deploying and using the web services.

BOF infrastructure

This section describes the infrastructure that supports the Business Object Framework.

Modules and registries

To understand BOF, first look at how it stores customizations in a repository. A *module* is a unit of executable business logic and its supporting material (for example, documentation, third party software, and so forth).

DFC uses a special type of repository folder (`dmc_module`) to contain a module. The *Content Server Object Reference Manual* describes the attributes of this type. The attributes identify the module type, its implementation class, the interfaces it implements, and the modules it depends on. Other attributes provide version information, a description of the module's functionality, and the developer's contact information.

Every repository has a System cabinet, which contains a top level folder called Modules. The folders in the Modules directory store the JAR files for the interface and implementation classes. Under Modules are the following folders, corresponding to the out-of-the-box module types:

- `/System/Modules/SBO`
Contains service based objects (SBOs). An SBO is a module whose executable business logic is Java code that extends `DfService`. Refer to [Service-based Business Objects \(SBOs\)](#), page 128 for more information about SBOs.
- `/System/Modules/TBO`
Contains type based objects (TBOs), that is, customizations of specific repository types. A TBO is a module in which the executable Java code extends a DFC repository type (normally, `DfDocument`) and implements the `IDfBusinessObject` interface. Refer to [Type-based Business Objects \(TBOs\)](#), page 135 for more information about TBOs.
- `/System/Modules/Aspect`
Contains aspects, a type of module used to apply behaviors and properties to system objects dynamically.

You can create other subfolders of `/System/Modules` to represent other types of module. For example, if the repository uses Java-based evaluation of validation expressions (see [Validation Expressions in Java](#), page 185), the associated modules appear under `/System/Modules/Validation`.

The bottom level folders under this hierarchy (except for aspects) are of type `dmc_module`. Each contains an individual module.

A module that is in other respects like an SBO but does not implement the `IDfService` interface is called a *simple module*. You can use simple modules to implement repository methods, such as those associated with workflows and document lifecycles. The implementation class of a simple module should implement the marker interface `IDfModule`. Use the `newModule` method of `IDfClient` to access a simple module.

The hierarchy of folders under `/System/Modules/` is the repository's module registry, or simply its *registry*.

Note: Earlier versions of DFC maintained a registry of TBOs and SBOs on each client machine. That registry is called the Documentum business object registry (DBOR). The DBOR form of registry is deprecated, but for now you can still use it, even in systems that contain repository based registries. Where both are present, DFC gives preference to the repository based registry.

Dynamic delivery mechanism

BOF delivers the implementation classes of TBOs, SBOs, and other modules dynamically from repository based registries to client machines. A TBO, an aspect, or a simple module is specific to its repository. An SBO is not. BOF can deliver SBO implementation classes to client machines from a single repository. [Global registry, page 127](#) explains how DFC does this.

Delivering implementation classes dynamically from a repository means that you don't need to register those classes on client machines. It also means that all client machines use the same version of the implementation class. You deploy the class to one place, and DFC does the rest.

The delivery mechanism supports *hot deployment*, that is, deployment of new implementations without stopping the application server. This means that applications can pick up changes immediately and automatically. You deploy the module to the global registry, and all clients quickly become aware of the change and start using the new version. DFC works simultaneously with the new version and existing instantiations of the old version until the old version is completely phased out.

The delivery mechanism relies on local caching of modules on client machines (where the term client machine often means the machine running an application server and, usually, WDK). DFC does not load TBOs, aspects, or simple modules into the cache until an application tries to use them. Once DFC has downloaded a module, it only reloads parts of the module that change.

DFC checks for updates to the modules in its local cache whenever an application tries to use one or after an interval specified in seconds in `dfc.bof.cache.currency_check_interval` in the `dfc.properties` file. The default value is 60 seconds.

If DFC tries to access a module registry and fails, it tries again after a specified interval. The interval, in seconds, is the value of `dfc.bof.registry.connect.attempt.interval` in the `dfc.properties` file. The default value is 60 seconds.

DFC maintains its module cache on the file system of the client machine. You can specify the location by setting `dfc.cache.dir` in the `dfc.properties` file. The default value is the cache subdirectory of the directory specified in `dfc.data.dir`. All applications that use the given DFC installation share the cache. Do not share the cache among more than one DFC installation. If you have more than one DFC installation, ensure that the value of `dfc.cache.dir` in the `dfc.properties` file does not point to the same directory.

Global registry

DFC delivers SBOs from a central repository. That repository's registry is called the *global registry*.

Global registry user

The global registry user, who has the user name of `dm_bof_registry`, is the repository user whose account is used by DFC clients to connect to the repository to access required service-based objects or network locations stored in the global registry. This user has Read access to objects in the `/System/Modules`, `/System/BocsConfig`, `/dm_bof_registry`, and `/System/NetworkLocations` only, and no other objects.

Accessing the global registry

The identity of the global registry is a property of the DFC installation. Different DFC installations can use different global registries, but a single DFC installation can have only one global registry.

In addition to efficiency, local caching provides backup if the global registry is unavailable. Only when the client application tries to instantiate an SBO for the first time, the corresponding implementation and interface jar files are downloaded from the global registry to the local cache. Thereafter, DFC reloads only those parts of the module that change. For that purpose, DFC checks for updates related to the cached modules whenever an application tries to use them or after the interval specified in seconds for the DFC property `dfc.bof.cache.currency_check_interval` in the `dfc.properties` file. Its default value is 60 seconds.

The `dfc.properties` file contains properties that relate to accessing the global registry. Refer to [BOF and global registry settings, page 17](#) for information about using these properties.

Service-based Business Objects (SBOs)

This section contains the following main sections:

- [SBO introduction, page 128](#)
- [SBO architecture, page 129](#)
- [Implementing SBOs, page 129](#)
- [Calling SBOs, page 149](#)
- [SBO Error Handling, page 134](#)
- [SBO Best Practices, page 134](#)

SBO introduction

A *service based object* (SBO) is a type of module designed to enable developers to access Documentum functionality by writing small amounts of relevant code. The underlying framework handles most of the details of connecting to Documentum repositories. SBOs are similar to session beans in an Enterprise JavaBean (EJB) environment.

SBOs can operate on multiple object types, retrieve objects unrelated to Documentum objects (for example, external email messages), and perform processing. You can use SBOs to implement

functionality that applies to more than one repository type. For example, a Documentum Inbox object is an SBO. It retrieves items from a user's inbox and performs operations like removing and forwarding items.

You can use SBOs to implement utility functions to be called by multiple TBOs. A TBO has the references it needs to instantiate an SBO.

You can implement an SBO so that an application server component can call the SBO, and the SBO can obtain and release repository sessions dynamically as needed.

SBOs are the basis for the web services framework.

SBO architecture

An SBO associates an interface with an implementation class. Each folder under /System/Modules/SBO corresponds to an SBO. The name of the folder is the name of the SBO, which by convention is the name of the interface.

SBOs are not associated with a repository type, nor are they specific to the repository in which they reside. As a result, each DFC installation uses a global registry (see [Global registry, page 127](#)). The dfc.properties file contains the information necessary to enable DFC to fetch SBO implementation classes from the global registry.

You instantiate SBOs with the newService method of IDfClient, which requires you to pass it a session manager. The newService method searches the registry for the SBO and instantiates the associated Java class. Using its session manager, an SBO can access objects from more than one repository.

You can easily design an SBO to be stateless, except for the reference to its session manager.

Note: DFC does not enforce a naming convention for SBOs, but we recommend that you follow the naming convention explained in [Follow the Naming Convention, page 134](#).

Implementing SBOs

This section explains how to implement an SBO.

An SBO is defined by its interface. Callers cannot instantiate an SBO's implementation class directly. The interface should refer only to the specific functionality that the SBO provides. A separate interface, IDfService, provides access to functionality common to all SBOs. The SBO's implementation class, however, should not extend IDfService. Instead, the SBO's implementation class must extend DfService, which implements IDfService. Extending DfService ensures that the SBO provides several methods for revealing information about itself to DFC and to applications that use the SBO.

To create an SBO, first specify its defining interface. Then create an implementation class that implements the defining interface and extends DfService. DfService is an abstract class that defines common methods for SBOs.

Override the following abstract methods of DfService to provide information about your SBO:

- getVersion returns the current version of the service as a string.

The version is a string and must consist of an integer followed by up to three instances of dot integers (for example, 1.0 or 2.1.1.36). The version number is used to determine installation options.

- `getVendorString` returns the vendor's copyright statement (for example, "Copyright 1994-2005 EMC Corporation. All rights reserved.") as a string.
- `isCompatible` checks whether the class is compatible with a specified service version

This allows you to upgrade service implementations without breaking existing code. Java does not support multiple versions of interfaces.

- `supportsFeature` checks whether the string passed as an argument matches a feature that the SBO supports.

The `getVersion` and `isCompatible` methods are important tools for managing SBOs in an open environment. The `getVendorString` method provides a convenient way for you to include your copyright information. The `supportsFeature` method can be useful if you develop conventions for naming and describing features.

SBO programming differs little from programming for other environments. The following sections address the principal additional considerations.

Stateful and stateless SBOs

SBOs can maintain state between calls, but they are easier to deploy to multithreaded and other environments if they do not do so. For example, a checkin service needs parameters like `retainLock` and `versionLabels`. A stateful interface for such a service provides `get` and `set` methods for such parameters. A stateless interface makes you pass the state as calling arguments.

Managing Sessions for SBOs

This section presents session manager related considerations for implementing SBOs.

Overview

When implementing an SBO, you normally use the `getSession` and `releaseSession` methods of `DfService` to obtain a DFC session and return it when finished. Once you have a session, use the methods of `IDfSession` and other DFC interfaces to implement the SBO's functionality.

If you need to access the session manager directly, however, you can do so from any method of a service, because the session manager object is a member of the `DfService` class. The `getSessionManager` method returns this object. To request a new session, for example, use the session manager's `newSession` method.

Structuring Methods to Use Sessions

Each SBO method that obtains a repository session must release the session when it is finished accessing the repository. The following example shows how to structure a method to ensure that it releases its session, even if exceptions occur:

```
public void doSomething( String strRepository, . . . ) {
    IDfSession session = getSession ( strRepository );
    try { /* do something */ }
    catch( Exception e ) { /* handle error */ }
    finally { releaseSession( session ); }
}
```

Managing repository names

To obtain a session, an SBO needs a repository name. To provide the repository name, you can design your code in any of the following ways:

- Pass the repository name to every service method.

This allows a stateless operation. Use this approach whenever possible.

- Store the repository name in an instance variable of the SBO, and provide a method to set it (for example, `setRepository (strRepository)`).

This makes the repository available from all of the SBO's methods.

- Extract the repository name from an object ID.

A method that takes an object ID as an argument can extract the repository name from the object ID (use the `getDocbaseNameFromId` method of `IDfClient`).

Maintaining State Beyond the Life of the SBO

The EMC | Documentum architecture enables SBOs to return persistent objects to the calling program. Persistent objects normally maintain their state in the associated session object. But an SBO must release the sessions it uses before returning to the calling program. At any time thereafter, the session manager might disconnect the session, making the state of the returned objects invalid.

The calling program must ensure that the session manager does not disconnect the session until the calling program no longer needs the returned objects.

Another reason for preserving state between SBO calls occurs when a program performs a query or accesses an object. It must obtain a session and apply that session to any subsequent calls requiring authentication and Content Server operations. For application servers, this means maintaining the session information between HTTP requests.

The main means of preserving state information are `setSessionManager` and transactions.

[Maintaining state in a session manager, page 46](#) describes the `setSessionManager` mechanism and its cost in resources. [Using Transactions With SBOs, page 132](#) provides details about using transactions with SBOs.

You can also use the `DfCollectionEx` class to return a collection of typed objects from a service. `DfCollectionEx` locks the session until you call its `close` method.

Obtaining Session Manager State Information

For testing or performance tuning you can examine such session manager state as reference counters, the number of sessions, and repositories currently connected. Use the `getStatistics` method of `IDfSessionManager` to retrieve an `IDfSessionManagerStatistics` object that contains the state information. The statistics object provides a snapshot of the session manager's internal data as of the time you call `getStatistics`. DFC does not update this object if the session manager's state subsequently changes.

The DFC Javadocs describe the available state information.

Using Transactions With SBOs

DFC supports two transaction processing mechanisms: session based and session manager based. [Transactions, page 35](#) describes the differences between the two transaction mechanisms. You cannot use session based transactions within an SBO method. DFC throws an exception if you try to do so.

Use the following guidelines for transactions within an SBO:

- Never begin a transaction if one is already active.

The `isTransactionActive` method returns true if the session manager has a transaction active.

- If the SBO does not begin the transaction, do not use `commitTransaction` or `abortTransaction` within the SBO's methods.

If you need to abort a transaction from within an SBO method, use the session manager's `setTransactionRollbackOnly` method instead, as described in the next paragraph.

When you need the flow of a program to continue when transaction errors occur, use the session manager's `setTransactionRollbackOnly`. Thereafter, DFC silently ignores attempts to commit the transaction. The owner of the transaction does not know that one of its method calls aborted the transaction unless it calls the `getTransactionRollbackOnly` method, which returns true if some part of the program ever called `setTransactionRollbackOnly`. Note that `setTransactionRollbackOnly` does not throw an exception, so the program continues as if the batch process were valid.

The following program illustrates this.

```
void serviceMethodThatRollsBack( String strRepository, IDfId idDoc )
    throws DfNoTransactionAvailableException, DfException {

    IDfSessionManager sMgr = getSessionManager();
    IDfSession = getSession( strRepository );
    if( ! sMgr.isTransactionActive() ) {
        throw new DfNoTransactionAvailableException();
    }

    try {
        IDfPersistentObject obj = session.getObject( idDoc );
        obj.checkout();
        modifyObject( obj );
        obj.save();
    }

    catch( Exception e ) {
        setTransactionRollbackOnly();
        throw new DfException();
    }
}
```

```

    }
}

```

When more than one thread is involved in session manager transactions, calling `beginTransaction` from a second thread causes the session manager to create a new session for the new thread.

The session manager supports transaction handling across multiple services. It does not disconnect or release sessions while transactions are pending.

For example, suppose one service creates folders and a second service stores documents in these folders. To make sure that you remove the folders if the document creation fails, place the two service calls into a transaction. The DFC session transaction is bound to one DFC session, so it is important to use the same DFC session across the two services calls. Each service performs its own atomic operation. At the start of each operation, they request a DFC session and at the end they release this session back to the session pool. The session manager holds on to the session as long as the transaction remains open.

Use the `beginTransaction` method to start a new transaction. Use the `commitTransaction` or `abortTransaction` method to end it. You must call `getSession` after you call `beginTransaction`, or the session object cannot participate in the transaction.

Use the `isTransactionActive` method to ask whether the session manager has a transaction active that you can join. DFC does not allow nested transactions.

The transaction mechanism handles the following issues:

- With multiple threads, transaction handling operates on the current thread only.
For example, if there is an existing session for one thread, DFC creates a new session for the second thread automatically. This also means that you cannot begin a transaction in one thread and commit it in a second thread.
- The session manager provides a separate session for each thread that calls `beginTransaction`.
For threads that already have a session before the transaction begins, DFC creates a new session.
- When a client starts a transaction using the `beginTransaction` method, the session manager does not allow any other DFC-based transactions to occur.

The following example illustrates a client application calling two services that must be inside a transaction, in which case both calls must succeed, or nothing changes:

```

IDfClient client = DfClientX.getLocalClient();
IDfSessionManager sMgr = client.newSessionManager();

sMgr.setIdentity(repo, loginInfo);

IMyService1 s1 = (IMyService1)
    client.newService(IMyService1.class.getName(), sMgr);

IMyService2 s2 = (IMyService2)
    client.newService(IMyService2.class.getName(), sMgr);

s1.setRepository( strRepository1 );
s2.setRepository( strRepository2 );

sMgr.beginTransaction();

try {
    s1.doRepositoryUpdate();
}

```

```
s2.doRepositoryUpdate();  
sMgr.commitTransaction();  
}  
  
catch (Exception e) {  
    sMgr.abortTransaction();  
}
```

If either of these service methods throws an exception, the program bypasses commit and executes abort.

Each of the `doRepositoryUpdate` methods calls the session manager's `getSession` method.

Note that the two services in the example are updating different repositories. Committing or aborting the managed transaction causes the session manager to commit or abort transactions with each repository.

Session manager transactions involving more than one repository have an inherent weakness that arises from their reliance on the separate transaction mechanisms of the databases underlying the repositories. Refer to [Transactions, page 35](#) for information about what session manager transactions can and cannot do.

SBO Error Handling

The factory method that instantiates SBOs throws a variety of exceptions. For example, it throws `DfServiceInstantiationException` if DFC finds the requested service but is unable to instantiate the specified Java class. This can happen if the Java class is not in the classpath or is an invalid data class. Security for Java classes on an application server can also cause this exception.

SBO Best Practices

This section describes best practices for using SBOs.

Follow the Naming Convention

DFC does not enforce a naming convention for SBOs, but we recommend that you give an SBO the same name as the fully qualified name of the interface it implements. For example, if you produce an SBO that implements an interface called `IContentValidator`, you might name it `com.myFirm.services.IContentValidator`. If you do this, the call to instantiate an SBO becomes simple. For example, to instantiate an instance of the SBO that implements the `IContentValidator` interface, simply write

```
IContentValidator cv = (IContentValidator)client.newService(  
    IContentValidator.class.getName(), sMgr);
```

The only constraint DFC imposes on SBO names is that names must be unique within a registry.

Don't Reuse SBOs

Instantiate a new SBO each time you need one, rather than reusing one. Refer to [Calling SBOs, page 149](#) for details.

Make SBOs Stateless

Make SBOs as close to stateless as possible. Refer to [Stateful and stateless SBOs, page 130](#) for details.

Rely on DFC to Cache Repository Data

DFC caches persistent repository data. There is no convenient way to keep a private cache synchronized with the DFC cache, so rely on the DFC cache, rather than implementing a separate cache as part of your service's implementation.

Type-based Business Objects (TBOs)

Use of Type-based Business Objects

Type-based Business Objects are used for modifying and extending the behavior of persistent repository object types, including core DFC types (such as documents and users) and other TBOs. TBOs extend DFC object types that inherit from `IDfPersistentObject`, and which map to persistent repository objects, such as `dm_document` or `dm_user`. The TBOs themselves map to custom repository object types.

For example, suppose you want to add or modify behaviors exhibited by a custom repository type derived from `dm_document`, which we will call `mycompany_sop`. Typically, you might want to extend behavior that occurs whenever a document of type `mycompany_sop` is checked in, by starting a workflow, validating or setting XML attributes, applying a lifecycle, or creating a rendition. Or you may want to add new behaviors to the object type that are called from a client application or from an SBO.

A TBO provides a client-independent, component-based means of implementing this type of business logic, using either or both of the following techniques:

- Override the methods of the parent class from which the TBO class is derived. This approach is typically used to add pre- or postprocessing to the parent method that will be invoked during normal operations such as checkin, checkout, or link.
- Add new methods to the TBO class that can be called by an SBO or by a DFC client application.

Creating a TBO

The following sections describe how to create a TBO. Here is a summary of the steps required:

1. Create a custom repository type.
2. Create the TBO interface.
3. Create the TBO class.
4. Implement the methods of IDfBusinessObject.
5. Code your business logic by adding new methods and overriding methods of the parent class.

The following sections provide more detailed instructions for building TBOs. The sample code provided is works from the assumption that the TBO is derived from the DfDocument class, and that its purpose is to extend the behavior of the custom object on checkin and save.

Create a custom repository type

Using Composer, create and configure your custom type. For an example implementation, see [Deploying the SBO and TBO, page 154](#).

Create the TBO interface

Creating an interface for the TBO is generally recommended, but optional if you do not intend to extend the parent class of the TBO by adding new methods. If you only intend to override methods inherited from the parent class, there is no strict need for a TBO interface, but use of such an interface may make your code more self-documenting, and make it easier to add new methods to the TBO should you have a need to add them in the future.

The design of the TBO interface should be determined by which methods you want to expose to client applications and SBOs. If your TBO needs to expose new public methods, declare their signatures in the TBO interface. Two other questions to consider are (1) whether to extend the interface of the TBO superclass (e.g. IDfDocument), and (2) whether to extend IDfBusinessObject.

While the TBO *class* will need to extend the base DFC class (for example DfDocument), you may want to make the TBO *interface* more restricted by redeclaring only those methods of the base class that your business logic requires you to expose to clients. This avoids polluting the custom interface with unnecessary methods from higher-level DFC interfaces. On the other hand, if your TBO needs to expose a large number of methods from the base DFC class, it may be more natural to have the TBO interface extend the interface of the superclass. This is a matter of design preference.

Although not a functional requirement of the BOF framework, it is generally accepted practice for the TBO interface to extend IDfBusinessObject, merging into the TBO's contract its concerns as a business object with its concerns as a persistent object subtype. This enables you to get an instance of the TBO class and call IDfBusinessObject methods without the complication of a cast to IDfBusinessObject:

```
IMySop mySop = (IMySop) session.getObject(id);
if (mySop.supportsFeature("some_feature"))
{
    mySop.mySopMethod();
}
```



```
}
```

The following sample TBO interface extends IDfBusinessObject and redeclares a few required methods of the TBO superclass (rather than extending the interface of the superclass):

```
import com.documentum.fc.common.DfException;
import com.documentum.fc.common.IDfId;
import com.documentum.fc.client.IDfBusinessObject;

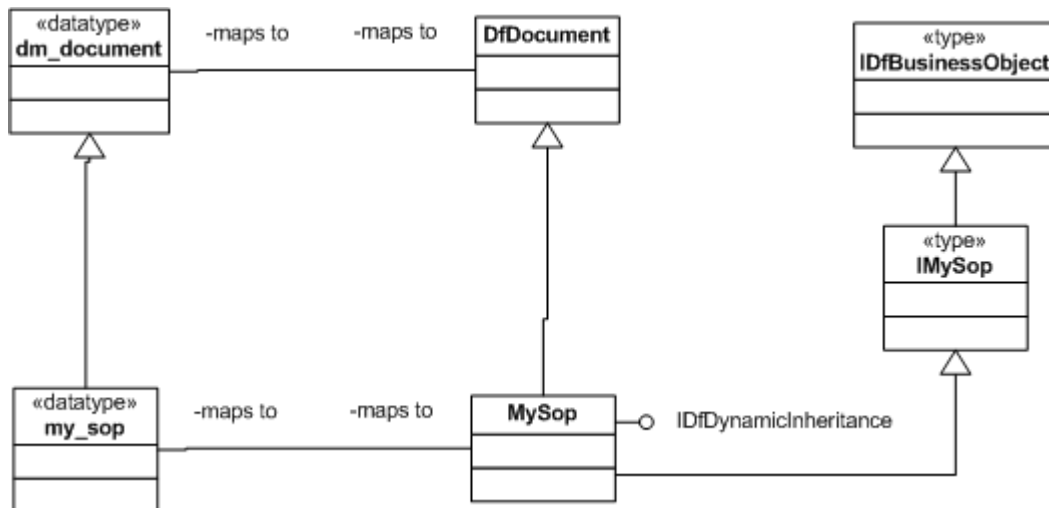
/**
 * TBO interface intended to override checkout and save behaviors of
 * IDfDocument. IDfDocument is not extended because only a few of its
 * methods are required IDfBusinessObject is extended to permit calling
 * its methods without casting the TBO instance to IDfBusinessObject
 */
public interface IMySop extends IDfBusinessObject
{
    public boolean isCheckedOut() throws DfException;
    public void checkout() throws DfException;
    public IDfId checkin(boolean fRetainLock, String versionLabels)
        throws DfException;
    public void save() throws DfException;
}
```

Define the TBO implementation class

The main class for your TBO is the class that will be associated with a custom repository object type when deploying the TBO. This class will normally extend the DFC type class associated with the repository type from which your custom repository type is derived. For example, if your custom repository type my_sop extends dm_document, extend the DfDocument class. In this case the TBO class must implement IDfBusinessObject (either directly or by implementing your custom TBO interface that extends IDfBusinessObject) and it must implement IDfDynamicInheritance.

```
public class MySop extends DfDocument implements IMySop,
    IDfDynamicInheritance
```

Figure 4. Basic TBO design

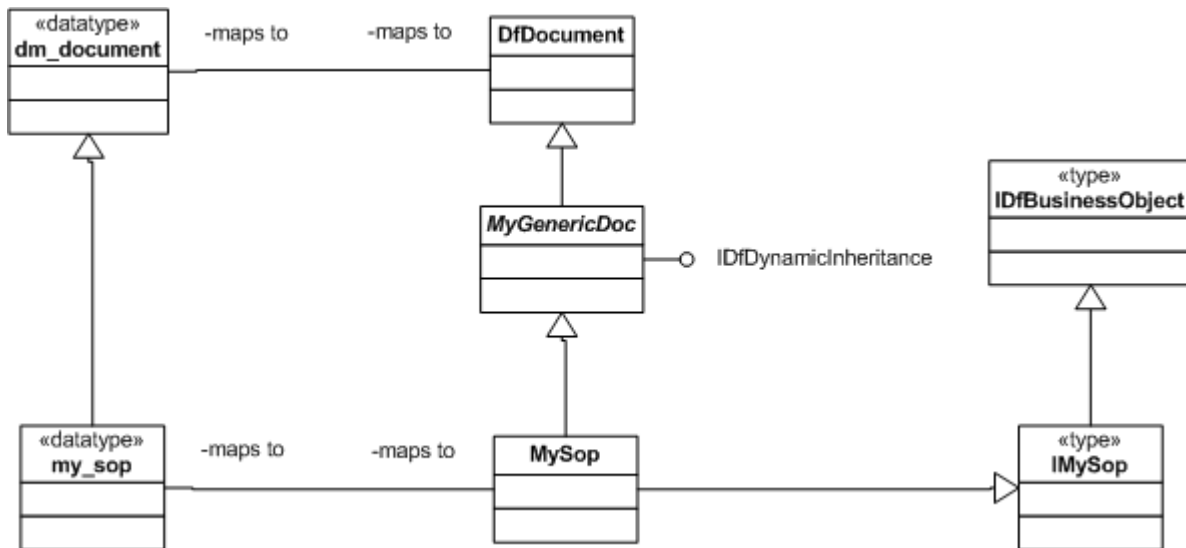


It is also an option to create more hierarchical levels between your main TBO class and the DFC superclass. For example, you may want to place generic methods used in multiple TBOs

in an abstract class. In this case the higher class will extend the DfDocument superclass and implement IDfDynamicInheritance, and the main TBO class would extend the abstract class and implement the TBO interface. This will result in the correct runtime behavior for dynamic inheritance.

```
public abstract class MyGenericDoc extends DfDocument
    implements IDfDynamicInheritance
public class MySop extends MyGenericDoc implements IMySop
```

Figure 5. TBO design with extended intervening class



Note that in this situation you would need to package both MyGenericDoc and MySop into the TBO class jar file, and specify MySop as the main TBO class when deploying the TBO in the repository. For an example of packaging and deploying business objects see [Deploying the SBO and TBO, page 154](#).

Implement methods of IDfBusinessObject

To fulfill its contract as a class of type IDfBusinessObject, the TBO class must implement the following methods:

- getVersion
- getVendorString
- isCompatible
- supportsFeature

The version support features getVersion and isCompatible must have functioning implementations (these are required and used by the Business Object Framework) and it is important to keep the TBO version data up-to-date. Functional implementation of the supportsFeature method is optional: you can provide a dummy implementation that just returns a Boolean value.

For further information see IDfBusinessObject in the Javadocs.

getVersion method

The getVersion method must return a string representing the version of the business object, using the format <major version>.<minor version> (for example 1.10), which can be extended to include as many as four total integers, separated by periods (for example 1.10.2.12).

getVendorString method

The getVendorString method returns a string containing information about the business object vendor, generally a copyright string.

isCompatible method

The isCompatible method takes a String argument in the format <major version>.<minor version> (for example 1.10), which can be extended to include as many as four total integers, separated by periods (for example 1.10.2.12). The isCompatible method, which is intended to be used in conjunction with getVersion, must return true if the TBO is compatible with the version and false if it is not.

supportsFeature method

The supportsFeature method is passed a string representing an application or service feature, and returns true if this feature is supported and false otherwise. Its intention is to allow your application to store lists of features supported by the TBO, perhaps allowing the calling application to switch off features that are not supported.

Support for features is an optional adjunct to mandatory version compatibility support. Features are a convenient way of advertising functionality that avoids imposing complicated version checking on the client. If you choose not to use this method, your TBO can provide a minimal implementation of supportsFeature that just returns a boolean value.

Code the TBO business logic

You can implement business logic in your TBO by adding new methods, or by adding overriding methods of the class that your TBO class extends. When overriding methods, you will most likely want to add custom behavior as pre- or postprocessing before or after a call to super.<methodName>. The following sample shows an override of the IDfSysObject.doCheckin method that writes an entry to the log.

```
protected IDfId doCheckin(boolean fRetainLock,
    String versionLabels,
    String oldCompoundArchValue,
    String oldSpecialAppValue,
    String newCompoundArchValue,
    String newSpecialAppValue,
    Object[] extendedArgs) throws DfException
{
    Date now = new Date();
    DfLogger.warn(this, now + " doCheckin() called", null, null);
```

```
// your preprocessing logic here
return super.doCheckin(fRetainLock,
    versionLabels,
    oldCompoundArchValue,
    oldSpecialAppValue,
    newCompoundArchValue,
    newSpecialAppValue,
    extendedArgs);
// your postprocessing logic here
}
```

Override only methods beginning with do (doSave, doCheckin, doCheckout, and similar). The signatures for these methods are documented in [Signatures of Methods to Override](#), page 145.

Using a TBO from a client application

A TBO is an extension of a core DFC typed object, so instantiating a TBO is no different from instantiating a core DFC typed object. To instantiate a TBO from a client application, use a method of IDfSession that fetches an object, such as getObject, newObject, or getObjectByQualification. The session object used to generate the TBO must be a managed session, that is, a session that was instantiated using a IDfSessionManager.getSession or newSession factory method.

The following test method exercises a TBO by getting a known object of the TBO type from the repository using getObjectByQualification and checking it in.

```
private void testCheckinOverride(String userName,
    String docName,
    String password,
    String docbaseName,
    String typeName) throws Exception
{
    IDfSessionManager sessionManager = null;
    IDfSession docbaseSession = null;
    IMyCompanySop mySop = null;
    try
    {
        // get a managed session
        IDfClient localClient = DfClient.getLocalClient();
        sessionManager = localClient.newSessionManager();
        IDfLoginInfo loginInfo = new DfLoginInfo();
        loginInfo.setUser(userName);
        loginInfo.setPassword(password);
        sessionManager.setIdentity(docbaseName, loginInfo);
        docbaseSession = sessionManager.getSession(docbaseName);

        // get the test document
        // the query string must uniquely identify it
        StringBuffer bufQual = new StringBuffer(32);
        bufQual.append(typeName)
            .append(" where object_name like '")
            .append(docName).append("'");
        mySop = (IMyCompanySop)
            docbaseSession.getObjectByQualification(bufQual.toString());
        if (mySop == null)
        {
            fail("Unable to locate object with name " + docName);
        }

        // check in document to see whether anything gets
        //written to the log
    }
}
```

```

    if (!mySop.isCheckedOut())
    {
        mySop.checkout();
    }
    if (mySop.isCheckedOut())
    {
        mySop.checkin(false, "MOD_TEST_FILE");
    }
}
catch (Throwable e)
{
    fail("Failed with exception " + e);
}
finally
{
    if ((sessionManager != null) && (docbaseSession != null))
    {
        sessionManager.release(docbaseSession);
    }
}
}

```

Using TBOs from SBOs

If you are instantiating a TBO from an SBO, use the `IDfService.getSession` method, which returns a session managed by the session manager associated with the SBO. The following sample SBO method gets and returns a TBO:

```

public IDfDocument getDoc( String strRepository, IDfId idDoc )
{
    IDfSession session = null;
    IDfDocument doc = null;
    try
    {
        // calls IDfService.getSession
        session = getSession ( strRepository );
        doc = (IDfDocument)session.getObject( idDoc );
        // note no call to setSessionManager
        // this is not needed in Documentum 6
    }
    finally
    {
        releaseSession( session );
    }
    return doc;
}

```

Note the absence of any call to `setSessionManager` in the preceding listing. In Documentum 5 a call to `setSessionManager` was required to *disconnect* the object from the session and place it in the state of the session manager. This allowed the SBO to release the session and return the TBO instance without the object becoming stale (that is, disassociated from its session context). In Documentum 6 there is transparent handling of the association of objects and sessions: if you return an object and release its session, DFC will create a new session for the object when it is required. Legacy calls to `setSessionManager` will continue to work, but are no longer required for this purpose.

Getting sessions inside TBOs

You can obtain a reference to the session that was originally used to fetch an object using the `IDfTypedObject.getSession` method. However, when you obtain a session in this way you cannot release it (if you attempt this an exception will be thrown), and you cannot change its repository scope (that is, the name of the repository to which the session maintains a connection). This is because no ownership of the session is implied when you get an existing session using `IDfTypedObject.getSession`. You can only release sessions that were obtained using a factory method of a session manager. This restriction prevents a misuse of sessions that could lead to abstruse bugs.

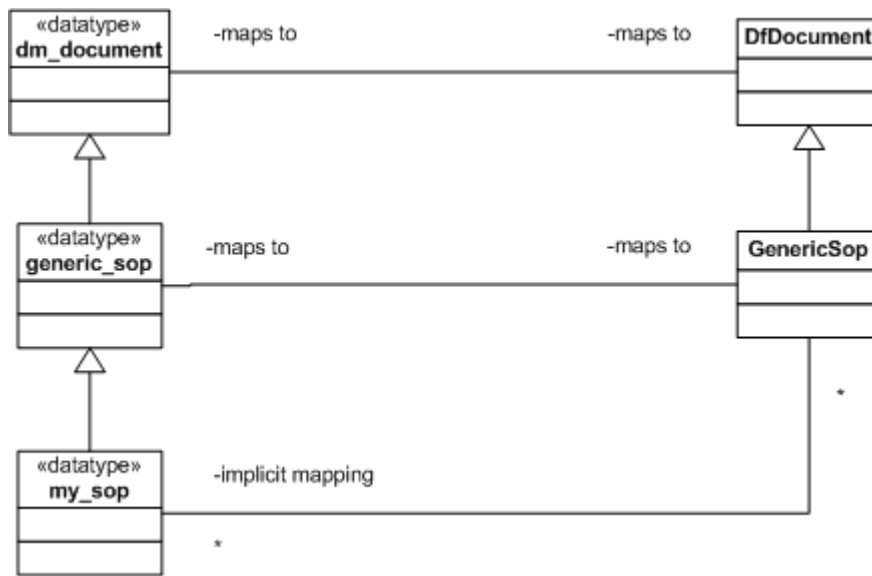
If you need to get a session independent of any session associated with the TBO object, you can use `IDfTypedObject.getSessionManager` to return the session manager associated with the TBO object. You can then get sessions using the factory methods of this session manager (which allows you to use any identities defined in the session manager) and release them (in a finally block) after you have finished using the session.

The `IDfTypedObject.getSession` method returns the session on which the object was originally fetched, which in most situations will be the session maintaining a connection to the repository that holds the object. However, in a distributed environment where DFC is doing work in multiple repositories, the session on which the object was originally fetched and the session to the object's repository may not be the same. In this case you will generally want to obtain the session maintaining a connection to the object's repository. To do this you can use the `IDfTypedObject.getObjectSession` method instead of `getSession`. If you specifically want to get the session that was originally used to fetch the object, you can use the `getOriginalSession` method. The `IDfTypedObject.getSession` method is a synonym for `getOriginalSession`.

Inheritance of TBO methods by repository subtypes without TBOs

If you create a repository object type B that does not have a TBO but which inherits from a repository type A that does have a TBO, object B will inherit the behaviors defined by A's TBO. This means that you do not have to create a TBO for each repository subtype unless you need to extend or override the behaviors associated with the parent type. This inheritance mechanism allows administrators to create repository subtypes that behave in a natural way, inheriting the behaviors of the parent type, without the developer having to write another TBO.

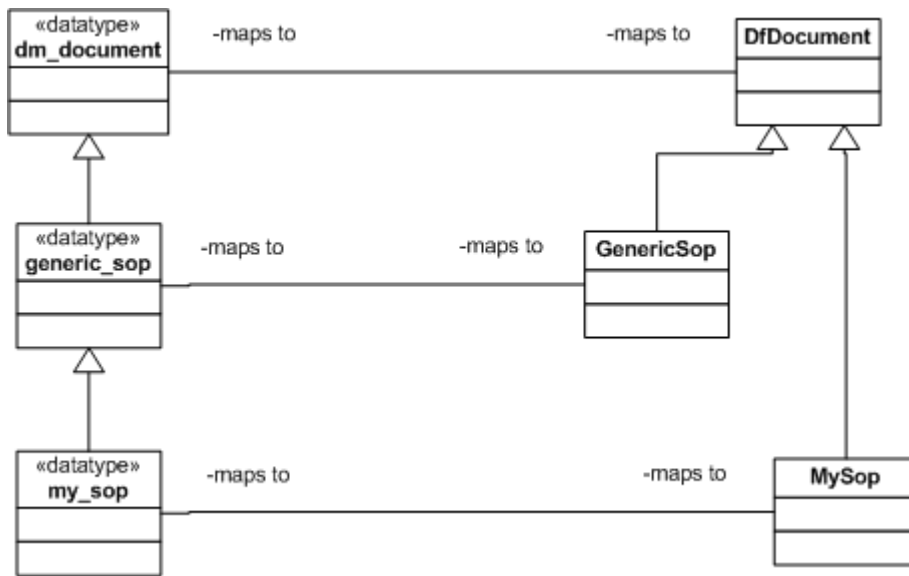
For example, suppose you have a repository type `generic_sop`, which has custom behaviors on `checkin` defined in a class `GenericSop`. If an administrator then created a new type `my_sop`, the new type would inherit the custom `checkin` behaviors of `generic_sop` automatically, without the developer needing to implement and test a new TBO for `my_sop`. Similarly, if a process fetches an object of type `my_sop` using `IDfSession.getObject` or a similar method, the `getObject` method will return an instance of `GenericSop`.

Figure 6. Inheritance by object subtype without associated TBO

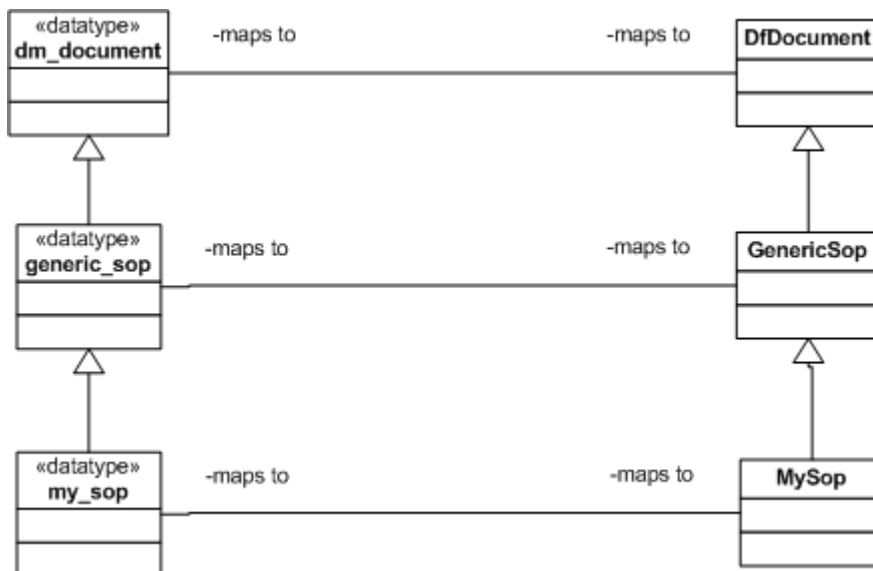
Dynamic inheritance

Dynamic inheritance is a BOF mechanism that modifies the class inheritance of a TBO dynamically at runtime, driven by the hierarchical relationship of associated repository objects. This mechanism enforces consistency between the repository object hierarchy and the associated class hierarchy. It also allows you to design polymorphic TBOs that inherit from different superclasses depending on runtime dynamic resolution of the class hierarchy.

For example, suppose you have the following TBO design, in which repository objects are related hierarchically, but in which the associated TBO classes each inherit from DfDocument:

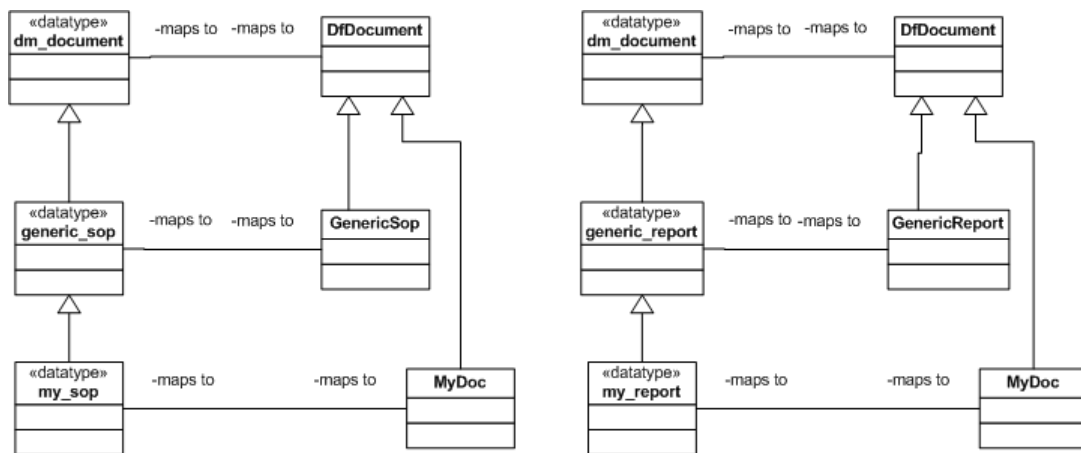
Figure 7. Design-time dynamic inheritance hierarchies

If dynamic inheritance is enabled, at runtime the class hierarchy is resolved dynamically to correspond to the repository object hierarchy, so that the MySop class inherits from GenericSop:

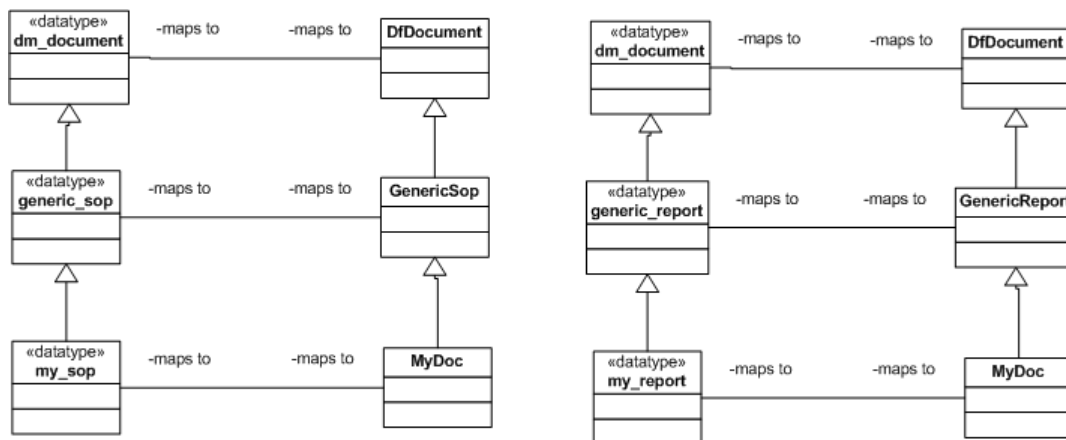
Figure 8. Runtime dynamic inheritance hierarchies

Exploiting dynamic inheritance with TBO reuse

The dynamic inheritance mechanism allows you to design reusable components that exhibit different behaviors at runtime inherited from their dynamically determined superclass. For example, in the following design-time configuration, the MyDoc class is packaged in two TBOs: one in which it is associated with type my_sop, and one in which it is associated with type my_report:

Figure 9. Design-time dynamic inheritance with TBO reuse

At runtime, MyDoc will inherit from GenericSop where it is associated with the my_sop repository object type, and from GenericReport where it is associated with the my_report repository object type.

Figure 10. Runtime dynamic inheritance with TBO reuse

Signatures of Methods to Override

This section contains a list of methods of DfSysObject, DfPersistentObject, and DfTypedObject that are designed to be overridden by implementors of TBOs. All of these methods are protected. You should make your overrides of these methods protected as well.

Table 2. Methods to override when implementing TBOs

Methods of DfSysObject
IDfId doAddESignature (String userName, String password, String signatureJustification, String formatToSign, String hashAlgorithm, String preSignatureHash, String signatureMethodName, String applicationProperties, String passThroughArgument1, String passThroughArgument2, Object[] extendedArgs) throws DfException

IDfld doAddReference (IDfld folderId, String bindingCondition, String bindingLabel, Object[] extendedArgs) throws DfException
void doAddRendition (String fileName, String formatName, int pageNumber, String pageModifier, String storageName, boolean atomic, boolean keep, boolean batch, String otherFileName, Object[] extendedArgs) throws DfException
void doAppendFile (String fileName, String otherFileName, Object[] extendedArgs) throws DfException
IDfCollection doAssemble (IDfld virtualDocumentId, int interruptFrequency, String qualification, String nodesortList, Object[] extendedArgs) throws DfException
IDfVirtualDocument doAsVirtualDocument (String lateBindingValue, boolean followRootAssembly, Object[] extendedArgs) throws DfException
void doAttachPolicy (IDfld policyId, String state, String scope, Object[] extendedArgs) throws DfException
void doBindFile (int pageNumber, IDfld srcId, int srcPageNumber, Object[] extendedArgs) throws DfException
IDfld doBranch (String versionLabel, Object[] extendedArgs) throws DfException
void doCancelScheduledDemote (IDfTime scheduleDate, Object[] extendedArgs) throws DfException
void doCancelScheduledPromote (IDfTime scheduleDate, Object[] extendedArgs) throws DfException
void doCancelScheduledResume (IDfTime schedule, Object[] extendedArgs) throws DfException
void doCancelScheduledSuspend (IDfTime scheduleDate, Object[] extendedArgs) throws DfException
IDfld doCheckin (boolean fRetainLock, String versionLabels, String oldCompoundArchValue, String oldSpecialAppValue, String newCompoundArchValue, String newSpecialAppValue, Object[] extendedArgs)
IDfld doCheckout (String versionLabel, String compoundArchValue, String specialAppValue, Object[] extendedArgs)
void doDemote (String state, boolean toBase, Object[] extendedArgs) throws DfException
void doDestroyAllVersions (Object[] extendedArgs) throws DfException
void doDetachPolicy (Object[] extendedArgs) throws DfException
void doDisassemble (Object[] extendedArgs) throws DfException
boolean doFetch (String currencyCheckValue, boolean usePersistentCache, boolean useSharedCache, Object[] extendedArgs) throws DfException
void doFreeze (boolean freezeComponents, Object[] extendedArgs) throws DfException
void doInsertFile (String fileName, int pageNumber, String otherFileName, Object[] extendedArgs) throws DfException
void doGrant (String accessorName, int accessorPermit, String extendedPermission, Object[] extendedArgs) throws DfException
void doGrantPermit (IDfPermit permit, Object[] extendedArgs) throws DfException
void doLink (String folderSpec, Object[] extendedArgs) throws DfException

void doLock (Object[] extendedArgs) throws DfException
void doMark (String versionLabels, Object[] extendedArgs) throws DfException
void doPromote (String state, boolean override, boolean fTestOnly, Object[] extendedArgs) throws DfException
void doPrune (boolean keepSLabel, Object[] extendedArgs) throws DfException
IDfId doQueue (String queueOwner, String event, int priority, boolean sendMail, IDfTime dueDate, String message, Object[] extendedArgs) throws DfException
void doRefreshReference (Object[] extendedArgs) throws DfException
void doRegisterEvent (String message, String event, int priority, boolean sendMail, Object[] extendedArgs) throws DfException
void doRemovePart (IDfId containmentId, double orderNo, boolean orderNoFlag, Object[] extendedArgs) throws DfException
void doRemoveRendition (String formatName, int pageNumber, String pageModifier, boolean atomic, Object[] extendedArgs) throws DfException
String doResolveAlias (String scopeAlias, Object[] extendedArgs) throws DfException
void doResume (String state, boolean toBase, boolean override, boolean fTestOnly, Object[] extendedArgs) throws DfException
void doRevert (boolean aclOnly, Object[] extendedArgs) throws DfException
void doRevoke (String accessorName, String extendedPermission, Object[] extendedArgs) throws DfException
void doRevokePermit (IDfPermit permit, Object[] extendedArgs) throws DfException
void doSave (boolean saveLock, String versionLabel, Object[] extendedArgs) throws DfException
IDfId doSaveAsNew (boolean shareContent, boolean copyRelations, Object[] extendedArgs) throws DfException
void doScheduleDemote (String state, IDfTime scheduleDate, Object[] extendedArgs) throws DfException
void doSchedulePromote (String state, IDfTime scheduleDate, boolean override, Object[] extendedArgs) throws DfException
void doScheduleResume (String state, IDfTime scheduleDate, boolean toBase, boolean override, Object[] extendedArgs) throws DfException
void doScheduleSuspend (String state, IDfTime scheduleDate, boolean override, Object[] extendedArgs) throws DfException
void doSetACL (IDfACL acl, Object[] extendedArgs) throws DfException
void doSetFile (String fileName, String formatName, int pageNumber, String otherFile, Object[] extendedArgs) throws DfException
Note: If you are using Accelerated Content Service (ACS) or Branch Office Caching Service (BOCS), the doSetFile method is not used for content transfer. If your TBO relies on this method, the content transfer does not happen via ACS or BOCS, but the content will be brought to the application server, and doSetFile will be executed.
void doSetIsVirtualDocument (boolean treatAsVirtual, Object[] extendedArgs) throws DfException

void doSetPath (String fileName, String formatName, int pageNumber, String otherFile, Object[] extendedArgs) throws DfException
void doSuspend (String state, boolean override, boolean fTestOnly, Object[] extendedArgs) throws DfException
void doUnfreeze (boolean thawComponents, Object[] extendedArgs) throws DfException
void doUnlink (String folderSpec, Object[] extendedArgs) throws DfException
void doUnmark (String versionLabels, Object[] extendedArgs) throws DfException
void doUnRegisterEvent (String event, Object[] extendedArgs) throws DfException
void doUpdatePart (IDfId containmentId, String versionLabel, double orderNumber, boolean useNodeVerLabel, boolean followAssembly, int copyChild, String containType, String containDesc, Object[] extendedArgs) throws DfException
void doUseACL (String aclType, Object[] extendedArgs) throws DfException
void doVerifyESignature (Object[] extendedArgs) throws DfException
Methods of DfPersistentObject
IDfRelation doAddChildRelative (String relationTypeName, IDfId childId, String childLabel, boolean isPermanent, String description, Object[] extendedArgs) throws DfException
IDfRelation doAddParentRelative (String relationTypeName, IDfId parentId, String childLabel, boolean isPermanent, String description, Object[] extendedArgs) throws DfException
void doDestroy (boolean force, Object[] extendedArgs) throws DfException
void doRemoveChildRelative (String relationTypeName, IDfId childId, String childLabel, Object[] extendedArgs) throws DfException
void doRemoveParentRelative (String relationTypeName, IDfId parentId, String childLabel, Object[] extendedArgs) throws DfException
void doRevert (boolean aclOnly, Object[] extendedArgs) throws DfException
void doSave (boolean saveLock, String versionLabel, Object[] extendedArgs) throws DfException
void doSignoff (String user, String password, String reason, Object[] extendedArgs) throws DfException
Methods of DfTypedObject
void doAppendString (String attrName, String value, Object[] extendedArgs) throws DfException
String doGetString (String attrName, int valueIndex, Object[] extendedArgs) throws DfException
void doInsertString (String attrName, int valueIndex, String value, Object[] extendedArgs) throws DfException
doSetString (String attrName, int valueIndex, String value, Object[] extendedArgs) throws DfException
void doRemove (String attrName, int beginIndex, int endIndex, Object[] extendedArgs) throws DfException
Methods of DfGroup
boolean doAddGroup (String groupName, Object[] extendedArgs) throws DfException

boolean doAddUser (String userName, Object[] extendedArgs) throws DfException
void doRemoveAllGroups (Object[] extendedArgs) throws DfException
void doRemoveAllUsers (Object[] extendedArgs) throws DfException
boolean doRemoveGroup (String groupName, Object[] extendedArgs) throws DfException
boolean doRemoveUser (String userName, Object[] extendedArgs) throws DfException
void doRenameGroup (String groupName, boolean isImmediate, boolean unlockObjects, boolean reportOnly, Object[] extendedArgs) throws DfException
Methods of DfUser
void doChangeHomeDocbase (String homeDocbase, boolean isImmediate, Object[] extendedArgs) throws DfException
void doRenameUser (String userName, boolean isImmediate, boolean unlockObjects, boolean reportOnly, Object[] extendedArgs) throws DfException
void doSetAliasSet (String aliasSetName, Object[] extendedArgs) throws DfException
void doSetClientCapability (int clientCapability, Object[] extendedArgs) throws DfException
void doSetDefaultACL (String aclName, Object[] extendedArgs) throws DfException
void doSetDefaultFolder (String folderPath, boolean isPrivate, Object[] extendedArgs) throws DfException
void doSetHomeDocbase (String docbaseName, Object[] extendedArgs) throws DfException
void doSetUserOSName (String accountName, String domainName, Object[] extendedArgs) throws DfException
void doSetUserState (int userState, boolean unlockObjects, Object[] extendedArgs) throws DfException

Calling TBOs and SBOs

This section describes special considerations for using TBOs and SBOs.

Calling SBOs

This section provides rules and guidelines for instantiating SBOs and calling their methods.

The client application should instantiate a new SBO each time it needs one, rather than reusing one. For example, to call a service during an HTTP request in a web application, instantiate the service, execute the appropriate methods, then abandon the service object.

This approach is thread safe, and it is efficient, because it requires little resource overhead. The required steps to instantiate a service are:

1. Prepare an IDfLoginInfo object containing the necessary login information.
2. Instantiate a session manager object.
3. Call the service factory method.

The following code illustrates these steps:

```
IDfClient client = DfClient.getLocalClient();
IDfLoginInfo loginInfo = new DfLoginInfo();
loginInfo.setUser( strUser );
loginInfo.setPassword( strPassword );
if( strDomain != null )
    loginInfo.setDomain( strDomain );

IDfSessionManager sMgr = client.newSessionManager();
sMgr.setIdentity( strRepository, loginInfo );
IAutoNumber autonumber = (IAutoNumber)
    client.newService( IAutoNumber.class.getName(), sMgr);
```

An SBO client application uses the `newService` factory method of `IDfClient` to instantiate a service:

```
public IDfService newService ( String name, IDfSessionManager sMgr )
    throws DfServiceException;
```

The method takes the service name and a session manager as parameters, and returns the service interface, which you must cast to the specific service interface. The `newService` method uses the service name to look up the Java implementation class in the registry. It stores the session manager as a member of the service, so that the service implementation can access the session manager when it needs a DFC session.

Returning a TBO from an SBO

The following example shows how to return a TBO, or any repository object, from within an SBO method.

```
public IDfDocument getDoc( String strRepository, IDfId idDoc ) {
    IDfSession session = null;
    IDfDocument doc = null;
    try {
        session = getSession ( strRepository );
        doc = (IDfDocument)session.getObject( idDoc );
        doc.setSessionManager (getSessionManager());
    }
    finally { releaseSession( session ); }
    return doc;
}
```

Because `getDoc` is a method of an SBO, which must extend `DfService`, it has access to the session manager associated with the service. The methods `getSession`, `getSessionManager`, and `releaseSession` provide this access.

Refer to [Maintaining state in a session manager, page 46](#) for information about the substantial costs of using the `setSessionManager` method.

Calling TBOs

Client applications and methods of SBOs can use TBOs. Use a factory method of `IDfSession` to instantiate a TBO's class. Release the session when you finish with the object.

Within a method of an SBO, use `getSession` to obtain a session from the session manager. DFC releases the session when the service method is finished, making the session object invalid.

Use the `setSessionManager` method to transfer a TBO to the control of the session manager when you want to:

- Release the DFC session but keep an instance of the TBO.
- Store the TBO in the SBO state.

Refer to [Maintaining state in a session manager, page 46](#) for information about the substantial costs of using the `setSessionManager` method.

Sample SBO and TBO implementation

This section presents a straightforward example of a SBO and TBO for you to use as reference for creating your own business objects. The example is trivial — the TBO and SBO work together to set an arbitrary value (flavor) when a document is saved or checked in to the repository. The `setFlavor()` method represents the location where you can add any business logic required for your application.

ITutorialSBO

Create an interface for the service-based object. This interface provides the empty `setFlavorSBO` method, to be overridden in the implementation class. All SBOs must extend the `IDfService` interface.

Example 7-1. ITutorialSBO.java

```
package com.documentum.tutorial

import com.documentum.fc.client.IDfService;
import com.documentum.fc.client.IDfSysObject;
import com.documentum.fc.common.DfException;

public interface ITutorialSBO extends IDfService
{
    // This is our empty setFlavor method.
    public void setFlavorSBO (IDfSysObject myObj, String flavor)
        throws DfException;
}
```

TutorialSBO

The `TutorialSBO` class extends the `DfService` class, which provides fields and methods to provide common functionality for all services.

Example 7-2. TutorialSBO.java

```
package com.documentum.tutorial

import com.documentum.fc.client.DfService;
import com.documentum.fc.client.IDfSysObject;
```

```
import com.documentum.fc.common.DfException;

public class TutorialSBO extends DfService implements ITutorialSBO {

    // Overrides to standard methods.

    public String getVendorString() {
        return "Copyright (c) Documentum, Inc., 2007";
    }

    public static final String strVersion = "1.0";

    public String getVersion() {
        return strVersion;
    }

    public boolean isCompatible(String str) {
        int i = str.compareTo( getVersion() );
        if(i <= 0 )
            return true;
        else
            return false;
    }

    // Custom method. This method sets a string value on the system object.
    // You can set any number of values of any type (for example, int, double,
    // boolean) using similar methods.

    public void setFlavorSBO(IDfSysObject myObj, String myFlavor)
        throws DfException
    {
        myObj.setString("flavor",myFlavor);
    }
}
```

ITutorialTBO

The interface for the TBO is trivial — its only function is to extend the IDfBusinessObject interface, a requirement for all TBOs.

Example 7-3. ITutorialTBO.java

```
package com.documentum.tutorial

import com.documentum.fc.client.IDfBusinessObject;

public interface ITutorialTBO extends IDfBusinessObject
{
    /**
     * No code required (just extends IDfBusinessObject)
     */
}
```


TutorialTBO

The TutorialTBO is the class that pulls the entire example together. This class overrides the doSave(), doSaveEx() and doCheckin() methods of DfSystemObject and uses the setFlavorSBO() method of TutorialSBO to add a string value to objects of our custom type.

Example 7-4. TutorialTBO.java

```
package com.documentum.tutorial

import com.documentum.com.DfClientX;
import com.documentum.com.IDfClientX;

import com.documentum.fc.client.DfDocument;
import com.documentum.fc.client.IDfClient;
import com.documentum.fc.client.IDfSession;
import com.documentum.fc.client.IDfSessionManager;

import com.documentum.fc.common.DfException;
import com.documentum.fc.common.IDfId;

/**
 * simple TBO that overrides the behavior of save(), saveLock() and checkinEx()
 */

public class TutorialTBO extends DfDocument implements ITutorialTBO {

    private static final String strCOPYRIGHT =
        "Copyright (c) Documentum, Inc., 2007";

    private static final String strVERSION = "1.0";

    // Instantiate a null client.
    IDfClientX clientx = new DfClientX();
    IDfClient client = null;

    // Override standard methods.
    public String getVersion() {
        return strVERSION;
    }

    public String getVendorString() {
        return strCOPYRIGHT;
    }

    public boolean isCompatible( String s ) {
        return s.equals("1.0");
    }

    public boolean supportsFeature( String s ) {
        String strFeatures = "createhtmlfile";

        if( strFeatures.indexOf( s ) == -1 )
            return false;

        return true;
    }

    /*
     * Overridden IDfSysObject methods. These methods intercept the save(),
     * saveLock(), and checkinEx() methods, use the local setFlavor method
     * to attach a String value to the current system object, then pass
     * control to the parent class to complete the operation.
     */
}
```

```

*/

public void doSave() throws DfException {
    setFlavor("Pistachio");
    super.save();
} //end save()

public void doSaveEx() throws DfException
{
    setFlavor("Banana Fudge");
    super.saveLock();
}

public IDfId doCheckin(boolean b, String s, String s1,
    String s2, String s3, String s4) throws DfException
{
    setFlavor("Strawberry Ripple");
    return super.checkinEx(b, s, s1, s2, s3, s4);
}

// The setFlavor gets a session, session manager, and local client
// instance, then uses them to access an instance of the custom
// service-based object ITutorialSBO. Once instantiated, we can use
// the setFlavorSBO method to attach the value to the current system
// object.

private void setFlavor(String myFlavor) throws DfException {
    IDfSession session = getSession();
    IDfSessionManager sMgr = session.getSessionManager();
    client = clientx.getLocalClient();
    ITutorialSBO tutSBOObj =
        (ITutorialSBO) client.newService(ITutorialSBO.class.getName(), sMgr);
    tutSBOObj.setFlavorSBO(this, myFlavor);
}
} //end class

```

Deploying the SBO and TBO

You use Documentum Composer to package your BOF modules. *EMC Documentum Composer BOF Development Tutorial* contains information on creating and deploying a simple BOF module.

Calling DFS, JAXB API, or other components from a BOF module

To call DFS, JAXB API, or any third-party component which relies on the ThreadContext ClassLoader to load classes from within a BOF module, you will need to set the module class loader to the ThreadContext class loader. The following code saves the current ThreadContext ClassLoader, sets up ThreadContext ClassLoader as the current class ClassLoader in a try block, calls DFS (or other component), and uses the finally block to reset the ThreadContext ClassLoader to the original ThreadContext ClassLoader.

```

ClassLoader oldContextClassLoader = Thread.currentThread().getContextClassLoader();

try

```

```

{
    ClassLoader classClassLoader = this.getClass().getClassLoader();
    Thread.currentThread().setContextClassLoader(classClassLoader);
    // call dfs service here
}
catch (ServiceException e)
{
    // catch block
}
finally
{
    Thread.currentThread().setContextClassLoader(oldContextClassLoader);
}

```

Aspects

Aspects are a mechanism for adding behavior and/or attributes to a Documentum object *instance* without changing its type definition. They are similar to TBOs, but they are not associated with any one document type. Aspects also are late-bound rather than early-bound objects, so they can be added to an object or removed as needed.

Aspects are a BOF type (`dmc_aspect_type`). Like other BOF types, they have these characteristics:

- Aspects are installed into a repository.
- Aspects are downloaded on demand and cached on the local file system.
- When the code changes in the repository, aspects are automatically detected and new code is “hot deployed” to the DFC runtime.

Examples of usage

One use for aspects would be to attach behavior and attributes to objects at a particular time in their lifecycle. For example, you might have objects that represent customer contact records. When a contact becomes a customer, you could attach an aspect that encapsulates the additional information required to provide customer support. This way, the system won’t be burdened with maintenance of empty fields for the much larger set of prospective customers.

If you defined levels of support, you might have an additional level of support for “gold” customers. You could define another aspect reflecting the additional behavior and fields for the higher level of support, and attach them as needed.

Another scenario might center around document retention. For example, your company might have requirements for retaining certain legal documents (contracts, invoices, schematics) for a specific period of time. You can attach an aspect that will record the date the document was created and the length of time the document will have to be retained. This way, you are able to attach the retention aspect to documents regardless of object type, and only to those documents that have retention requirements.

You will want to use aspects any time you are introducing cross-type functionality. You can use them when you are creating elements of a common application infrastructure. You can use them when upgrading an existing data model and you want to avoid performing a database upgrade. You can use them any time you are introducing functionality on a per-instance basis.

General characteristics of aspects

Applications attach aspects to an object instance. They are not a “per application” customization — an aspect attached by one application will customize the instance across all applications. They have an affinity for the object instance, not for any particular application.

A persistent object instance may have multiple uniquely named aspects with or without attributes. Different object instances of the same persistent type may have different sets of aspects attached.

Aspects define attributes and set custom values that can be attached to a persistent object. There are no restrictions on the attributes that an aspect can define: they can be single or repeating, and of any supported data type. An aspect with an attribute definition can be attached to objects of any type — they provide a natural extension to the base type. Aspect attributes should be fully qualified as *aspect_name.attribute_name* in all setters and getters.

Attributes defined by one aspect can be accessed by other aspects. All methods work on aspect attributes transparently: fetching an object retrieves both the basic object attributes and any aspect attributes; destroying an object deletes any attached aspect attributes.

If you attach an aspect that appends a new attribute when creating an object, you must have the Write permission or higher. By contrast, the creation of an object that does not have an aspect attached only requires the Version permission.

Creating an aspect

Aspects are created in a similar fashion to other BOF modules.

1. Decide what your aspect will provide: behavior, attributes, or both.
2. Create the interface and implementation classes. Write any new behavior, override existing behavior, and provide getters and setters to your aspect attributes.
3. Deploy the aspect module using the Composer tool. For details, see the *Composer User Guide*.

As an example, we’ll walk through the steps of implementing a simple aspect. Our aspect is designed to be attached to a document that stores contact information. The aspect identifies the contact as a customer and indicates the level of service (three possible values — customer, silver, gold). It will also track the expiration date of the customer’s subscription.

Creating the aspect interface

Define the new behavior for your aspect in an interface. In this case, we’ll add getters and setters for two attributes: *service_level* and *expiration_date*.

Example 7-5. ICustomerServiceAspect.java

```
package dfctestenvironment;

import com.documentum.fc.common.DfException;
import com.documentum.fc.common.IDfTime;

public interface ICustomerServiceAspect {
```

```
// Behavior for extending the expiration date by <i>n</i> months.
public String extendExpirationDate(int months)
    throws DfException;

// Getters and setters for custom attributes.
public abstract IDfTime getExpirationDate()
    throws DfException;
public abstract String getServiceLevel()
    throws DfException;
public abstract void setExpirationDate(IDfTime expirationDate)
    throws DfException;
public abstract void setServiceLevel(String level)
    throws DfException;
}
```

Creating the aspect class

Now that we have our interface, we can implement it with a custom class.

Example 7-6. CustomerServiceAspect.java

```
package dfctutorialenvironment;

import com.documentum.fc.client.DfDocument;
import com.documentum.fc.common.DfException;
import com.documentum.fc.common.DfTime;
import com.documentum.fc.common.IDfTime;
import dfctestenvironment.ICustomerServiceAspect;

import java.util.GregorianCalendar;

public class CustomerServiceAspect
    extends DfDocument
    implements ICustomerServiceAspect
{
    public String extendExpirationDate(int months) {
        try {

            // Get the current expiration date.
            IDfTime startDate = getExpirationDate();

            // Convert the expiration date to a calendar object.
            GregorianCalendar cal = new GregorianCalendar(
                startDate.getYear(),
                startDate.getMonth(),
                startDate.getDay()
            );

            // Add the number of months -1 (months start counting from 0).
            cal.add(GregorianCalendar.MONTH, months-1);

            // Convert the recalculated date to a DfTime object.
            IDfTime endDate = new DfTime (cal.getTime());

            // Set the expiration date and return results.
            setExpirationDate(endDate);
            return "New expiration date is " + endDate.toString();
        }
        catch (Exception ex) {
            ex.printStackTrace();
        }
    }
}
```

```
        return "Exception thrown: " + ex.toString();
    }
}
// Getters and setters for the expiration_date and service_level custom attributes.
public IDfTime getExpirationDate() throws DfException
{
    return getTime("customer_service_aspect.expiration_date");
}
public String getServiceLevel() throws DfException
{
    return getString("customer_service_aspect.service_level");
}
public void setExpirationDate(IDfTime expirationDate) throws DfException {
    setTime("customer_service_aspect.expiration_date", expirationDate);
}
public void setServiceLevel (String serviceLevel) throws DfException {
    setString("customer_service_aspect.service_level", serviceLevel);
}
}
```

Deploy the customer service aspect

For details on deploying aspect modules, please see the *Documentum Foundation Classes Release Notes Version 6*

TestCustomerServiceAspect

Once you have compiled and deployed your aspect classes and defined the aspect on the server, you can use the class to set and get values in the custom aspect, and to test the behavior for adjusting the expiration date by month. This example is compatible with the sample environment described in [Chapter 4, Creating a Test Application](#).

Example 7-7. TestCustomerServiceAspect.java

```
package dfctestenvironment;

import com.documentum.com.DfClientX;
import com.documentum.com.IDfClientX;

import com.documentum.fc.client.IDfDocument;
import com.documentum.fc.client.IDfSession;
import com.documentum.fc.client.aspect.IDfAspects;

import com.documentum.fc.common.DfId;
import com.documentum.fc.common.DfTime;
import com.documentum.fc.common.IDfTime;

public class TestCustomerServletAspect {
    public TestCustomerServletAspect() {
    }

    public String attachCustomerServiceAspect(
        IDfSession mySession,
        String docId,
        String serviceLevel,
```

```

        String expirationDate
    )
    {
// Instantiate a client.
        IDfClientX clientx = new DfClientX();

        try {
            String result = "";

// Get the document instance using the document ID.
            IDfDocument doc =
                (IDfDocument) mySession.getObject(new DfId(docId));

// Convert the expirationDate string to an IDfTime object.
// DF_TIME_PATTERN1 is "mm/dd/yy"
            IDfTime ed =
                clientx.getTime(expirationDate,DfTime.DF_TIME_PATTERN1);

// Attach the aspect.
            ((IDfAspects)doc).attachAspect("customer_service_aspect",null);

// Save the document.
            doc.save();

// Get the document with its newly attached aspect.
            doc = (IDfDocument)mySession.getObject(doc.getObjectId());

// Set the aspect values.
            doc.setString("customer_service_aspect.service_level",
                serviceLevel);
            doc.setTime("customer_service_aspect.expiration_date", ed);

// Save the document.
            doc.save();

            result = "Document " + doc.getObjectName() +
                " set to service level " +
                doc.getString("customer_service_aspect.service_level") +
                " and will expire " +
                doc.getTime("customer_service_aspect.expiration_date").toString()
                + ".";
            return result;
        }
        catch (Exception ex) {
            ex.printStackTrace();
            return "Exception thrown:" + ex.toString();
        }
    }

    public String extendExpirationDate(
        IDfSession mySession,
        String docId,
        int months)
    {
// Instantiate a client.
        IDfClientX clientx = new DfClientX();

        try {
            String result = "";

// Get the document instance using the document ID.
            IDfDocument doc =
                (IDfDocument) mySession.getObject(new DfId(docId));

// Call the extendExpirationDate method
            result = ((ICustomerServiceAspect)doc).extendExpirationDate(months);

```

```
// Save the document.
    doc.save();

    return result;
}
catch (Exception ex) {
    ex.printStackTrace();
    return "Exception thrown.";
}
}

public String setExpirationDate(
    IDfSession mySession,
    String docId,
    String expDate)
{
    // Instantiate a client.
    IDfClientX clientx = new DfClientX();

    try {
        String result = "";
        IDfTime expirationDate = clientx.getTime(
            expDate,
            IDfTime.DF_TIME_PATTERN1
        );
    // Get the document instance using the document ID.
        IDfDocument doc =
            (IDfDocument) mySession.getObject(new DfId(docId));

    // Set the date using the time object.
        doc.setTime("customer_service_aspect.expiration_date",
            expirationDate);

    // Save the document and return the result.
        doc.save();
        result = "Expiration date set to " +
            doc.getTime("customer_service_aspect.expiration_date").toString();
        return result;
    }
    catch (Exception ex) {
        ex.printStackTrace();
        return "Exception thrown.";
    }
}

public String setServiceLevel(
    IDfSession mySession,
    String docId,
    String serviceLevel
)
{
    String result = "";
    // Instantiate a client.
    IDfClientX clientx = new DfClientX();

    try {

    // Get the document instance using the document ID.
        IDfDocument doc =
            (IDfDocument) mySession.getObject(new DfId(docId));

    // Set the date using the time object.
        doc.setString("customer_service_aspect.service_level",
            serviceLevel);
    }
```



```
// Save the document.
doc.save();
result = "Service Level set to " +
    doc.getString("customer_service_aspect.service_level")
    + ".";
return result;
}
catch (Exception ex) {
    ex.printStackTrace();
    return "Exception thrown.";
}
}
}
```

Using aspects in a TBO

Attaching and detaching aspects from within a TBO or aspect requires the use of callbacks to execute aspect methods or access an aspect's attributes. The attach/detach is encapsulated within the object instance. Aspect behavior and attributes are not available until the object has been fully initialized.

To use aspects within a TBO or aspect, you must

- Create a class that implements IDfAttachAspectCallback or IDfDetachAspectCallback.
- Implement the doPostAttach() or doPostDetach() method of the callback interface.

In the following examples, the attachAspect() method is called inside the overridden version of the doSave() method, but only if it is not already attached. Attempting to attach an aspect more than once will result in an exception. The MyAttachCallback callback implementation sets the attributes *retained_since* and *years_to_retain* in the new aspect.

Example 7-8. Code snippet from MySopTBO.java

```
...
//Override doSave()
protected synchronized void doSave(
    boolean saveLock,
    String v,
    Object[] args)
{
    if (this.getAspects().findString("my_retention_aspect") < 0) {
        MyAttachCallback myCallback = new MyAttachCallback();
        this.attachAspect("my_retention_aspect", myCallback);
    }
    super.doSave(saveLock, v, args);
}
```

Example 7-9. MyAttachCallback.java

```
public class MyAttachCallback implements IDfAttachAspectCallback
{
    public void doPostAttach(IDfPersistentObject obj) throws Exception
    {
        obj.setTime("my_retention_aspect.retained_since", Datevalue);
        obj.setInt("my_retention_aspect.years_to_retain", Years);
        obj.save();
    }
}
```

```
}  
}
```

Using DQL with aspects

Once an aspect has been defined in the repository, you can use DQL (Documentum Query Language) instructions to add, modify, or drop attributes. Aspects can be modified using the following commands:

```
ALTER ASPECT aspect_name ADD attribute_def[,attribute_def] [OPTIMIZEFETCH|  
NO OPTIMIZEFETCH] [PUBLISH]
```

```
ALTER ASPECT aspect_name MODIFY attribute_modifier_clause[, attribute_modifier_clause]  
[PUBLISH]
```

```
ALTER ASPECT aspect_name DROP attribute_name[, attribute_name] [PUBLISH]
```

```
ALTER ASPECT aspect_name DROP ALL [PUBLISH]
```

The syntax for the *attribute_def* and *attribute_modifier_clause* is the same as the syntax for the ALTER TYPE statement.

The OPTIMIZEFETCH keyword in the ALTER ASPECT statement causes the aspect attribute values to be stored alongside the object to which the aspect is attached. This results in reduced database queries and can improve performance. The aspect attributes are duplicated into the object's property bag (a new Documentum 6 feature). The trade off is the increased storage cost of maintaining more than one instance of the attribute value.

There are some limitations when using the DQL SELECT statement. If you are selecting a repeating aspect attribute, *r_object_id* should be included in the selected list. Repeating aspect attributes cannot be in the select list of a sub_query. Repeating aspect attributes from different aspects cannot be referenced in an expression. If the select list contains an aggregate function on a repeating aspect attribute, then the 'GROUP BY' clause, if any, must be on *r_object_id*.

Data Dictionary-specific clauses are available in the ALTER ASPECT statement, but the semantics (validations, display configuration, etc.) are not supported in the Documentum 6 release.

Enabling aspects on object types

By default, *dm_sysobject* and its sub-types are enabled for aspects. This includes any custom object sub-types. Any non-sysobject application type can be enabled for use with aspects using the following syntax.

```
ALTER TYPE type_name ALLOW ASPECTS
```

Default aspects

Type definitions can include a default set of aspects. This allows you to modify the data model and behavior for future instances. It also ensures that specific aspects are attached to all selected object instances, no matter which application creates the object. The syntax is

```
ALTER TYPE type_name [SET | ADD | REMOVE] DEFAULT ASPECTS aspect_list
```

the *aspect_list* value is a comma-separated list of *dmc_aspect_type* *object_name* values. No quotes are necessary, but if you choose to use quotes they must be single quotes and surround the entire list. For example, *aspect_list* could be a single value such as *my_retention_aspect*, or it could be multiple values specified as '*my_aspect_name1*, *my_aspect_name2*' or *my_aspect_name1*, *my_aspect_name2*.

When you attach a default aspect to a type, all the subtypes of that type will also have that aspect attached. If you attach a default aspect to *dm_sysobject*, that aspect will be inherited by all the *dm_sysobject* subtypes. However, some of the *dm_sysobject* subtypes are implemented as final types and will throw a runtime error, "java.lang.VerifyError: Cannot inherit from final class". Do not attach default aspects to *dm_sysobject*. Instead, attach it to the specific types that will use the aspect. Alternatively, create a subtype of *dm_sysobject*, attach the default aspect to that, and derive your other subtypes from that subtype of *dm_sysobject*, rather than *dm_sysobject* itself.

Referencing aspect attributes from DQL

All aspect attributes in a DQL statement must be fully qualified as *aspect_name.attribute_name*. For example:

```
SELECT r_object_id, my_retention_aspect.retained_since
FROM my_sop WHERE my_retention_aspect.years_to_retain = 10
```

If more than one type is specified in the FROM clause of a DQL statement, aspect attributes should be further qualified as *type_name.aspect_name.attribute_name* OR *alias_name.aspect_name.attribute_name*.

Aspect attributes specified in a DQL statement appear in a DQL statement like a normal attribute, wherever legally permitted by the syntax.

Full-text index

By default, full-text indexing is turned off for aspects. You can control which aspect attributes have full-text indexing using the following DQL syntax.

```
ALTER ASPECT aspect_name FULLTEXT SUPPORT ADD | DROP a1, a2,...
```

```
ALTER ASPECT aspect_name FULLTEXT SUPPORT ADD | DROP ALL
```

Object replication

Aspect attributes can be replicated in a second repository just as normal attributes are replicated (“dump and load” procedures). However, the referenced aspects must be available on the target repository.

Working with Virtual Documents

This chapter covers some of the basics of working programmatically with virtual documents. It includes the following sections:

- [Understanding virtual documents, page 165](#)
- [Setting version labels, page 166](#)
- [Getting version labels, page 169](#)
- [Creating a virtual document, page 171](#)
- [Traversing the virtual document structure, page 174](#)
- [Binding to a version label, page 176](#)
- [Clearing a version label binding, page 179](#)
- [Removing a virtual document child, page 180](#)

Understanding virtual documents

A virtual document is a data-driven construction that allows you to treat more than one system object as a single document. The components, or *nodes*, of the document remain their individual identities, but are related to one another using `dmr_containment` objects to link a parent node to child nodes in a hierarchical structure.

Multiplicitas componatis res simplex (Complexity is composed of simple things). While the structure of a virtual document can be quite complex taken as a whole, the relationship of any one node to its parent is straightforward and easy to follow.

Virtual documents are commonly created in three ways. When you import a document with Microsoft OLE links, it can be added as a parent, with the linked objects imported as child nodes. XML documents can be automatically converted to virtual documents on import using an XML application. When working with documents created automatically from OLE or XML documents, you are able to check out and edit individual nodes of a virtual document. When you check out the parent document, the document is reassembled into one cohesive unit.

The third way is to “manually” convert a document to a virtual document (programmatically or through an interface) and add arbitrary repository objects as children. The virtual document imposes a hierarchical structure that can help to manage the components as a single unit (for example, you can check out a parent document and all descendants in a single operation).

Virtual documents can be linked to what is always the CURRENT version of a child node, or they can be linked to a specific version by its version label (either the system-supplied version number or a user-supplied symbolic version label). For example, you could label a version of each document in a publication *August* and create a virtual document linking to the versions of the nodes that have that label. This way, regardless of the updates made to the document for September, the August virtual document would point to the nodes used at the time of its publication.

Setting version labels

To link to a particular version of a document, you first need to set a version label. Setting a version label requires that you check out the document, then check it in again with the version label you want to use. You can set more than one label, using a single, comma-delimited String value. If you want the document to also be the CURRENT version, you must explicitly include that label, as well.

To add a Set Version Labels button to the DFC Base Tutorial Frame

1. Create a JTextField control named jTextField_versionLabel.
2. Create a JButton control named jButton_setVersionLabel.
3. Update the [TutorialCheckIn](#) class to set the version label on check in.
4. Create the class [TutorialSetVersion](#).
5. Add a button handler method for [Set Version Label](#).

Example 8-1. Updated TutorialCheckIn class

```
package com.emc.tutorial;

import com.documentum.com.DfClientX;
import com.documentum.com.IDfClientX;
import com.documentum.fc.client.IDfDocument;
import com.documentum.fc.client.IDfSession;
import com.documentum.fc.client.IDfSessionManager;
import com.documentum.fc.client.IDfSysObject;
import com.documentum.fc.common.DfId;
import com.documentum.fc.common.IDfId;
import com.documentum.operations.IDfCheckinNode;
import com.documentum.operations.IDfCheckinOperation;

public class TutorialCheckIn
{
    public TutorialCheckIn()
    {
    }
    public String checkinLabelExample(
        IDfSessionManager sessionManager,
        String repositoryName,
        String docId,
        // Add a variable for incoming version labels, entered as
        // a single comma-delimited string.
        String versionLabels
    )
    {
        IDfSession mySession = null;
        try
        {
```

```

        mySession = sessionManager.getSession(repositoryName);

// Get the object ID based on the object ID string.
        IDfId idObj =
            mySession.getIdByQualification(
                "dm_sysobject where r_object_id='" + docId + "'"
            );

// Instantiate an object from the ID.

        IDfSysObject sysObj = (IDfSysObject) mySession.getObject(idObj);

// Instantiate a client.
        IDfClientX clientx = new DfClientX();

// Use the factory method to create an IDfCheckinOperation instance.
        IDfCheckinOperation cio = clientx.getCheckinOperation();

// Set the version increment. In this case, the next major version
// ( version + 1)
        cio.setCheckinVersion(IDfCheckinOperation.NEXT_MAJOR);

// Set one or more comma-delimited labels for this version.
// When updating to the next major version, you need to explicitly
// set the version label for the new object to "CURRENT".
        cio.setVersionLabels(versionLabels);

// Create a document object that represents the document being
// checked in.
        IDfDocument doc =
            (IDfDocument) mySession.getObject(new DfId(docId));

// Create a checkin node, adding it to the checkin operation.
        IDfCheckinNode node = (IDfCheckinNode) cio.add(doc);

// Execute the checkin operation and return the result.
        if (!cio.execute())
        {
            return "Checkin failed.";
        }

// After the item is created, you can get it immediately using the
// getNewObjectId method.

        IDfId newId = node.getNewObjectId();
        return "Checkin succeeded - new object ID is: " + newId;
    }
    catch (Exception ex)
    {
        ex.printStackTrace();
        return "Checkin failed.";
    }
    finally
    {
        sessionManager.release(mySession);
    }
}
}

```

Example 8-2. The TutorialSetVersion class

```
package com.emc.tutorial;
```

```
import com.documentum.fc.client.IDfSession;
import com.documentum.fc.client.IDfSessionManager;

public class TutorialSetVersion
{
    public TutorialSetVersion()
    {
    }

    public String setVersion(
        IDfSessionManager sessionManager,
        String repositoryName,
        String documentIdString,
        String versionLabels
    )
    {
        IDfSession mySession = null;
        StringBuffer sb = new StringBuffer("");
        try
        {
            mySession = sessionManager.getSession(repositoryName);

            // Check out the document.
            TutorialCheckOut tco = new TutorialCheckOut();
            tco.checkoutExample(
                sessionManager,
                repositoryName,
                documentIdString
            );

            // Check in the document, specifying the version labels.
            TutorialCheckIn tci = new TutorialCheckIn();
            sb.append( "\n" +
                tci.checkinLabelExample(
                    sessionManager,
                    repositoryName,
                    documentIdString,
                    versionLabels
                )
            );
            return "Success!\n*****\n\n" + sb.toString();
        }
        catch (Exception ex)
        {
            ex.printStackTrace();
            return ("Failure!\n*****\n\n" + sb.toString());
        }
        finally
        {
            sessionManager.release(mySession);
        }
    }
}
```

Example 8-3. Handler for the Set Version Label button

```
private void jButton_setVersionLabel_actionPerformed(ActionEvent e)
{
    if (m_parentId == null)
    {
        jLabel_messages.setText("Please reset the virtual document parent.");
    }
}
```



```

else
{
    String repositoryName = jTextField_repositoryName.getText();
    String childIdString =
        m_fileIDs.elementAt(list_id.getSelectedIndex()).toString();
    String versionLabel = jTextField_versionLabel.getText();
    TutorialSetVersion tsv =
        new TutorialSetVersion();
    JTextArea_results.setText(
        tsv.setVersion(
            m_sessionManager,
            repositoryName,
            childIdString,
            versionLabel
        )
    );
    initDirectory();
    getDirectory();
}
}

```

Getting version labels

You can query to get the labels from a document in the system in preparation for linking to a specific version of the document. There is, in fact, an `IDfVersionLabels` interface specially designed to let you query and manipulate the labels on a system object.

To add a Get Version Labels button to the DFC Base Tutorial Frame

1. If you haven't done so already, create a `JTextField` control named `jTextField_versionLabel`.
2. Create a `JButton` control named `jButton_getVersionLabel`.
3. Create the class [TutorialGetVersion](#).
4. Add a button handler method for [Get Version Label](#).

Example 8-4. The `TutorialGetVersion` class

```

package com.emc.tutorial;

import com.documentum.fc.client.IDfSession;
import com.documentum.fc.client.IDfSessionManager;
import com.documentum.fc.client.IDfSysObject;
import com.documentum.fc.client.IDfVersionLabels;
import com.documentum.fc.common.IDfId;

public class TutorialGetVersion
{
    public TutorialGetVersion()
    {
    }

    public String getVersion(
        IDfSessionManager sessionManager,
        String repositoryName,
        String documentIdString
    )
    {
        IDfSession mySession = null;
    }
}

```

```
        StringBuffer sb = new StringBuffer("");
        try
        {
            mySession = sessionManager.getSession(repositoryName);

// Get the object ID based on the object ID string.
            IDfId idObj =
                mySession.getIdByQualification(
                    "dm_sysobject where r_object_id='" + documentIdString + "'"
                );

// Instantiate an object from the ID.

            IDfSysObject sysObj = (IDfSysObject) mySession.getObject(idObj);

            IDfVersionLabels vl = sysObj.getVersionLabels();
            for (int i=0; i<vl.getVersionLabelCount();i++)
            {
                if (i==0) {
                    sb.append(vl.getVersionLabel(i));
                }
                else
                {
                    sb.append(", " + vl.getVersionLabel(i));
                }
            }
            return sb.toString();
        }
        catch (Exception ex)
        {
            ex.printStackTrace();
            return ("Failed to get version labels for selected object.");
        }
        finally
        {
            sessionManager.release(mySession);
        }
    }
}
```

Example 8-5. Handler for the Get Version Label button

```
private void jButton_getVersionLabel_actionPerformed(ActionEvent e)
{
    String repositoryName = jTextField_repositoryName.getText();
    String documentIdString =
        m_fileIDs.elementAt(list_id.getSelectedIndex()).toString();
    TutorialGetVersion tgv = new TutorialGetVersion();
    jTextField_versionLabel.setText(
        tgv.getVersion(
            m_sessionManager,
            repositoryName,
            documentIdString
        )
    );
    initDirectory();
    getDirectory();
}
```

Creating a virtual document

Creating a virtual document is a fairly straightforward exercise. If you select any system object (with the exception of folders and cabinets), you can get use the `DfClient.asVirtualDocument()` method to indicate that you want to convert the object into a virtual document. It is not persisted as a virtual document in the repository, though, until you add one or more child nodes and save the document. You can, and should, set the `r_is_virtual_document` property on the file to `TRUE` as a signal to other applications that the document should be treated as a virtual document, then set the value to `FALSE` if the children are subsequently removed.

In this example, we will add a field for the parent name, and a button that populates the field and captures the ID of the selected item in a global variable. You can then select another document to add it as a child node.

By default, the child node is added as the first node of the virtual document. If you add another child node it becomes the first node, and the existing node moves down to become the second child node. If you want to control the placement of the node in the virtual document hierarchy, you have the option of passing a sibling node as an argument to the `addNode()` method. The new node is placed after the sibling node in the virtual document hierarchy. We will add a field to capture the name of an optional virtual document sibling.

The only new class created for this example is the `TutorialAddVirtualDocumentNode` class. The remaining button handlers (`Set Virtual Document Parent`, `Set Preceding Sibling`, `Clear Preceding Sibling`) capture supporting information and set member variables for use by the `Add Virtual Document Node` button handler.

To add an Add Virtual Document Node button to the DFC Base Tutorial Frame

1. Create a `TextField` control named `textField_virtualDocumentParent`.
2. Create a `Button` control named `button_setVirtualDocumentParent`.
3. Create a `TextField` control named `textField_precedingSibling`.
4. Create a `Button` control named `button_setPrecedingSibling`.
5. Create a `Button` control named `button_clearSibling`.
6. Create the class `TutorialAddVirtualDocumentNode`.
7. Add a button handler method for `Set Virtual Document Parent`.
8. Add a button handler method for `Set Preceding Sibling`.
9. Add a button handler method for `Clear Sibling`.
10. Add a button handler method for `Add Virtual Document Node`.

Example 8-6. Handler for the Set Virtual Document Parent button

```
private void setVirtualDocumentParent()
{
    // Store the parent ID in a global variable for use by other operations.
    m_parentId = m_fileIDs.elementAt(list_id.getSelectedIndex()).toString();
    textField_virtualDocumentParent.setText(
        textField_cwd.getText() + "/" + list_id.getSelectedItem()
    );
}
```

```
jTextArea_results.setText(m_parentId + "\n" + list_id.getSelectedItem());
}
```

Example 8-7. Handler for the Set Preceding Sibling button

```
private void jButton_setPrecedingSibling_actionPerformed(ActionEvent e)
{
    m_siblingId = m_fileIDs.elementAt(list_id.getSelectedIndex()).toString();
    jTextField_precedingSibling.setText(
        list_id.getSelectedItem()
    );
    jTextArea_results.setText(m_siblingId + "\n" + list_id.getSelectedItem());
}
```

Example 8-8. Handler for the Clear Sibling button

```
private void jButton_clearSibling_actionPerformed(ActionEvent e)
{
    m_siblingId = "";
    jTextField_precedingSibling.setText("");
}
```

Example 8-9. Handler for the Add Virtual Document Node button

```
private void jButton_addVirtualDocumentNode_actionPerformed(ActionEvent e)
{
    if (m_parentId == null)
    {
        jLabel_messages.setText("Please reset the virtual document parent.");
    }
    else
    {
        String repositoryName = jTextField_repositoryName.getText();
        String childIdString =
            m_fileIDs.elementAt(list_id.getSelectedIndex()).toString();
        String bindingLabel = jTextField_versionLabel.getText();
        TutorialAddVirtualDocumentNode tavdn =
            new TutorialAddVirtualDocumentNode();
        jTextArea_results.setText(
            tavdn.addNode(
                m_sessionManager,
                repositoryName,
                m_parentId,
                childIdString,
                m_siblingId,
                bindingLabel
            )
        );
        initDirectory();
        getDirectory();
        m_parentId = null;
    }
}
```

Example 8-10. The TutorialAddVirtualDocumentNode class

```
package com.emc.tutorial;

import com.documentum.com.DfClientX;
import com.documentum.com.IDfClientX;
```

```

import com.documentum.fc.client.IDfSession;
import com.documentum.fc.client.IDfSessionManager;
import com.documentum.fc.client.IDfSysObject;
import com.documentum.fc.client.IDfVirtualDocument;
import com.documentum.fc.client.IDfVirtualDocumentNode;
import com.documentum.fc.common.IDfId;

public class TutorialAddVirtualDocumentNode
{
    public TutorialAddVirtualDocumentNode()
    {
    }
    public String addNode(
        IDfSessionManager sessionManager,
        String repositoryName,
        String parentIdString,
        String childIdString,
        String siblingIdString,
        String versionLabel
    )
    {
        IDfSession mySession = null;
        StringBuffer sb = new StringBuffer("");
        try
        {
            mySession = sessionManager.getSession(repositoryName);

// Check out the parent object
            TutorialCheckOut tco = new TutorialCheckOut();
            tco.checkoutExample(
                sessionManager,
                repositoryName,
                parentIdString
            );

            IDfClientX clientx = new DfClientX();

//Instantiate the parent object.
            IDfVirtualDocument vDoc = null;
            IDfId parentIdObject = clientx.getId(parentIdString);
            IDfSysObject sysObj =
                (IDfSysObject) mySession.getObject(parentIdObject);
            if (versionLabel.equals("")) versionLabel = null;
            if (sysObj != null)
            {

// Instantiate the parent as a virtual document.
                vDoc = sysObj.asVirtualDocument(null,false);

// Instantiate the root node of the virtual document.
                IDfVirtualDocumentNode root = vDoc.getRootNode();

// Create an ID object for the child node.
                IDfId childId = mySession.getIdByQualification(
                    "dm_sysobject where r_object_id='" + childIdString + "'"
                );

//Instantiate the child as a sysobject.
                IDfSysObject childObj =
                    (IDfSysObject) mySession.getObject(childId);

// Instantiate the sibling object (if not null) as a virtual document node.
                IDfVirtualDocumentNode siblingNode = null;
                if (!siblingIdString.equals(null) & !siblingIdString.equals("")){
                    IDfId siblingId = mySession.getIdByQualification(

```

```
        "dm_sysobject where r_object_id='" + siblingIdString + "'"
    );
    IDfSysObject siblingObj =
        (IDfSysObject) mySession.getObject(siblingId);
    String siblingChronId = siblingObj.getChronicleId().toString();
    siblingNode = vDoc.find(root,siblingChronId,"i_chronicle_id",0);
}

// Add the child to the virtual document.
IDfVirtualDocumentNode childVDNode = vDoc.addNode(
    root,
    siblingNode,
    childObj.getChronicleId(),
    versionLabel,
    false,
    false
);
}

// Check in the parent document.
TutorialCheckIn tci = new TutorialCheckIn();
sb.append( "\n" +
    tci.checkinExample(
        sessionManager,
        repositoryName,
        parentIdString
    )
);
return "Success!\n*****\n\n" + sb.toString();
}
catch (Exception ex)
{
    ex.printStackTrace();
    return ("Failure!\n*****\n\n" + sb.toString());
}
finally
{
    sessionManager.release(mySession);
}
}
```

Traversing the virtual document structure

Once you have created a virtual document, you can use a recursive routine to traverse the hierarchical structure and display information about each node.

One side note: in the past, there have been other methods requiring several statements in order to get the underlying system object represented by a virtual document node. Going forward, you should use the method `IDfVirtualDocumentChild.getSelectedObject()` to get the object represented by a virtual document node.

To add a Traverse Virtual Document button to the DFC Base Tutorial Frame

1. Create a JButton control named `JButton_traverseVirtualDocument`.
2. Create the [TutorialTraverseVirtualDocument](#) class.

3. Create the handler for the [Traverse Virtual Document](#) button.

Example 8-11. The TutorialTraverseVirtualDocument class

```
package com.emc.tutorial;

import com.documentum.fc.client.IDfSession;
import com.documentum.fc.client.IDfSessionManager;
import com.documentum.fc.client.IDfSysObject;
import com.documentum.fc.client.IDfVirtualDocument;
import com.documentum.fc.client.IDfVirtualDocumentNode;
import com.documentum.fc.common.IDfId;

public class TutorialTraverseVirtualDocument
{
    public TutorialTraverseVirtualDocument()
    {
    }
    public String traverseVirtualDocument(
        IDfSessionManager sessionManager,
        String repositoryName,
        String docId
    )
    {
        IDfSession mySession = null;
        StringBuffer results = new StringBuffer("");
        try
        {
            mySession = sessionManager.getSession(repositoryName);

// Get the object ID based on the object ID string.
            IDfId idObj =
                mySession.getIdByQualification(
                    "dm_sysobject where r_object_id='" + docId + "'"
                );

// Instantiate an object from the ID.

            IDfSysObject sysObj = (IDfSysObject) mySession.getObject(idObj);
            IDfVirtualDocument parent = sysObj.asVirtualDocument(null,false);
            IDfVirtualDocumentNode node = parent.getRootNode();

// Capture useful information about the root node.
            results.append( "\nNode Name: " + sysObj.getObject_name() +
                "\nVersion: " + sysObj.getVersionLabel(0) );

            int childCount = node.getChildCount();
            IDfVirtualDocumentNode child = null;
            for( int i = 0; i < childCount; ++i )
            {
                child = node.getChild( i );

// If the child is also a virtual document, make a nested
// call to this method.
                if( child.getSelectedObject().isVirtualDocument() )
                {
                    sysObj = child.getSelectedObject();
                    results.append(
                        traverseVirtualDocument(
                            sessionManager,
                            repositoryName,
                            sysObj.getObjectId().toString()
                        )
                    );
                }
            }
        }
    }
}
```

```
        }
        else
        {
// Otherwise, capture interesting information about the child node.
            sysObj = child.getSelectedObject();
            results.append( "\nNode Name: " + sysObj.getObject_name() +
                "\n Version: " + sysObj.getVersionLabel(0) +
                "\n VDM Number: " + child.getVDMNumber() +
                "\n Binding Label: " + child.getBinding());
        }
    }
    return results.toString();

}
catch (Exception ex)
{
    ex.printStackTrace();
    return ("Traversal failed.");
}
finally
{
    sessionManager.release(mySession);
}
}
```

Example 8-12. Handler for the Traverse Virtual Document button

```
private void jButton_traverse_actionPerformed(ActionEvent e)
{
    String repositoryName = jTextField_repositoryName.getText();
    String docId =
        m_fileIDs.elementAt(list_id.getSelectedIndex()).toString();

    TutorialTraverseVirtualDocument ttvd =
        new TutorialTraverseVirtualDocument();
    jTextArea_results.setText(
        ttvd.traverseVirtualDocument(
            m_sessionManager,
            repositoryName,
            docId
        )
    );
}
```

Binding to a version label

In practical use, you may want to bind to a particular version label after the virtual document has been created. To do this, you check out the parent node, instantiate the child node based on its chronicle ID, use the `IDfVirtualDocumentNode.setBinding(String bindingLabel)` method to assign the binding value, then check in the parent node.

To add a Set Binding button to the DFC Base Tutorial Frame

1. Create a `jButton` control named `jButton_setBinding`.
2. Create the [TutorialSetBinding](#) class.

3. Create the handler for the [Set Binding](#) button.

Example 8-13. The TutorialSetBinding class

```
package com.emc.tutorial;

import com.documentum.com.DfClientX;
import com.documentum.com.IDfClientX;
import com.documentum.fc.client.IDfSession;
import com.documentum.fc.client.IDfSessionManager;
import com.documentum.fc.client.IDfSysObject;
import com.documentum.fc.client.IDfVirtualDocument;
import com.documentum.fc.client.IDfVirtualDocumentNode;
import com.documentum.fc.common.IDfId;
import com.documentum.operations.IDfCheckinNode;

public class TutorialSetBinding
{
    public TutorialSetBinding()
    {
    }

    public String setBinding (
        IDfSessionManager sessionManager,
        String repositoryName,
        String parentIdString,
        String childIdString,
        String bindingLabel // An existing label already assigned to the child node.
    )
    {
        IDfSession mySession = null;
        StringBuffer sb = new StringBuffer("");
        try
        {
            // Start by checking out the parent document.
            mySession = sessionManager.getSession(repositoryName);
            TutorialCheckOut tco = new TutorialCheckOut();
            tco.checkoutExample(
                sessionManager,
                repositoryName,
                parentIdString
            );

            IDfClientX clientx = new DfClientX();

            IDfVirtualDocument vDoc = null;

            // Use the parent ID to instantiate a sysobject.
            IDfId parentIdObject = clientx.getId(parentIdString);
            IDfSysObject sysObj =
                (IDfSysObject) mySession.getObject(parentIdObject);

            if (sysObj != null)
            {
                // Get a virtual document instance of the parent.
                vDoc = sysObj.asVirtualDocument(null, false);
                IDfVirtualDocumentNode root = vDoc.getRootNode();

                // Get an instance of the child document.
                IDfId childId = mySession.getIdByQualification(
                    "dm_sysobject where r_object_id='" + childIdString + "'"
                );
                IDfSysObject childObj =
```

```
(IDfSysObject) mySession.getObject(childId);

// Use the current instance to get the chronicle ID of the child.
String childChronId = childObj.getChronicleId().toString();

// Instantiate a virtual document node based on the child's chronicle ID.
IDfVirtualDocumentNode childNode =
    vDoc.find(root,childChronId,"i_chronicle_id",0);

// Set the child's binding to the label provided.
childNode.setBinding(bindingLabel);
}

// Check in the parent document.
TutorialCheckIn tci = new TutorialCheckIn();
sb.append( "\n" +
    tci.checkinExample(
        sessionManager,
        repositoryName,
        parentIdString
    )
);
return "Success!\n*****\n\n" + sb.toString();
}
catch (Exception ex)
{
    ex.printStackTrace();
    return ("Failure!\n*****\n\n" + sb.toString());
}
finally
{
    sessionManager.release(mySession);
}
}
```

Example 8-14. Handler for the Set Binding button

```
private void jButton_setBinding_actionPerformed(ActionEvent e)
{
    if (m_parentId == null)
    {
        jLabel_messages.setText("Please reset the virtual document parent.");
    }
    else
    {
        String repositoryName = jTextField_repositoryName.getText();
        String childIdString =
            m_fileIDs.elementAt(list_id.getSelectedIndex()).toString();
        String bindingLabel = jTextField_versionLabel.getText();
        TutorialSetBinding tsb = new TutorialSetBinding();
        jTextArea_results.setText(
            tsb.setBinding(
                m_sessionManager,
                repositoryName,
                m_parentId,
                childIdString,
                bindingLabel
            )
        );
        initDirectory();
        getDirectory();
    }
}
```

```
// Nullify the parent ID global variable, as a new version of the
// parent document has been saved, and the existing value is no
// longer current.
    m_parentId = null;
}
}
```

Clearing a version label binding

To clear a binding based on version label, you can set the binding to the value `CURRENT`, which will restore the default behavior of always including the `CURRENT` version of the child object as a virtual document node.

To add a Set Binding button to the DFC Base Tutorial Frame

1. Create a `jButton` control named `jButton_clearBinding` .
2. If you haven't done so already, create the [TutorialSetBinding](#) class.
3. Create the handler for the [Clear Binding](#) button.

Example 8-15. Handler for the Clear Binding button

```
private void jButton_clearBinding_actionPerformed(ActionEvent e)
{
    if (m_parentId == null)
    {
        jLabel_messages.setText("Please reset the virtual document parent.");
    }
    else
    {
        String repositoryName = jTextField_repositoryName.getText();
        String childIdString =
            m_fileIDs.elementAt(list_id.getSelectedIndex()).toString();
        String bindingLabel = "CURRENT";
        TutorialSetBinding tsb = new TutorialSetBinding();
        jTextArea_results.setText(
            tsb.setBinding(
                m_sessionManager,
                repositoryName,
                m_parentId,
                childIdString,
                bindingLabel
            )
        );
        initDirectory();
        getDirectory();
        // Nullify the parent ID global variable, as a new version of the
        // parent document has been saved, and the existing value is no
        // longer current.
        m_parentId = null;
    }
}
```

Removing a virtual document child

Removing a virtual document child is another operation that takes a little extra wangling, because you need to remove the child node based on its chronicle ID. Your code needs to check out the parent document, instantiate it as a virtual document, locate the child node, get its ID, use that to find the child node's chronicle ID, instantiate the child virtual document node using the chronicle ID, pass the node as an argument to the `IDfVirtualDocument.remove(IDfVirtualDocumentNode childNode)` method, then check in the parent document.

You should also test to see if you are removing the last node from the virtual document (if the child count is now 0). If so, use the method `IDfSysObject.setIsVirtualDocument()` to set the property to `FALSE`.

To add a Remove Virtual Document Child button to the DFC Base Tutorial Frame

1. Create a `jButton` control named `jButton_removeVirtualDocumentChild` .
2. Create the [TutorialRemoveVirtualDocumentNode](#) class.
3. Create the handler for the [Remove Virtual Document Child](#) button.

Example 8-16. The TutorialRemoveVirtualDocumentNode class

```
package com.emc.tutorial;

import com.documentum.com.DfClientX;
import com.documentum.com.IDfClientX;
import com.documentum.fc.client.IDfSession;
import com.documentum.fc.client.IDfSessionManager;
import com.documentum.fc.client.IDfSysObject;
import com.documentum.fc.client.IDfVirtualDocument;
import com.documentum.fc.client.IDfVirtualDocumentNode;
import com.documentum.fc.common.IDfId;

public class TutorialRemoveVirtualDocumentNode
{
    public TutorialRemoveVirtualDocumentNode()
    {
    }
    public String removeNode(
        IDfSessionManager sessionManager,
        String repositoryName,
        String parentIdString,
        String childIdString
    )
    {
        IDfSession mySession = null;
        StringBuffer sb = new StringBuffer("");
        try
        {
            mySession = sessionManager.getSession(repositoryName);

            // Check out the parent document.
            TutorialCheckOut tco = new TutorialCheckOut();
            tco.checkoutExample(
                sessionManager,
                repositoryName,
                parentIdString
            );

            IDfClientX clientx = new DfClientX();
```

```

// Create an empty virtual document object.
IDfVirtualDocument vDoc = null;

// Get the parent object
IDfId parentIdObject = clientx.getId(parentIdString);
IDfSysObject sysObj =
    (IDfSysObject) mySession.getObject(parentIdObject);

    if (sysObj != null)
    {

// Get the parent document as a virtual document.
        vDoc = sysObj.asVirtualDocument(null,false);

// Get the root node of the virtual document.
        IDfVirtualDocumentNode root = vDoc.getRootNode();

// Create an ID object based on the child ID string.
        IDfId childId = mySession.getIdByQualification(
            "dm_sysobject where r_object_id='" + childIdString + "'"
        );
// Instantiate the child object.
        IDfSysObject childObj =
            (IDfSysObject) mySession.getObject(childId);

// Get the child object's chronicle ID
        String childChronId = childObj.getChronicleId().toString();

// Get the child object as a virtual document node.
        IDfVirtualDocumentNode childNode =
            vDoc.find(root,childChronId,"i_chronicle_id",0);
        sb.append("Child ID: " + childId.toString());
        sb.append("\nChild Chron ID: " + childChronId);
        sb.append("\nChild Node: " + childNode.toString());

// And here is the actual command that removes the node.
        vDoc.removeNode(childNode);

// If you have removed the last virtual document node,
// set the r_is_virtual_document property to FALSE.
        if (root.getChildCount()==0)
        {
            sysObj.setIsVirtualDocument(false);
        }
    }

// Check in the parent virtual document to save the change.
TutorialCheckIn tci = new TutorialCheckIn();
sb.append( "\n" +
    tci.checkinExample(
        sessionManager,
        repositoryName,
        parentIdString
    )
);
return "Success!\n*****\n" + sb.toString();
}
catch (Exception ex)
{
    ex.printStackTrace();
    return ("Failure!\n*****\n" + sb.toString());
}
finally
{

```

```
        sessionManager.release(mySession);
    }
}
```

Example 8-17. Handler for the Remove Virtual Document Child button

```
private void jButton_removeVirtualDocumentChild_actionPerformed(ActionEvent e)
{
    if (m_parentId == null)
    {
        jLabel_messages.setText("Please reset the parent field.");
    }
    else
    {
        String repositoryName = jTextField_repositoryName.getText();
        String childIdString =
            m_fileIDs.elementAt(list_id.getSelectedIndex()).toString();
        TutorialRemoveVirtualDocumentNode trvdn =
            new TutorialRemoveVirtualDocumentNode();
        jTextArea_results.setText(
            trvdn.removeNode(
                m_sessionManager,
                repositoryName,
                m_parentId,
                childIdString
            )
        );
        initDirectory();
        getDirectory();

        // Nullify the parent ID global variable, as a new version of the
        // parent document has been saved, and the existing value is no
        // longer current.
        m_parentId = null;
    }
}
```

Support for Other Documentum Functionality

Because DFC is the principal low level interface to all Documentum functionality, there are many DFC interfaces that this manual covers only superficially. They provide access to features that other documentation covers in more detail. For example, the *Server Fundamentals* manual describes virtual documents and access control. The DFC Javadocs provide the additional information necessary to enable you to take advantage of those features. Similarly, the Federated Search Services (FS2) product includes extensive capabilities for searching Documentum repositories. The DFC Javadocs provide information about how to use DFC to access that functionality.

This chapter introduces some of the Documentum functionality that you can use DFC to access. It contains the following major sections:

- [Security Services, page 183](#)
- [XML, page 184](#)
- [Virtual Documents, page 184](#)
- [Workflows, page 184](#)
- [Document Lifecycles, page 185](#)
- [Validation Expressions in Java, page 185](#)
- [Unified Login, page 186](#)
- [Search Service, page 186](#)
- [WORM Support for Filestores, page 186](#)

Security Services

Content Server provides a variety of security features. From the DFC standpoint, they fall into the following categories:

- User authentication
Refer to [Chapter 3, Sessions and Session Managers](#) for more information.
- Object permissions

Refer to the *Server Fundamentals* manual and the DFC Javadocs for IDfACL, IDfPermit, and other interfaces for more information.

DFC also provides a feature related to folder permissions. Users may have permission to view an object but not have permission to view all of the folders to which it is linked. The IDfObjectPath interface and the getObjectPaths method of IDfSession provide a powerful and flexible mechanism for finding paths for which the given user has the necessary permissions. Refer to the Javadocs for more details.

XML

[Chapter 6, Working with Document Operations](#) provides some information about working with XML. DFC provides substantial support for the Documentum XML capabilities. Refer to *XML Application Development Guide* for details of how to use these capabilities.

Virtual Documents

[Chapter 6, Working with Document Operations](#) provides some information about working with virtual documents. Refer to *Server Fundamentals* and the DFC Javadocs for the IDfVirtualDocument interface for much more detail.

Workflows

The *Server Fundamentals* manual provides a thorough treatment of the concepts underlying workflows. DFC provides interfaces to support the construction and use of workflows, but there is almost no reason to use those interfaces directly. The workflow manager and business process manager software packages handle all of those details.

Individual workflow tasks can have methods associated with them. You can program these methods in Java and call DFC from them. These methods run on the method server, an application server that resides on the Content Server machine and is dedicated to running Content Server methods. The code for these methods resides in the repository's registry as modules. [Modules and registries, page 126](#) provides more information about registries and modules.

The com.documentum.fc.lifecycle package provides the following interfaces for use by modules that implement lifecycle actions:

- IDfLifecycleUserEntryCriteria to implement userEntryCriteria.
- IDfLifecycleUserAction to implement userAction.
- IDfLifecycleUserPostProcessing to implement userPostProcessing

There is no need to extend DfService, but you can do so. You need only implement IDfModule, because lifecycles are modules, not SBOs.

Document Lifecycles

The *Server Fundamentals* manual provides information about lifecycles. There are no DFC interfaces for constructing document lifecycles. You can define actions to take place at various stages of a document's lifecycle. You can code these in Java to run on the Content Server's method server. Such Java methods must implement the appropriate interfaces from the following:

- IDfLifecycleAction.java
- IDfLifecycleUserAction.java
- IDfLifecycleUserEntryCriteria.java
- IDfLifecycleUserPostProcessing.java
- IDfLifecycleValidate.java

The code for these methods resides in the repository's registry (see [Modules and registries](#), page 126) as modules.

Validation Expressions in Java

When you create a repository type, you can associate constraints with it. Some of these constraints are expressions involving either a single attribute or combinations of attributes of the type. These reside in the repository's data dictionary, where client programs can access them to enforce the constraints. Content Server does not enforce constraints defined in the data dictionary.

In earlier versions of DFC these constraints were exclusively Docbasic expressions. Since DFC 5.3, you can provide Java translations of the constraint expressions. If a Java version of a constraint exists, DFC uses it in preference to the Docbasic version of the same constraint. This usually results in a substantial performance improvement.

The key points about this feature are the following:

- You must take steps to use it.

If you do nothing, DFC continues to use the existing Docbasic expressions. A server script enables the feature by creating the necessary types and the methods for creating Java translations.

- The migration is incremental and non-destructive.

You can migrate all, none, or any portion in between of the Docbasic expression evaluation in a Content Server repository to Java. The Docbasic versions remain in place. You can disable any specific Java translations and revert to using Docbasic for that function or that object type.

- Translations are available for all Docbasic functions that you are likely to use in validation expressions.

We do not provide Java translations of operating system calls, file system access, COM and DDE functions, print or user interface functions, and other similar functions. We do not provide Java translations of financial functions.

Unified Login

If you are using Unified Login with your DFC client there is an important requirement for authentication to work. Unified Login relies on the Windows SID authentication mechanism. DFC forms an SID based on the login info from the client machine and sends it to the server. The server authenticates it using the Windows Domain Controller (DC). If the DC does not recognize the client SID, then authentication will fail. The client SID should be recognized by the DC and the user should be registered using the same domain as the repository. So the requirement is that client and repository machine should be in the same workgroup or domain.

Search Service

The DFC search service replaces prior mechanisms for building and running queries. You can use the IDfQuery interface, which is not part of the search service, for simple queries. The search service provides the ability to run searches across multiple Documentum repositories and, in conjunction with the Federated Search Services (FS2) product, external repositories as well.

The Javadocs for the `com.documentum.fc.client.search` package provide a description of how to use this capability.

WORM Support for Filestores

WORM support is available for filestores.

Run the following command to enable or disable WORM for a filestore:

```
EXECUTE set_storage_state with store='<store>',worm=[true | false]
```

Any new content that is created on a file store with the WORM feature enabled is saved with the read-only flag on the filesystem. When the read-only file attribute is enabled, the filesystem that is WORM-enabled must keep this file under infinite retention. Any edit operation on the content at the filesystem level is denied.

Objects with content stored on the WORM-enabled filestore cannot be deleted or destroyed.

You can perform the `MIGRATE_CONTENT` operation on worm content. However, the `remove_original` parameter is ignored and is considered as `false`.

A

- abort method, 122, 124
- abortTransaction method, 133
- add method. *See* operations, add method
- addNode method, 171
- aspects, 126
- assemblies. *See* virtual documents, assemblies
- assembly objects. *See* virtual documents, assembly objects
- asVirtualDocument method, 84, 99, 112, 118, 171
- authentication
 - Kerberos SSO, 39
 - principal, 44

B

- beginTransaction method, 133
- best practices
 - caching repository objects, 135
 - reusing SBOs, 135
 - SBO state information, 135
- binding. *See* virtual documents, binding

C

- caches, 135
 - See also* persistent caching
- casting, 88, 90, 122 to 123, 150
- children. *See* virtual documents, terminology
- chronicle ID, 176
- chronicle IDs, 83
- chunking rules, 85
- classes
 - extending, 24
 - instantiating, 24
- classpaths, 21
- collections
 - leaks, 18

- COM (Microsoft component object model), 21, 24
- com package, 24
- commitTransaction method, 133
- config directory, 15, 21
- containment objects. *See* virtual documents, containment objects
- copy_child attribute, 84
- CURRENT, 179
- CURRENT version, 166

D

- datatypes, 68
- DBOR (Dokumentum business object registry), 127
- dctm.jar, 21
- Desktop client program, 19
- destroyAllVersions method, 80
- dfc.bof.cache.currency_check_interval, 17
- dfc.bof.cacheconsistency.interval
 - property, 17
- dfc.bof.registry.connect.attempt.interval
 - property, 17
- dfc.config.check_interval property, 19
- dfc.data.dir, 18
- dfc.diagnostics.resources.enable, 18
- dfc.globalregistry.password property, 17
- dfc.globalregistry.repository property, 17
- dfc.globalregistry.username property, 17
- dfc.housekeeping.cleanup.interval
 - property, 19
- dfc.properties, 20
- dfc.properties file, 16
- dfc.registry.mode property, 19
- dfc.session.allow_trusted_login, 20
- dfcfull.properties file, 16
- DfClientX class, 25
- DfCollectionEx class, 131
- DfDborNotFoundException class, 134
- DfException class, 26, 122

- DfNoTransactionAvailableException
 - class, 132
- DfService class, 126, 129 to 130
- DfServiceCriticalException class, 134
- DfServiceException class, 134
- DfServiceInstantiationException class, 134
- DfServiceNotFoundException class, 134
- directories (file system), 105
- dmc_module type, 126
- DMCL process, 25
- docbroker, 18
- document manipulation. *See* operations
- DONT_RECORD_IN_REGISTRY field, 106
- DTDs (document type definitions), 118
- dump method, 64

E

- edges. *See* virtual documents, terminology
- enableDeepDeleteVirtualDocumentsInFolders
 - method, 114
- error handlers, 26, 122
- execute method. *See* operations, execute method
- External applications
 - XML support, 85

F

- factory methods, 23 to 24, 129
- Federated Search Services (FS2), 19
- FileOutputStream class, 119

G

- getCancelCheckoutOperation method, 99
- getCheckinOperation method, 96
- getChildren method, 89
- getDefaultDestinationDirectory
 - method, 108
- getDeleteOperation method, 114
- getDestinationDirectory method, 108
- getDocbaseNameFromId method, 131
- getErrors method, 89, 118, 122
- getExportOperation method, 106
- getFilePath method, 108
- getId method, 90
- getImportOperation method, 103
- getLocalClient method, 25, 150
- getMoveOperation method, 112
- getNewObjects method, 98, 103, 105

- getNodes method, 89
- getObjectID method, 90
- getSelectedObject method, 174
- getSession method, 25, 130, 132, 134, 150
- getSessionManager method, 130, 150
- getStatistics method, 132
- getTransactionRollbackOnly method, 132
- getValidationOperation method, 118
- getVendorString method, 129
- getVersion method, 129
- getXMLTransformOperation method, 119 to 121
- getYesNoAnswer method, 124
- global registry, 17

H

- HTML (Hypertext Markup Language), 120 to 121

I

- IDfBusinessObject interface, 126
- IDfCancelCheckoutNode interface, 99
- IDfCancelCheckoutOperation interface, 99
- IDfCheckinNode interface, 96
- IDfCheckinOperation interface, 96
- IDfCheckoutOperation interface, 92
- IDfClient interface, 23, 25
- IDfClientX interface, 25, 86
- IDfDeleteNode interface, 114
- IDfDeleteOperation interface, 114
- IDfDocument interface, 25, 90
- IDfExportNode interface, 108
- IDfExportOperation interface, 106
- IDfFile interface, 105
- IDfFolder interface, 58, 90
- IDfImportNode interface, 103, 105
- IDfImportOperation interface, 103, 120
- IDfList interface, 122
- IDfLoginInfo interface, 150
- IDfModule interface, 126
- IDfMoveOperation interface, 112
- IDfOperation interface, 87
- IDfOperationError interface, 89, 122
- IDfOperationMonitor interface, 123
- IDfOperationNode interface, 88 to 89
- IDfPersistentObject interface, 24 to 25
- IDfProperties interface, 120
- IDfService interface, 126, 129

IDfSession interface, 23, 25
 IDfSessionManagerStatistics interface, 132
 IDfSysObject interface, 24
 IDfValidationOperation interface, 118
 IDfVersionLabels, 169
 IDfXMLTransformNode interface, 119 to 121
 IDfXMLTransformOperation interface, 119 to 121
 ignore (to turn off XML processing), 105
 interface inheritance, 24, 26
 interfaces, 24
 isCompatible method, 129
 isTransactionActive method, 132 to 133

J

Java language, 21, 24
 setup, 21
 supported versions, 15
 java.util.Properties class, 16
 javac compiler, 21
 Javadocs. *See* online reference documentation

K

Kerberos authentication, 39

L

leaks, 18
 local files, 98
 log4j.properties, 20

M

manifests, 21
 metadata, 57
 modules, 126
 Modules folder, 126

N

naming conventions, 23, 129, 134
 .NET platform, 21, 125
 newObject method, 25
 newService method, 129, 150
 newSessionManager method, 25, 150
 NEXT_MAJOR field, 96

nodes. *See* virtual documents, terminology;
 operations, nodes
 null returns, 122

O

OLE (object linking and embedding)
 links, 103
 OLE (Object Linking and Embedding)
 links, 20
 online reference documentation, 16
 operation monitors, 123
 operations
 aborting, 122
 add method, 88, 122
 cancel checkout, 99
 checkin, 96
 checkout, 92
 delete, 114
 errors, 122
 execute method, 89, 122
 export, 106
 factory methods, 86
 import, 103
 move, 112
 nodes, 88 to 89
 parameters, 88
 procedure for using, 87
 steps, 89
 transactions, 124
 validate, 118
 operations package, 24 to 25, 85
 orphan documents, 98

P

packages, 23
 parents. *See* virtual documents, terminology
 Predictive caching, 117
 principal authentication, 44
 progressReport method, 124

R

Reader class, 121
 release method, 25
 releaseSession method, 130, 150
 renditions, 121
 reportError method, 124
 repositories, 126

S

SBOs. *See* service based objects
schemas. *See* XML schemas
service based objects (SBOs), 128
 architecture, 129
 implementing, 129
 instantiating, 135, 149
 returning TBOs, 150
 session manager, 150
 specifying a repository, 131
 threads, 149
 transactions, 132
session leaks, 18
session managers, 23, 25, 129
 internal statistics, 132
 transactions, 133
setBinding, 176
setCheckinVersion method, 96
setDestination method, 119 to 121
setDestinationDirectory method, 118
setDestinationFolderId method, 103, 105, 120
setDomain method, 150
setFilePath method, 106
setIdentity method, 150
setKeepLocalFile method, 99
setOperationMonitor method, 123
setOutputFormat method, 120 to 121
setPassword method, 150
setRecordInRegistry method, 106, 108
setSession method, 103, 105, 119 to 121
setSessionManager method, 46, 150 to 151
setTransactionRollbackOnly method, 132
setTransformation method, 119 to 121
setUser method, 150
setXMLApplicationName method, 103, 105
simple modules. *See* modules
state information, 46, 129 to 131

T

threads, 133
transactions
 nested, 133
type based objects (TBOs)
 returning from SBO, 150

U

URLs (universal resource locators), 121

V

version trees, 83
virtual documents, 83
 assemblies, 84
 assembly objects, 84
 asVirtualDocument method, 84, 99, 112, 118, 171
 binding, 84
 cancelling checkout, 99, 101
 containment objects, 84
 copy behavior, 84
 CURRENT version, 166
 deleting, 114
 exporting, 106
 Microsoft OLE links, 165
 terminology, 84
 versioning, 84
 XML applications, 165

X

Xalan transformation engine, 119
Xerces XML parser, 118
XML schemas, 118
XML support, 85, 118 to 119
 See also ignore
XSLT stylesheets, 119 to 121